

LTTng

The *Linux Trace Toolkit: next generation* is an open source software toolkit which you can use to simultaneously trace the Linux kernel, user applications, and user libraries.

<http://lttng.org/features/>

LTTng supports architectures such as IA-32 (x86), x86-64, PowerPC, ARM, and MIPS, among others. ...

[*The LTTng Documentation*](#)

<http://lttng.org/docs/v2.10/#doc-tracing-your-own-user-application>

[Howto tracing with LTTng](#) – tutorial covers installation and basic usage.

Kernel Trace with LTTng

Source: <http://lttng.org/docs/v2.10/#doc-tracing-the-linux-kernel>

Before you follow this guide, make sure to [install](#) LTTng. [including kernel modules for LTTng kernel tracing].

...

[Basic required pkgs:

lttng-tools

lttng-modules-dkms

]

The following command lines start with the # prompt because you need root privileges to trace the Linux kernel. You can also trace the kernel as a regular user if your Unix user is a member of the [tracing group](#).

1. Create a [tracing session](#) which writes its traces to /tmp/my-kernel-trace:

```
# lttng create my-kernel-session --output=/tmp/my-kernel-trace
```

2. List the available kernel tracepoints and system calls:

```
# lttng list --kernel
```

```
# lttng list --kernel --syscall
```

3. Create [event rules](#) which match the desired instrumentation point names, for example the `sched_switch` and `sched_process_fork` tracepoints, and the [open\(2\)](#) and [close\(2\)](#) system calls:

```
# lttng enable-event --kernel sched_switch,sched_process_fork
# lttng enable-event --kernel --syscall open,close
```

You can also create an event rule which matches *all* the Linux kernel tracepoints (this will generate a lot of data when tracing):

```
# lttng enable-event --kernel --all
```

4. [Start tracing](#):

```
# lttng start
```

5. Do some operation on your system for a few seconds. For example, load a website, or list the files of a directory.

6. [Stop tracing](#) and destroy the tracing session:

```
# lttng stop
# lttng destroy
```

The [lttng-destroy\(1\)](#) command does not destroy the trace data; it only destroys the state of the tracing session.

7. For the sake of this example, make the recorded trace accessible to the non-root users:

```
# chown -R $(whoami) /tmp/my-kernel-trace
```

See [View and analyze the recorded events](#) to view the recorded events.

Also See:

- [How to use LTTng to diagnose problems](#)
- <http://archive.eclipse.org/tracecompass/doc/stable/org.eclipse.tracecompass.doc.user/LTTng-Kernel-Analysis.html>
- LTTng built-in help system: <http://127.0.0.1:42105/help/index.jsp>

A lttng kernel trace session example

```
# lttng create k1-lttng --output=/tmp/k1-lttng.trc
# lttng list --kernel << see all possible kernel events >>
Kernel events:
-----
    scsi_dispatch_cmd_start (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    scsi_dispatch_cmd_error (loglevel: TRACE_EMERG (0)) (type: tracepoint)
```

```
...
    sched_pi_setprio (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    signal_generate (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    signal_deliver (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    skb_kfree (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    skb_consume (loglevel: TRACE_EMERG (0)) (type: tracepoint)
...
    block_rq_issue (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    block_bio_bounce (loglevel: TRACE_EMERG (0)) (type: tracepoint)
...
    gpio_value (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    irq_handler_entry (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    irq_handler_exit (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    irq_softirq_entry (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    irq_softirq_exit (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    irq_softirq_raise (loglevel: TRACE_EMERG (0)) (type: tracepoint)
...
    kmem_kmalloc (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    kmem_cache_alloc (loglevel: TRACE_EMERG (0)) (type: tracepoint)
...
    kvm_async_pf_completed (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    module_load (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    module_free (loglevel: TRACE_EMERG (0)) (type: tracepoint)
...
    writeback_wait_iff_congested (loglevel: TRACE_EMERG (0)) (type: tracepoint)
    writeback_single_inode (loglevel: TRACE_EMERG (0)) (type: tracepoint)
#
# lttng enable-event --kernel --all    << trace all! Big files.. >>
# lttng start ; sleep 1; lttng stop
```

Viewing the trace

```
# babeltrace /tmp/k1-lttng.trc/kernel/
[11:40:08.863740605] (+?.????????) kaiwan-T460 kmem_kmalloc: { cpu_id = 0 },
{ call_site = 0xFFFFFFFFC10AA457, ptr = 0xFFFFF930CC5098000, bytes_req = 739, bytes_alloc
= 1024, gfp_flags = 20971712 }
[11:40:08.863747488] (+0.000006883) kaiwan-T460 kmem_kfree: { cpu_id = 0 }, { call_site
= 0xFFFFFFFFC10AA52A, ptr = 0xFFFFF930CC5098000 }
[11:40:08.863749701] (+0.000002213) kaiwan-T460 sched_waking: { cpu_id = 0 }, { comm =
"lttng-consumerd", tid = 32437, prio = 20, target_cpu = 3 }
[11:40:08.863752608] (+0.000002907) kaiwan-T460 sched_wakeup: { cpu_id = 0 }, { comm =
"lttng-consumerd", tid = 32437, prio = 20, target_cpu = 3 }
[11:40:08.863754831] (+0.000002223) kaiwan-T460 power_cpu_idle: { cpu_id = 3 }, { state
= 4294967295, cpu_id = 3 }
[11:40:08.863756344] (+0.000001513) kaiwan-T460 kmem_kmalloc: { cpu_id = 0 },
{ call_site = 0xFFFFFFFFC10AA457, ptr = 0xFFFFF930C7FAEAA00, bytes_req = 275, bytes_alloc
= 512, gfp_flags = 20971712 }
...

[11:40:09.948720350] (+0.000001293) kaiwan-T460 power_cpu_idle: { cpu_id = 1 }, { state
= 1, cpu_id = 1 }
[11:40:09.948720370] (+0.000000020) kaiwan-T460 syscall_entry_fcntl: { cpu_id = 0 },
{ fd = 46, cmd = 2, arg = 1 }
[11:40:09.948721513] (+0.000001143) kaiwan-T460 syscall_exit_fcntl: { cpu_id = 0 },
{ ret = 0, arg = 1 }
[11:40:09.948722298] (+0.000000785) kaiwan-T460 syscall_entry_setsockopt: { cpu_id =
0 }, { fd = 46, level = 1, optname = 16, optval = 139758116252252, optlen = 4 }
[11:40:09.948723921] (+0.000001623) kaiwan-T460 syscall_exit_setsockopt: { cpu_id = 0 },
{ ret = 0 }
[11:40:09.948728342] (+0.000004421) kaiwan-T460 syscall_entry_recvmsg: { cpu_id = 0 }, {
fd = 46, msg = 139758116252128, flags = 0 }
```

...

#

It's all there, just not very human readable...

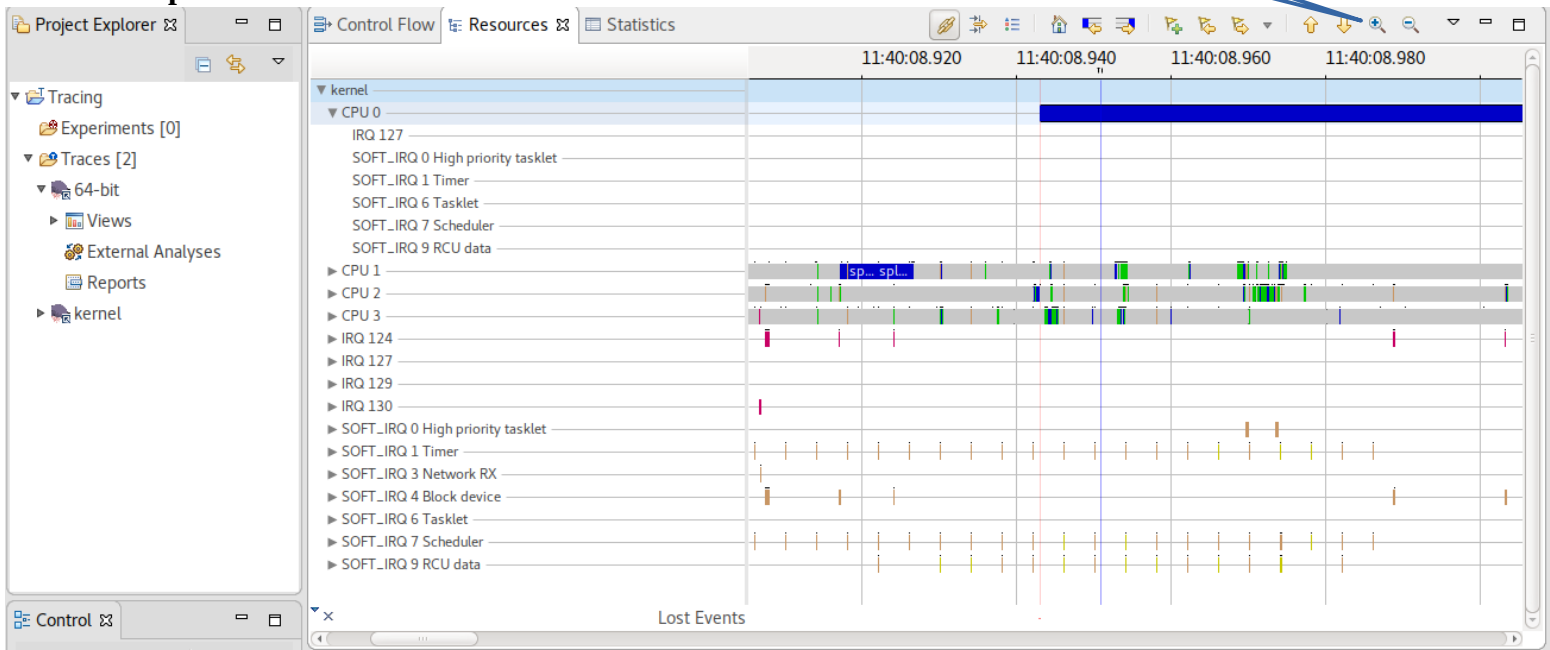
So we use the [TraceCompass](#) *GUI*. Much, much better!

Using the TraceCompass GUI!

Doc: <https://archive.eclipse.org/tracecompass/doc/stable/org.eclipse.tracecompass.doc.user/LTTng-Kernel-Analysis.html>

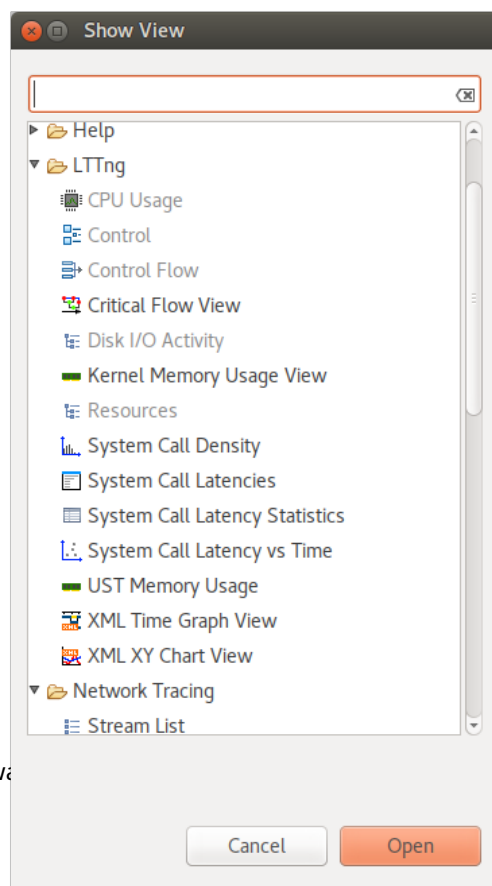
Tips:


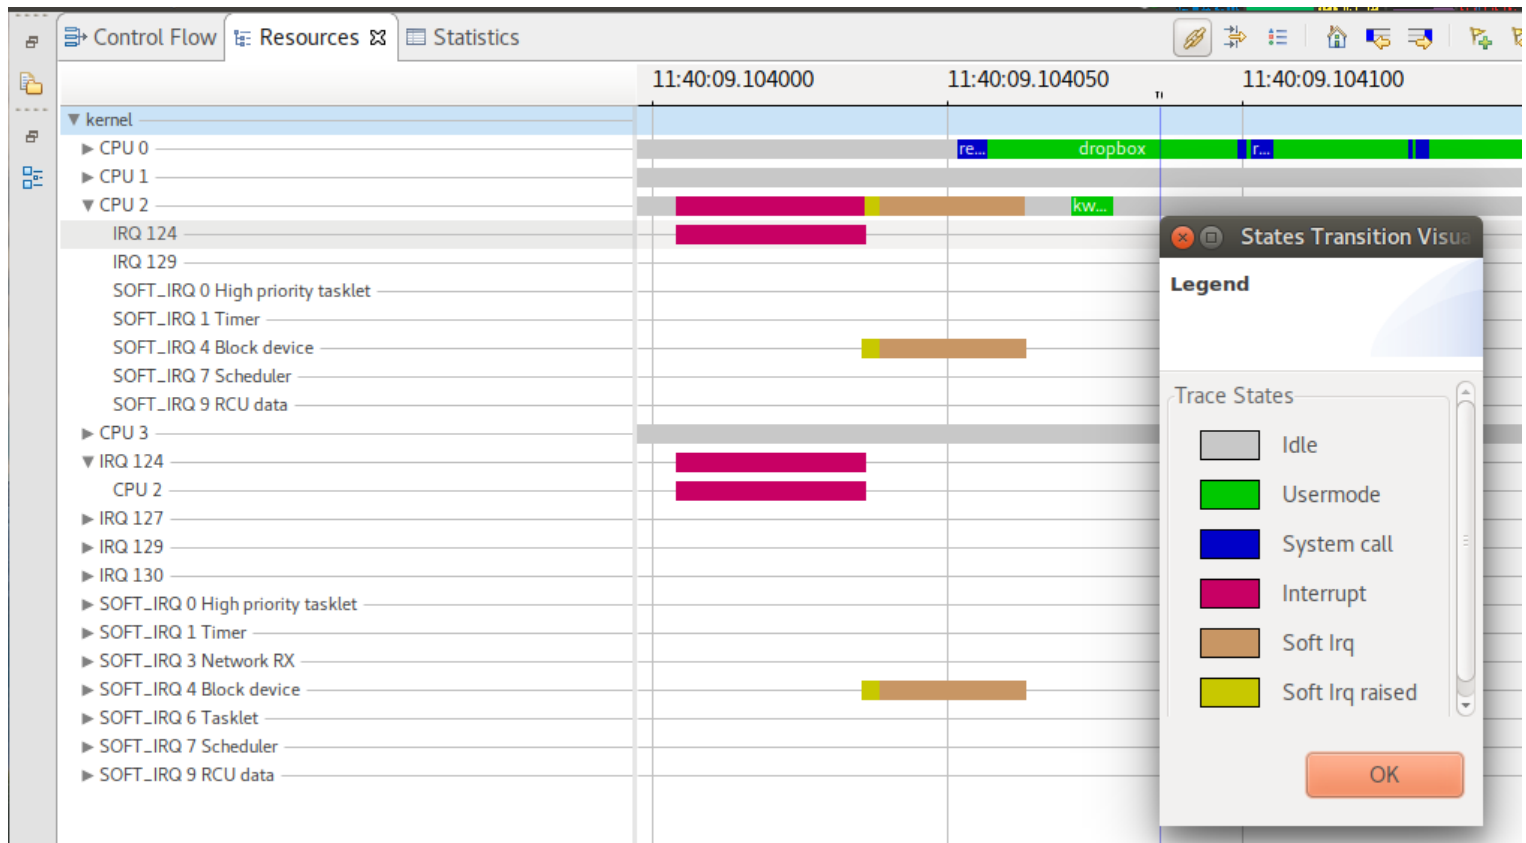
Zoom in/out
on timeline



Window / Show Perspective / LTTng Kernel (default)


Window / Show View :



'Control Flow' tab**'Resources' tab**


(Tip: in the first screenshot, an arrow can be seen; it represents a scheduling switch. From the [docs](#): “... The arrows indicate when the scheduler switches from one process to another for a given CPU...”).

(On a project: seeing CPU utilization by the web server (lighttpd) and the ‘cgi’ scripts that it runs)



(We can perhaps think: why's 'ps' taking an inordinately long time here?)

‘Statistics’ tab



Below: we're following the *ping* process; we have found it issuing the `sendto()` syscall. Within it, spotted the `kmem_cache_alloc_node()`; the return pointer is seen here (obtained by rt-click on the line / *Copy to Clipboard*):

Timestamp	Channel	CPU	Event type	Contents
<srch>	<srch>	<srch>	<srch>	0xffff90a4dbfe1200
16:54:10.364 297 744	channel0_0 0	0	sched_migrate_task	comm=EMT, tid=10202, prio=20, orig_cpu=1, dest_cpu=2
16:54:10.364 302 685	channel0_1 1	1	kmem_cache_alloc_node	call_site=0xfffffffff8774b227, ptr=0xffff90a4dbfe1200, bytes_req=232, bytes_alloc=256, gfp_flags=20971712
16:54:10.364 306 088	channel0_1 1	1	kmem_kmalloc_node	call_site=0xfffffffff8774b252, ptr=0xffff90a627532200, bytes_req=448, bytes_alloc=512, gfp_flags=21037760
16:54:10.364 306 713	channel0_0 0	0	sched_wakeup	comm=EMT, tid=10202, prio=20, target_cpu=2
16:54:10.364 313 386	channel0_0 0	0	kmem_kfree	call_site=0xffffffffc1240e2f, ptr=0xffff90a79259f200
16:54:10.364 314 613	channel0_2 2	2	power_cpu_idle	state=4294967295, cpu_id=2
16:54:10.364 315 232	channel0_1 1	1	kmem_kfree	call_site=0xfffffffff877ba73e, ptr=0x0
16:54:10.364 317 630	channel0_0 0	0	syscall_exit_ioctl	ret=0, arg=140451819113536
16:54:10.364 324 011	channel0_2 2	2	timer_hrtimer_cancel	hrtimer=0xffff90a80251b5a0
16:54:10.364 324 850	channel0_0 0	0	syscall_entry_rt_sigtimedwait	uts=0, sigsetsize=8
16:54:10.364 326 971	channel0_2 2	2	timer_hrtimer_start	hrtimer=0xffff90a80251b5a0, function=0xfffffffff871308f0, expires=51664519000000, softexpires=51664519
16:54:10.364 332 046	channel0_0 0	0	rcu_utilization	s=Start context switch
16:54:10.364 333 553	channel0_0 0	0	rcu_utilization	s=End context switch
16:54:10.364 335 303	channel0_2 2	2	rcu_utilization	s=Start context switch
16:54:10.364 336 963	channel0_2 2	2	rcu_utilization	s=End context switch
16:54:10.364 337 975	channel0_0 0	0	sched_stat_runtime	comm=Timer, tid=10214, runtime=143279, vruntime=313237138711
16:54:10.364 344 080	channel0_2 2	2	sched_switch	prev_comm=swapper/2, prev_tid=0, prev_prio=20, prev_state=0, next_comm=EMT, next_tid=10202, next_pri
16:54:10.364 349 531	channel0_1 1	1	net_dev_queue	skbaddr=0xffff90a4dbfe1200, len=98, name=enp0s31f6, network_header_type=_ipv4, network_header=ipv4

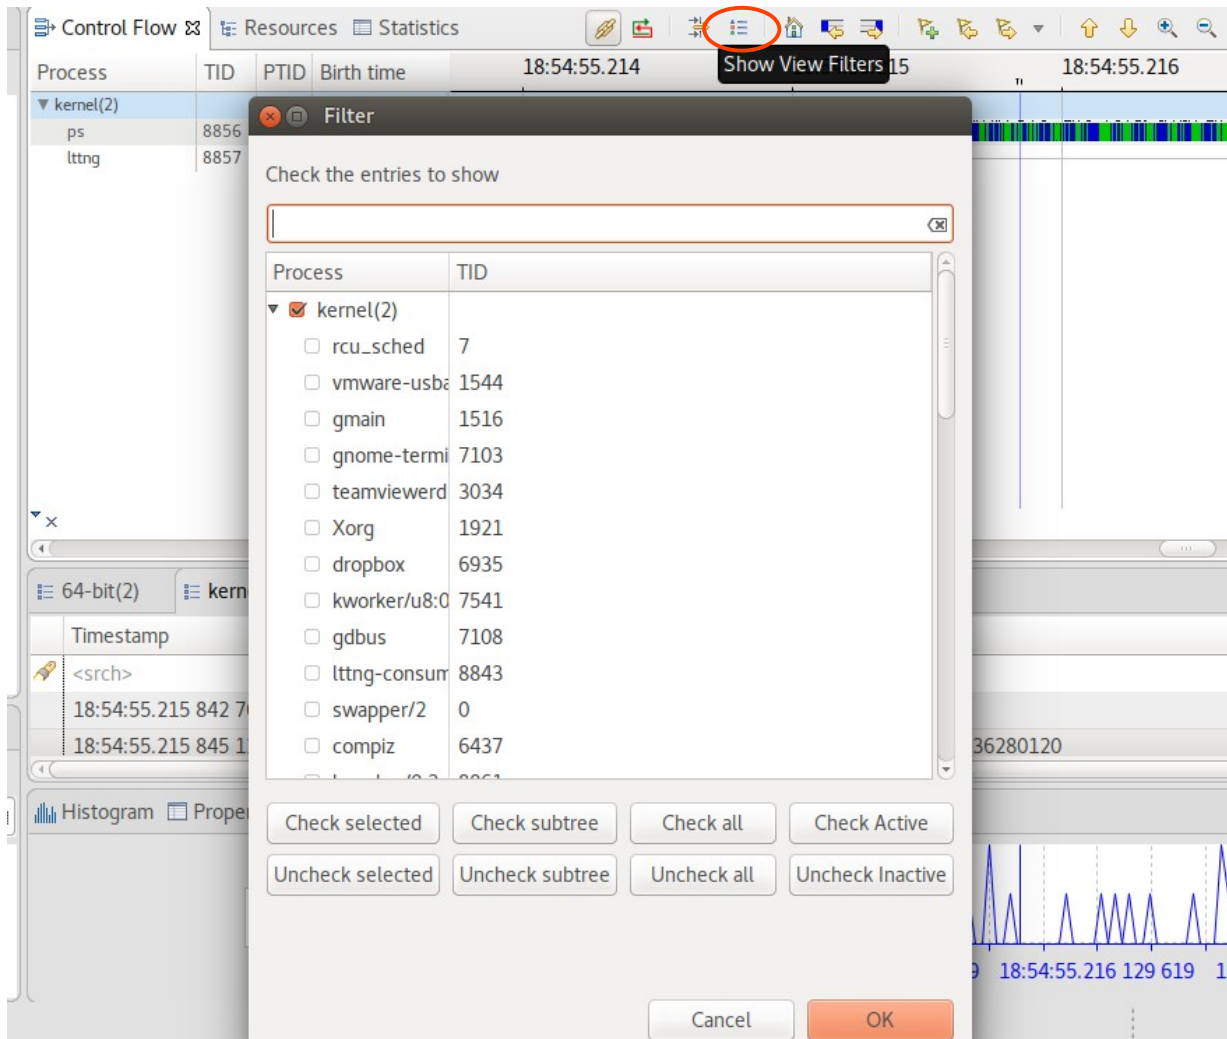
```
Timestamp      Channel CPU   Event type  Contents      TID   Prio  PID
16:54:10.364 302 685      channel0_1  1    kmem_cache_alloc_node
               call_site=0xfffffffff8774b227, ptr=0xffff90a4dbfe1200, bytes_req=232,
bytes_alloc=256, gfp_flags=20971712, node=-1    6737  20    6737
```

The ptr value is filtered upon (above screenshot); we can now literally *see* for ourselves where it's used next (its highlighted); in `net_dev_queue()`; following it down further, one comes to the free call site (the `kmem_cache_free()` on this same ptr)!

Also, the parameters are displayed ! Note that they aren't necessarily parameters, but very very useful information!

Also:

Click the "Show View Filters" button; a dialog pops up allowing one to filter entries. Select entries of interest and check them out! (in the screenshot below, among processes alive at the time of the trace, we selected only 'ps' and 'ltnng'):



Additional CLI Profiling Tools

Very powerful CLI tools available too!

Check this out:

```
$ lttng<tab><tab>
lttng          lttng-iolatencytop      lttng-irqlog-mi
lttng-periodstats  lttng-schedstats
lttng-analyses-record  lttng-iolatencytop-mi  lttng-irqstats
lttng-periodstats-mi  lttng-schedstats-mi
lttng-cputop        lttng-iolog              lttng-irqstats-mi
lttng-periodtop      lttng-schedtop
lttng-cputop-mi       lttng-iolog-mi           lttng-memtop
lttng-periodtop-mi    lttng-schedtop-mi
lttng-crash           lttng-iousagetop         lttng-memtop-mi
lttng-relayd          lttng-sessiond
lttng-iolatencyfreq   lttng-iousagetop-mi     lttng-periodfreq
lttng-schedfreq       lttng-syscallstats
lttng-iolatencyfreq-mi lttng-irqfreq           lttng-periodfreq-mi
lttng-schedfreq-mi    lttng-syscallstats-mi
lttng-iolatencystats  lttng-irqfreq-mi        lttng-periodlog
lttng-schedlog        lttng-track-process
lttng-iolatencystats-mi lttng-irqlog            lttng-periodlog-mi
lttng-schedlog-mi

    << mi = machine interface >>
$ lttng^C
```

```
$ lttng-cputop /tmp/k1-lttng.trc/
```

Checking the trace for lost events...

```
[warning] Tracer discarded 12989 events between [11:40:08.873895341] and
[11:40:08.943080657] in trace UUID 72e8b8b633cc44f90422b6f534500, at path: "/tmp/k1-
lttng.trc/kernel", within stream id 0, at relative path: "channel0_0". You should
consider recording a new trace with larger buffers or with fewer events enabled.
[warning] Tracer discarded 992 events between [11:40:08.943080657] and
[11:40:09.485410272] in trace UUID 72e8b8b633cc44f90422b6f534500, at path: "/tmp/k1-
lttng.trc/kernel", within stream id 0, at relative path: "channel0_0". You should
consider recording a new trace with larger buffers or with fewer events enabled.
```

Processing the trace: 100%

```
[#####] Time: 0:00:01
```

```
#####] Time: 0:00:01
```

Timerange: [2017-08-25 11:40:08.863740605, 2017-08-25 11:40:09.941728585]

Per-TID Usage

Process	Migrations	Priorities	
lttng-consumerd (32438)	1	[20]	4.81 %
indicator-multi (3794)	2	[20]	4.13 %
(3707)	5	[]	1.80 %
dropbox (4357)	5	[20]	1.77 %
kworker/u8:1 (26729)	2	[20]	1.29 %
Xorg (1875)	0	[20]	0.81 %

lttng-consumerd (32437)	2	[20]	0.46 %
teamviewerd (2632)	5	[20]	0.26 %
chrome (4494)	3	[20]	0.22 %
unity-panel-ser (3715)	1	[20]	0.20 %

Per-CPU Usage

```
#####  
CPU 0 24.62 %  
CPU 1 4.04 %  
CPU 2 6.44 %  
CPU 3 6.01 %
```

Total CPU Usage: 10.28%

\$

\$ lttng-memtop /tmp/k1-lttng.trc/

Checking the trace for lost events...

[warning] Tracer discarded 12989 events between [11:40:08.873895341] and [11:40:08.943080657] in trace UUID 72e8b8b633cc44f90422b6f534500, at path: "/tmp/k1-lttng.trc/kernel", within stream id 0, at relative path: "channel0_0". You should consider recording a new trace with larger buffers or with fewer events enabled.

[warning] Tracer discarded 992 events between [11:40:08.943080657] and [11:40:09.485410272] in trace UUID 72e8b8b633cc44f90422b6f534500, at path: "/tmp/k1-lttng.trc/kernel", within stream id 0, at relative path: "channel0_0". You should consider recording a new trace with larger buffers or with fewer events enabled.

Processing the trace: 100%

```
#####  
#####] Time: 0:00:01
```

Timerange: [2017-08-25 11:40:08.863740605, 2017-08-25 11:40:09.941728585]

Per-TID Memory Allocations

Process

```
#####  
pages lttng-consumerd (32438) 1024  
pages dropbox (4357) 233  
pages bash (32463) 50  
pages Xorg (1875) 25  
pages indicator-multi (3794) 21  
pages lttng-consumerd (32437) 18  
pages in:imklog (1615) 9  
pages rs:main Q:Reg (1616) 8  
pages lttng (32462) 3  
pages NetworkManager (1404) 3
```

Per-TID Memory Deallocations

Process

```
#####
```

```

[REDACTED] 960
pages lttng-consumerd (32438)
[REDACTED] 146
pages lttng (32462)
[REDACTED] 54
pages bash (32463)
[REDACTED] 21
pages indicator-multi (3794)
[REDACTED] 15
pages lttng-consumerd (32437)
[REDACTED] 4
pages kworker/2:3 (26920)
[REDACTED] 3
pages NetworkManager (1404)
[REDACTED] 2
pages jbd2/sda6-8 (232)
[REDACTED] 2
pages wpa_supplicant (1977)
[REDACTED] 2
pages gdbus (3258)

Total memory usage:
- 1405 pages allocated
- 1214 pages freed
$
```

Similarly, try out the [lttng-iolatencytop](#) / [lttng-iousetop](#) / [lttng-schedtop](#) / [lttng-periodtop](#) tools.


Userspace Tracing

Requires adding instrumentation and tracepoints into the project codebase.

UST = User Space Tracer

[***lttng-ust***](#) — [Linux Trace Toolkit Next Generation User-Space Tracer 2.x](#)

Linux Operating System Specialized




kaiwanTECH


The highest quality Training on:

Linux Fundamentals, CLI and Scripting
Linux Systems Programming
Linux Kernel Internals
Linux Device Drivers
Embedded Linux
Linux Debugging Techniques
New! *Linux OS for Technical Managers*

Please do visit our website for details:

<http://kaiwantech.in>



<https://kaiwantech.com>

kaiwanTECH Linux OS Corporate Training Programs

Please do check out our current offering of world-class, seriously-valuable, high on returns, technical Linux OS corporate training programs [here](#).