



Sanitizing the Linux kernel

On KASAN and other
Dynamic Bug-finding Tools

Andrey Konovalov, xairy.io

Linux Security Summit Europe
September 16th 2022

Fuzzing is awesome?

Fuzzing is useless

Fuzzing is useless★

★without dynamic bug detectors



Sanitizing the Linux kernel

On KASAN and other
Dynamic Bug-finding Tools

Andrey Konovalov, xairy.io

Linux Security Summit Europe
September 16th 2022

Agenda

1. What are Sanitizers
2. KASAN and its Generic mode
3. More Sanitizers: KMSAN, KCSAN, and UBSAN
4. In-field Sanitizers: KFENCE and SW_TAGS KASAN
5. Sanitizers as a mitigation: HW_TAGS KASAN
6. Extending Sanitizers on the example of KASAN and KMSAN
7. Notes on implementing custom bug detectors

Sanitizers

- A family of bug detectors
- Initially implemented for userspace (ASan, MSan, etc.)
- Later ported to the Linux kernel (KASAN, KMSAN, etc.)
- Features:
 - Easy to use
 - Fast
 - Precise: no false positives
 - Detailed reports

Generic KASAN

KASAN – Kernel Address Sanitizer

- Dynamic memory corruption detector for the Linux kernel
- Finds out-of-bounds, use-after-free, and double/invalid-free bugs
- Supports slab, page_alloc, vmalloc, stack, and global memory
- Requires compiler support: implemented in both Clang and GCC
- Has 3 modes, we'll focus on Generic mode (CONFIG_KASAN_GENERIC)
- google.github.io/kernel-sanitizers/KASAN

KASAN parts

- Compiler module ([llvm/lib/Transforms/Instrumentation/AddressSanitizer.cpp](#))
- Runtime part ([mm/kasan/](#) + [include/linux/kasan.h](#) + ...)

KASAN parts

- Compiler module ([llvm/lib/Transforms/Instrumentation/AddressSanitizer.cpp](https://llvm.org/docs/HowToBuildKASAN.html))
 - Instruments memory accesses when building the kernel
 - Inserts redzones for stack and global variables
- Runtime part ([mm/kasan/](#) + [include/linux/kasan.h](#) + ...)


KASAN parts

- Compiler module ([llvm/lib/Transforms/Instrumentation/AddressSanitizer.cpp](https://llvm.org/docs/HowToBuildKASAN.html))
 - Instruments memory accesses when building the kernel
 - Inserts redzones for stack and global variables
- Runtime part ([mm/kasan/](#) + [include/linux/kasan.h](#) + ...)
 - Maintains shadow memory to track memory state
 - Hooks into kernel allocators to track alloc/free events
 - Prints bug reports

Shadow memory


- Compiler module
 - Instruments memory accesses when building the kernel
 - Inserts redzones for stack and global variables
- Runtime part
 - Maintains **shadow memory** to track memory state
 - Hooks into kernel allocators to track alloc/free events
 - Prints bug reports

Shadow byte



Any 8 aligned bytes usually have only 9 states:

N good bytes and $8 - N$ bad bytes ($0 \leq N \leq 8$)



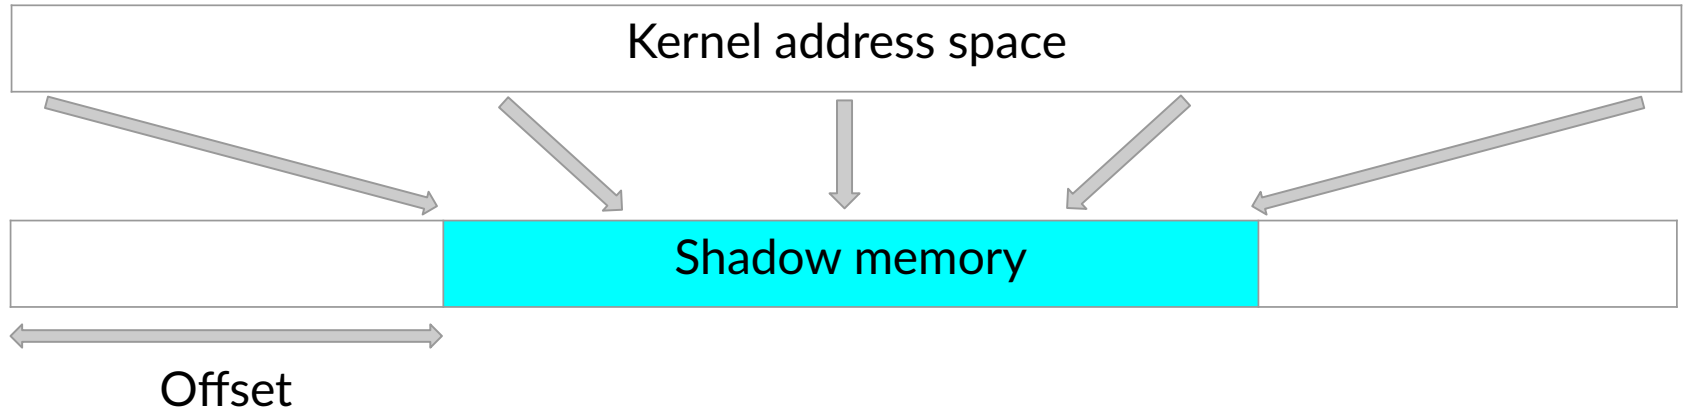
Shadow byte values for inaccessible memory

```
#define KASAN_PAGE_FREE          0xFF  /* freed page */
#define KASAN_PAGE_REDZONE      0xFE  /* redzone for kmalloc_large allocation */
#define KASAN_SLAB_REDZONE      0xFC  /* redzone for slab object */
#define KASAN_SLAB_FREE         0xFB  /* freed slab object */
#define KASAN_VMALLOC_INVALID   0xF8  /* inaccessible space in vmalloc area */
#define KASAN_SLAB_FREETRACK    0xFA  /* freed slab object with free track */
#define KASAN_GLOBAL_REDZONE    0xF9  /* redzone for global variable */
#define KASAN_STACK_LEFT       0xF1
#define KASAN_STACK_MID        0xF2
#define KASAN_STACK_RIGHT      0xF3
#define KASAN_STACK_PARTIAL     0xF4
```

Shadow memory region

- Contains shadow bytes for each mapped region of kernel memory
- Memory-to-shadow mapping scheme:

$$\text{Shadow} = (\text{Addr} \gg 3) + \text{Offset}$$



x86-64 kernel memory layout (4-level page tables)

...

ffff800000000000	ffff87ffffffffffff	8 TB	... guard hole, also reserved for hpv.
ffff880000000000	ffff887ffffffffffff	0.5 TB	LDT remap for PTI
ffff888000000000	ffffc87ffffffffffff	64 TB	mapping of phys. memory (page_offset_base)
ffffc88000000000	ffffc87ffffffffffff	0.5 TB	... unused hole
ffffc90000000000	ffffe87ffffffffffff	32 TB	vmalloc/ioremap space (vmalloc_base)
ffffe90000000000	ffffe97ffffffffffff	1 TB	... unused hole
ffffea0000000000	ffffea7ffffffffffff	1 TB	virtual memory map (vmemmap_base)
ffffeb0000000000	ffffeb7ffffffffffff	1 TB	... unused hole
ffffec0000000000	 fffffb7ffffffffffff	 16 TB	 KASAN shadow memory

...

Shadow mapping routines for x86-64

- During early boot, shadow [mapped to a zero page](#)
- Once page tables are initialized, proper shadow [is mapped](#)
- After boot, shadow [gets mapped](#) for vmalloc/vmap areas

Instrumentation of memory accesses

- Compiler module
 - **Instruments memory accesses** when building the kernel
 - Inserts redzones for stack and global variables
- Runtime part
 - Maintains shadow memory to track memory state
 - Hooks into kernel allocators to track alloc/free events
 - Prints bug reports

Compiler instrumentation: 8-byte access

```
*a = ...;
```



```
char *shadow = (a >> 3) + Offset;  
if (*shadow)  
    kasan_report(a);  
*a = ...;
```

Instrumentation: N-byte access (N = 1, 2, 4)

```
*a = ...;
```



```
char *shadow = (a >> 3) + Offset;  
if (*shadow && *shadow < (a & 7) + N)  
    kasan_report(a);  
*a = ...;
```

Allocators hooks

- Compiler module
 - Instruments memory accesses when building the kernel
 - Inserts redzones for stack and global variables
- Runtime part
 - Maintains shadow memory to track memory state
 - **Hooks into kernel allocators** to track alloc/free events
 - Prints bug reports

Allocators hooks

- KASAN need to keep shadow up-to-date
- This requires tracking of alloc/free events
- KASAN adds hooks to kernel allocators
 - [SLUB](#)/[SLAB](#), [page_alloc](#), [vmalloc](#) (grep code for "kasan_")


Slab layout (for SLUB)

- Slab is fully poisoned (marked as inaccessible) when allocated
- Objects in slab are unpoisoned when allocated
- Always-poisoned redzones between objects (\Rightarrow fewer objects fit in slab)

Slab layout without KASAN:




Slab layout with KASAN:



Additional redzones for kmalloc'ed objects

- kmalloc chooses the best fitting cache to serve objects
 - An allocation of 100 bytes is served from kmalloc-128
- KASAN poisons the unused tail (but [unpoisons](#) if ksize is called)



Quarantine for freed memory

- When memory is freed, it's typically immediately reallocated
 - ⇒ Detecting use-after-free is hard
- KASAN implements [quarantine](#) for slab objects
 - Freed objects are not returned to allocator immediately
 - Instead, they are put into a delayed reuse queue
 - ⇒ Higher chance to detect use-after-free

Redzones for stack and globals

- Compiler module
 - Instruments memory accesses when building the kernel
 - Inserts **redzones for stack and global variables**
- Runtime part
 - Maintains shadow memory to track memory state
 - Hooks into kernel allocators to track alloc/free events
 - Prints bug reports

Compiler instrumentation: stack variables [1/3]

```
void foo() {
```

```
    char x[10];
```

```
    <----- Original function code ----->
```

```
}
```

Compiler instrumentation: stack variables [2/3]

```
void foo() {  
    char rz1[32];  
    char x[10];  
    char rz2[22];
```

```
<----- Original function code ----->
```

```
}
```

Compiler instrumentation: stack variables [3/3]

```
void foo() {  
    char rz1[32];  
    char x[10];  
    char rz2[22];  
    // Compiler-inserted code to unpoison x and poison redzones.  
    <----- Original function code ----->  
    // Compiler-inserted code to unpoison stack frame.  
}
```

Compiler instrumentation: global variables

```
char x[10];
```



```
struct {  
    char original[10];  
    char redzone[54]; // Poisoned via constructors  
} x;
```

Reports for found bugs

- Compiler module
 - Instruments memory accesses when building the kernel
 - Inserts redzones for stack and global variables
- Runtime part
 - Maintains shadow memory to track memory state
 - Hooks into kernel allocators to track alloc/free events
 - **Prints bug reports**

Reporting bugs

- When KASAN detects a bug, it prints a bug report
- Reports should be detailed to be useful
- Among other things, KASAN keeps and prints alloc and free stack traces
 - For slab (and page_alloc via page owner)
 - Stack traces saved in [stack depot](#)
 - Alloc stack trace handle is stored [in redzone](#)
 - Free stack trace handle is stored either [in object](#) or [in redzone](#)

Report example: slab out-of-bounds [1/5]

```
void kmalloc_oob_right(void)
{
    char *ptr;
    size_t size = 115;

    ptr = kmalloc(size, GFP_KERNEL);
    ptr[size] = 'x';
}
```

Report example: slab out-of-bounds [2/5]

BUG: KASAN: **slab-out-of-bounds** in `kmalloc_oob_right+0x408/0x454`

Write of size 1 at addr **ffff00000d039a73** by task `kunit_try_catch/96`

CPU: 0 PID: 96 Comm: `kunit_try_catch` Tainted: G

N 6.0.0-rc3

#32

Hardware name: `linux,dummy-virt` (DT)

Call trace:

...

`__asan_report_store1_noabort+0x2c/0x38`

`kmalloc_oob_right+0x408/0x454`

Report example: slab out-of-bounds [3/5]

Allocated by task 96:

kasan_set_track+0x3c/0x6c

kasan_save_alloc_info+0x24/0x30

__kasan_kmalloc+0x90/0xa8

kmem_cache_alloc_trace+0x1e8/0x340

kmalloc_oob_right+0x100/0x454

kunit_try_run_case+0xec/0x2a4

kunit_generic_run_threadfn_adapter+0x54/0x80

kthread+0x220/0x2f4

ret_from_fork+0x10/0x20

Report example: slab out-of-bounds [4/5]

The buggy address belongs to the object at ffff00000d039a00

which belongs to the cache **kmalloc-128** of size 128

The buggy address is located **115 bytes inside** of

128-byte region [ffff00000d039a00, ffff00000d039a80)

The buggy address **belongs** to the **physical page**:

page:(____ptrval____) refcount:1 mapcount:0 mapping:0000000000000000 index:0x0

pfn:0x4d039

flags: 0xffff00000000200(slab|node=0|zone=0|lastcpupid=0xffff)

raw: 0ffff00000000200 0000000000000000 dead00000000122 ffff000006c01300

raw: 0000000000000000 0000000080100010 00000001ffffffff 0000000000000000

page dumped because: kasan: bad access detected

Report example: slab out-of-bounds [5/5]

Memory state around the buggy address:

```
ffff00000d039900: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
ffff00000d039980: fc fc fc fc fc fc fc fc fc fc fc fc fc fc fc
>ffff00000d039a00: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 03 fc
                                                         ^
ffff00000d039a80: fc fc fc fc fc fc fc fc fc fc fc fc fc fc fc
ffff00000d039b00: fc fc fc fc fc fc fc fc fc fc fc fc fc fc fc
```

Other assorted parts

- Compiler module
 - Instruments memory accesses when building the kernel
 - Inserts redzones for stack and global variables
- Runtime part
 - Maintains shadow memory to track memory state
 - Hooks into kernel allocators to track alloc/free events
 - Prints bug reports

Generic KASAN notes and parts

- Some kernel parts are not instrumented by compiler
 - Assembly; early boot code, allocators (grep for KASAN_SANITIZE)
- CONFIG_KASAN_OUTLINE: outline instrumentation mode
 - Compiler adds function calls; slower, but image is smaller

KASAN notes and parts

- KASAN has 3 modes
 - Generic: just covered, supported by many architectures
 - [Software Tag-Based](#): arm64-only, software Memory Tagging
 - [Hardware Tag-Based](#): arm64-only, Arm Memory Tagging Extension
- Test suite
 - [KUnit-compatible](#) tests
 - [KUnit-incompatible](#) tests
- Bugs and features are tracked in [Bugzilla](#)
 - Some modes are missing certain features

Generic KASAN summary

- Dynamic memory corruption detector for the Linux kernel
- Finds out-of-bounds, use-after-free, and double/invalid-free bugs
- Supports slab, page_alloc, vmalloc, stack, and global memory
- Requires compiler support: implemented in both Clang and GCC
- google.github.io/kernel-sanitizers/KASAN
- Relatively fast: ~x2 slowdown
- RAM impact: shadow (1/8 RAM) + quarantine (1/32 RAM) + ~x1.5 for slab
- Basic usage: enable and run tests or fuzzer

More Sanitizers

More Sanitizers

- KASAN takes care of out-of-bounds and use-after-free bugs
- What about the rest?

KMSAN – Kernel Memory Sanitizer

- Detects uses of uninitialized memory
 - Including information leaks to userspace or to hardware
- Uses compiler instrumentation and shadow memory
 - Shadow memory describes which memory is initialized
- On the way [into mainline](#)
- google.github.io/kernel-sanitizers/KMSAN

KCSAN — Kernel Concurrency Sanitizer

- Detects data races
- Uses compiler instrumentation and "soft" watchpoints
 - Stalls on access watchpoints and checks if memory value changed
- google.github.io/kernel-sanitizers/KCSAN
- Also see [KTSAN](#)
 - Failed attempt to implement a Happens-Before data race detector

UBSAN – [Kernel] Undefined Behavior Sanitizer

- Detects certain types of undefined behaviour
- Uses compiler instrumentation
- No K in its CONFIG_ name
- kernel.org/doc/html/latest/dev-tools/ubsan.html
- [After-talk update:] Some UBSAN checks can be enabled in production

In-field Sanitizers

In-field Sanitizers

- So far, mentioned Sanitizers were intended for testing and fuzzing
- What about running Sanitizers with real workloads?
 - In dogfood (internal beta testing with real-world usage)
 - In production

KFENCE — Kernel Electric-Fence

- Detects out-of-bounds, use-after-free, and double/invalid-free for slab
- Places allocated object next to a protected guard page
- Uses sampling to make the overhead unnoticeable
 - ⇒ Probabilistics, need to be deployed across fleet of machines
- Can be used in production or in dogfood
- Not a "Sanitizer" by name but is one by spirit
- google.github.io/kernel-sanitizers/KFENCE

Software Tag-Based KASAN


- [KASAN mode](#) based on software Memory Tagging approach
 - Tag checks are inserted via compiler instrumentation
- Similar performance impact to Generic: ~x2
- Less RAM impact than Generic: less shadow, no quarantine
 - Recommended for arm64 devices with limited RAM, e.g. Android
- Bigger performance/RAM impact than KFENCE
 - OK for dogfood
 - Prepares for Arm Memory Tagging Extension
- [Memory Tagging for the kernel: Tag-Based KASAN](#) [video]

Sanitizers as a mitigation

Hardware Tag-Based KASAN

- [KASAN mode](#) based on Arm Memory Tagging Extension (MTE)
 - Tag checks performed by CPU
- Either a production mitigation or a fast in-field bug detector
- RAM impact: ~3% for storing memory tags
- Performance impact: still unknown, no CPUs released yet
 - <10% expected in Sync mode
 - ~0% expected in Async mode

LSS 2021 talk on Hardware Tag-Based KASAN



Memory Tagging + Linux kernel =



or Mitigating Linux kernel memory corruptions with Arm Memory Tagging

Andrey Konovalov, xairy.io

Linux Security Summit
October 1st 2021

xairy.io/talks/:

- [Slides](#)
- [Video](#)

Production stack traces for KASAN

- Having proper HW_TAGS KASAN reports from production requires:
 - Fast stack trace collection
 - Memory-bounded stack trace storage
- Plan:
 - Use Shadow Call Stack for collecting traces (arm64 maintainers [resist](#))
 - Store stack trace handles in stack ring instead of redzones ([in mm](#))
 - Memory-bounded mode for stack depot
 - Potentially, use sampling to limit the performance impact

Extending Sanitizers

Extending Sanitizers

- Basic usage: just enable
- Advanced usage: extend
- Sanitizers provide frameworks
 - KASAN — for marking memory accessible/inaccessible
 - KMSAN — for checking whether memory is initialized

Extending KASAN: custom redzones

- `sk_buff.head` data buffer allocated as a single chunk
 - Has `skb_shared_info` [placed at the end](#)
- Overflowing data buffer into `skb_shared_info` is not caught by KASAN
 - So called intra-object-overflow
 - I targeted this `skb` behavior in my [CVE-2017-1000112 exploit](#)
- Possible KASAN extension:
 - Add a redzone between data buffer and `skb_shared_info`
 - Related issue tracked in [KASAN: poison skb linear data tail](#)

Extending KASAN: add support for other allocators

- Example: percpu allocator is [not supported](#) by KASAN

Percpu allocator: fairly used

block/bdev.c, line 500	net/caif/caif_dev.c, line 99	net/netfilter/nf_conntrack_core.c, line 2791
block/bio.c, line 1716	net/caif/cffrml.c, line 40	
block/blk-iocost.c, line 2845	net/can/raw.c, line 352	net/netfilter/nf_synproxy_core.c, line 348
block/blk-mq.c, line 3840	net/core/dev.c, line 10594	net/netfilter/nft_set_pipapo.c, 4 times
crypto/cryptd.c, line 105	net/core/gro_cells.c, line 73	net/openvswitch/actions.c, 2 times
fs/aio.c, line 782	net/core/ neighbour.c, line 1755	net/packet/af_packet.c, line 1230
fs/ext4/mballocc.c, line 3484	net/core/page_pool.c, line 184	net/sched/act_api.c, line 733
fs/gfs2/ops_fstype.c, line 82	net/core/sock.c, line 3668	net/sched/cls_basic.c, line 213
fs/namespace.c, line 214	net/dccp/proto.c, line 1081	net/sched/cls_matchall.c, line 221
fs/nfs/iostat.h, line 68	net/ipv4/af_inet.c, 6 times	net/sched/cls_u32.c, line 1054
fs/squashfs/decompressor_multi_percpu.c, line 35	net/ipv4/fib_trie.c, line 2427	net/sched/sch_generic.c, line 957
fs/xfs/xfs_super.c, 3 times	net/ipv6/af_inet6.c, 4 times	net/sctp/protocol.c, line 1227
include/scsi/libfc.h, line 836	net/ipv6/seg6_hmac.c, 2 times	net/smc/smc_stats.c, line 26
kernel/bpf/bpf_lru_list.c, 2 times	net/mpls/af_mpls.c, line 1459	net/tipc/crypto.c, line 535
kernel/bpf/devmap.c, line 1075	net/mptcp/mib.c, line 73	net/tls/tls_main.c, line 1009
kernel/bpf/ percpu_freelist.c, line 10	net/netfilter/ipvs/ip_vs_ctl.c, 3 times	net/xfrm/xfrm_ipcomp.c, 2 times
		net/xfrm/xfrm_policy.c, line 3994

Percpu allocator: no sanitization (out-of-bounds)

```
void percpu_oob(struct kunit *test) {  
    char __percpu *ptr = __alloc_percpu(128, 8);  
    char *c_ptr = this_cpu_ptr(ptr);  
    KUNIT_EXPECT_KASAN_FAIL(test, c_ptr[128] = 0x42);  
    free_percpu(ptr);  
}
```

percpu_oob: EXPECTATION FAILED at mm/kasan/kasan_test.c:1386

KASAN **failure expected** in "c_ptr[128] = 0x42", but **none occurred**

not ok 57 - percpu_oob

Percpu allocator: no sanitization (use-after-free)

```
void percpu_uaf(struct kunit *test) {  
    char __percpu *ptr = __alloc_percpu(128, 8);  
    char *c_ptr = this_cpu_ptr(ptr);  
    free_percpu(ptr);  
    KUNIT_EXPECT_KASAN_FAIL(test, c_ptr[0] = 0x42);  
}
```

percpu_uaf: EXPECTATION FAILED at mm/kasan/kasan_test.c:1398

KASAN **failure expected** in "c_ptr[0] = 0x42", but **none occurred**

not ok 58 - percpu_uaf

Percpu allocator: need KASAN annotations

- Tracked in [KASAN: sanitize per-cpu allocations](#)
- Annotations will allow detecting more bugs
- Similar extensions also applicable to other allocators
 - Including custom allocators in drivers
 - Android allocators?

Extending KMSAN: USB annotations

- KMSAN was extended to detect information leaks over [USB](#)
- As a result, [syzbot](#) found bugs:
 - [KMSAN: kernel-usb-infoleak in hif_usb_send](#)
 - [KMSAN: kernel-usb-infoleak in usbnet_write_cmd](#)
 - [KMSAN: kernel-usb-infoleak in hid_submit_ctrl](#)
 - [KMSAN: kernel-usb-infoleak in pcan_usb_wait_rsp](#)
 - [KMSAN: kernel-usb-infoleak in ttusb_dec_send_command](#)
 - [KMSAN: kernel-usb-infoleak in pcan_usb_pro_send_req](#)
 - [KMSAN: kernel-usb-infoleak in pcan_usb_pro_init](#)

Outro

Finding more bugs

1. Improving the fuzzer
2. Improving the bug detectors

Know Sanitizers

- Testing and fuzzing
 - KASAN — out-of-bounds, use-after-free, invalid/double-free
 - KMSAN — uses of uninitialized memory, information leaks
 - KCSAN — data races
 - UBSAN — undefined behavior
- Dogfood or production
 - KFENCE — slab memory corruptions, sampling-based
 - SW_TAGS KASAN — memory corruptions, arm64-only
- As a mitigation
 - HW_TAGS KASAN — memory corruptions, arm64-only, requires Arm MTE

Make better bug detectors

- Extend Sanitizers
 - Add KASAN redzones to detect intra-object-overflows
 - Add KASAN annotations to more allocators
 - Add KMSAN annotations to detect leaks across security boundaries
- Make your own bug detectors
 - For other types of bugs (type confusions?)
 - For logical bugs (missing TLB flushes?)
 - Take inspiration from Sanitizers (compiler instrumentation)

Thank you!

google.github.io/kernel-sanitizers

Andrey Kononov, xairy.io