



Kernel-Level Debuggers : KGDB

Important Notice : Courseware - Legal

This courseware is both the product of the author and of freely available opensource and/or public domain materials. Wherever external material has been shown, it's source and ownership have been clearly attributed. We acknowledge all copyrights and trademarks of the respective owners.

The contents of the **courseware PDFs are considered proprietary** and thus cannot be copied or reproduced in any form whatsoever without the explicit written consent of the author.

Only the programs - **source code** and binaries (where applicable) - that form part of this courseware, and that are made available to the participant, are released under the terms of the [permissive MIT license](#).

Under the terms of the MIT License, you can certainly use the source code provided here; you must just attribute the original source (author of this courseware and/or other copyright/trademark holders).

VERY IMPORTANT :: Before using this source(s) in your project(s), you ***MUST*** check with your organization's legal staff that it is appropriate to do so.

The courseware PDFs are ***not*** under the MIT License, they are to be kept confidential, non-distributable without consent, for your private internal use only.

The duration, contents, content matter, programs, etc. contained in this courseware and companion participant VM are subject to change at any point in time without prior notice to individual participants.

Care has been taken in the preparation of this material, but there is no warranty, expressed or implied of any kind, and we can assume no responsibility for any errors or omissions. No liability is assumed for incidental or consequential damages in connection with or arising out of the use of the information or programs contained herein.

2000-2023 Kaiwan N Billimoria
kaiwanTECH, Bangalore, India.

kaiwanTECH Linux OS Corporate Training Programs

Please do check out our current offering of world-class, seriously-valuable, high on returns, technical Linux OS corporate training programs here: <http://bit.ly/ktcorp>

<< Source: “Linux Device Drivers” by J Corbet, A Rubini & GK Hartman, 3rd Ed., O'Reilly >>

Using gdb (for kernel-space)

`gdb` can be quite useful for looking at the system internals. Proficient use of the debugger at this level requires some confidence with `gdb` commands, some understanding of assembly code for the target platform, and the ability to match source code and optimized assembly.

The debugger must be invoked as though the kernel were an application. In addition to specifying the filename for the ELF kernel image, you need to provide the name of a core file on the command line. **For a running kernel, that core file is the kernel core image, `/proc/kcore`.** A typical invocation of `gdb` looks like the following:

```
gdb /usr/src/linux/vmlinux /proc/kcore
```

[Note that you should be running the kernel that the `vmlinux` image corresponds to].

The first argument is the name of the uncompressed ELF kernel executable, not the `zImage` or `bzImage` or anything built specifically for the boot environment.

The second argument on the `gdb` command line is the name of the core file. Like any file in `/proc`, `/proc/kcore` is generated when it is read. When the read system call executes in the `/proc` filesystem, it maps to a data-generation function rather than a data retrieval one; we've already exploited this feature in the section “Using the `/proc` Filesystem” earlier in this chapter. `kcore` is used to represent the kernel “executable” in the format of a core file; it is a huge file, because it represents the whole kernel address space, which corresponds to all physical memory. From within `gdb`, you can look at kernel variables by issuing the standard `gdb` commands. For example, `p jiffies` prints the number of clock ticks from system boot to the current time.

When you print data from `gdb`, the kernel is still running, and the various data items have different values at different times; `gdb`, however, optimizes access to the core file by caching data that has already been read. If you try to look at the `jiffies` variable once again, you'll get the same answer as before. Caching values to avoid extra disk access is a correct behavior for conventional core files but is inconvenient when a “dynamic” core image is used. The solution is to issue the command **`core-file /proc/kcore`** whenever you want to flush the `gdb` cache; the debugger gets ready to use a new core file and discards any old information. You won't, however, always need to issue `core-file` when reading a new datum; `gdb` reads the core in chunks of a few kilobytes and caches only chunks it has already referenced.

Numerous capabilities normally provided by `gdb` are not available when you are working with the kernel. For example, `gdb` is not able to modify kernel data; it expects to be running a program to be debugged under its own control before playing with its memory image. It is also not possible to set breakpoints or watchpoints, or to single-step through kernel functions.

Note that, in order to have symbol information available for `gdb`, you must compile your kernel with the `CONFIG_DEBUG_INFO` option set. The result is a far larger kernel image on disk, but, without that information, digging through kernel variables is almost impossible.

<<

Sample Session:

```

# gdb <...>/linux-2.6.17/vmlinux /proc/kcore
GNU gdb Red Hat Linux (6.3.0.0-1.96rh)
[...]
Core was generated by `ro root=LABEL=/1 rhgb quiet 1'.
#0  0x00000000 in ?? ()
(gdb) p jiffies
$1 = 3700920
(gdb)

## <wait a bit...>

(gdb) p jiffies
$2 = 3700920                                     ## doesn't change
(gdb) core-file /proc/kcore                     ## refresh
Core was generated by `ro root=LABEL=/1 rhgb quiet 1'.
#0  0x00000000 in ?? ()
(gdb) p jiffies
$3 = 3717432
(gdb)
(gdb) p tasklist_lock
$4 = {raw_lock = {<No data fields>}, magic = 3736018669, owner_cpu =
4294967295, owner = 0xffffffff}
(gdb) set print pretty
(gdb) p tasklist_lock
$5 = {
  raw_lock = {<No data fields>},
  magic = 3736018669,
  owner_cpu = 4294967295,
  owner = 0xffffffff
}
(gdb) p task_state_array
$8 = {0xc02c6b9c "R (running)", 0xc02c6ba8 "S (sleeping)", 0xc02c6bb5 "D
(disk sleep)",
      0xc02c6bc4 "T (stopped)", 0xc02c6bd0 "T (tracing stop)", 0xc02c6be1 "Z
(zombie)", 0xc02c6bec "X (dead)"}
(gdb)
>>

```

Hardware Watchpoints

A hardware watchpoint has GDB use CPU hardware registers to “watch” a variable or expression. It’s much faster than a software watchpoint but obviously requires processor support.

How do we know?

(gdb) show can-use-hw-watchpoints

Debugger's willingness to use watchpoint hardware is 1.

(gdb)

There are three types of (hardware) watchpoint GDB commands:

- **watch** <var/expression> : sets up a hardware watchpoint that triggers on the variable/expression changing (being **written** to)
- **rwatch** <var/expression> : sets up a hardware watchpoint that triggers on the variable/expression being **read**
- **awatch** <var/expression> : sets up a hardware watchpoint that triggers on **either a read or write** on the variable/expression

On triggering, GDB will show the hardware watchpoint number, name and the old and new value of the variable/expression being watched.

An example snippet (when debugging the Linux kernel with KGDB):

(gdb) info watchpoints

No watchpoints.

(gdb) watch jiffies<tab><tab>

jiffies	jiffies_64_to_clock_t	jiffies_till_next_fqs
jiffies_to_timespec		
jiffies.c	jiffies_lock	jiffies_till_sched_qs
jiffies_to_timeval		
jiffies.h	jiffies_read	jiffies_to_clock_t
jiffies_to_usecs		
jiffies_64	jiffies_till_first_fqs	jiffies_to_msecs

(gdb) watch jiffies

Hardware watchpoint 5: jiffies

(gdb) c

Continuing.

[...]

Hardware watchpoint 5: jiffies << we hit the watchpoint => 'jiffies' has been written to >>

Old value = 4294938708

New value = 4294938709

do_timer (ticks=1) at kernel/time/timekeeping.c:1674

1674 calc_global_load(ticks);

(gdb) bt

#0 do_timer (ticks=1) at kernel/time/timekeeping.c:1674

#1 0x8007fe7c in tick_periodic (cpu=<optimized out>) at kernel/time/tick-common.c:86

```
#2 0x8008004c in tick_handle_periodic (dev=0x8fdfl00) at
kernel/time/tick-common.c:103
#3 0x80015da4 in twd_handler (irq=<optimized out>, dev_id=<optimized out>)
at arch/arm/kernel/smp_twd.c:236
#4 0x80067904 in handle_percpu_devid_irq (irq=29, desc=0x8f805600) at
kernel/irq/chip.c:711
#5 0x800635dc in generic_handle_irq_desc (desc=<optimized out>,
irq=<optimized out>)
    at include/linux/irqdesc.h:128
#6 generic_handle_irq (irq=29) at kernel/irq/irqdesc.c:351
#7 0x80063920 in __handle_domain_irq (domain=0x8f803000, hwirq=<optimized
out>, lookup=true,
    regs=<optimized out>) at kernel/irq/irqdesc.c:388
#8 0x800086e4 in handle_domain_irq (regs=<optimized out>, hwirq=<optimized
out>, domain=<optimized out>)
    at include/linux/irqdesc.h:146
#9 gic_handle_irq (regs=0x1 <__vectors_start>) at drivers/irqchip/irq-
gic.c:273
#10 0x80013844 in __irq_svc () at arch/arm/kernel/entry-armv.S:205
Backtrace stopped: frame did not save the PC
```

(gdb) c

Continuing.

Hardware watchpoint 5: jiffies << we hit the same watchpoint again >>

Old value = 4294938709

New value = 4294938710

do_timer (ticks=1) at kernel/time/timekeeping.c:1674

1674 calc_global_load(ticks);

(gdb) info watchpoints

Num	Type	Disp	Enb	Address	What
5	hw watchpoint	keep y			jiffies

breakpoint already hit 2 times

(gdb) delete 5

(gdb) rwatch jiffies << setup a 'read' watchpoint on 'jiffies' >>

Hardware read watchpoint 6: jiffies

(gdb) c

Continuing.

Hardware read watchpoint 6: jiffies << we hit the read watchpoint =>
'jiffies' has been read from >>

Value = 4294938710

0x8004f70c in calc_global_load (ticks=1) at kernel/sched/proc.c:347

347 if (time_before(jiffies, calc_load_update + 10))

(gdb) bt

```
#0 0x8004f70c in calc_global_load (ticks=1) at kernel/sched/proc.c:347
```

```
#1 0x8007aa2c in do_timer (ticks=<optimized out>) at
```

```
kernel/time/timekeeping.c:1674
```

```
#2 0x8007fe7c in tick_periodic (cpu=<optimized out>) at kernel/time/tick-
common.c:86
```

```
#3 0x8008004c in tick_handle_periodic (dev=0x8fdfl00) at
```

```
kernel/time/tick-common.c:103
```

```
#4 0x80015da4 in twd_handler (irq=<optimized out>, dev_id=<optimized out>)
at arch/arm/kernel/smp_twd.c:236
```

```

#5 0x80067904 in handle_percpu_devid_irq (irq=29, desc=0x8f805600) at
kernel/irq/chip.c:711
#6 0x800635dc in generic_handle_irq_desc (desc=<optimized out>,
irq=<optimized out>)
    at include/linux/irqdesc.h:128
#7 generic_handle_irq (irq=29) at kernel/irq/irqdesc.c:351
#8 0x80063920 in __handle_domain_irq (domain=0x8f803000, hwirq=<optimized
out>, lookup=true,
    regs=<optimized out>) at kernel/irq/irqdesc.c:388
#9 0x800086e4 in handle_domain_irq (regs=<optimized out>, hwirq=<optimized
out>, domain=<optimized out>)
    at include/linux/irqdesc.h:146
#10 gic_handle_irq (regs=0x1 <__vectors_start>) at drivers/irqchip/irq-
gic.c:273
#11 0x80013844 in __irq_svc () at arch/arm/kernel/entry-armv.S:205
Backtrace stopped: frame did not save the PC
(gdb) info watchpoints
Num      Type          Disp Enb Address      What
6        read watchpoint keep y          jiffies
        breakpoint already hit 1 time
(gdb) delete 6
(gdb) c
Continuing.

```

Ref [SO]: [Can I set a breakpoint on 'memory access' in GDB?](#)

Kernel Modules and gdb

<<

Note- recent kernel doc:

[Debugging kernel and modules via gdb](#)

>>

With the debugging information available, you can learn a lot about what is going on inside the kernel. gdb happily prints out structures, follows pointers, etc. One thing that is harder, however, is examining modules. Since modules are not part of the vmlinux image passed to gdb, the debugger knows nothing about them. Fortunately, as of kernel 2.6.7, it is possible to teach gdb what it needs to know to examine loadable modules.

Linux loadable modules are ELF-format executable images; as such, they have been divided up into numerous sections. A typical module can contain a dozen or more sections, but there are typically three that are relevant in a debugging session:

.text

This section contains the executable code for the module. The debugger must know where this section is to be able to give tracebacks or set breakpoints. (Neither of these operations is relevant when running the debugger on /proc/kcore, but they can be useful when working with kgdb).

.bss

.data

These two sections hold the module's variables. Any variable that is not initialized at compile time ends up in .bss, while those that are initialized go into .data.

Making gdb work with loadable modules **requires informing the debugger about where a given module's sections have been loaded**. That information is available in sysfs, under `/sys/module/<module-name>/sections`. For example, after loading the scull module, the directory `/sys/module/scull/sections` contains files with names such as `.text`; the content of each file is the base address for that section.

<<

Note!

You require **root** access to retrieve data from `/sys/module/<module-name>/sections/<file>`

>>

We are now in a position to issue a gdb command telling it about our module. The command we need is **add-symbol-file**; this command takes as parameters the name of the module object file, the `.text` base address, and a series of optional parameters describing where any other sections of interest have been put. After digging through the module section data in sysfs, we can construct a command such as:

(gdb) help add-symbol-file


```
Load symbols from FILE, assuming FILE has been dynamically loaded.
Usage: add-symbol-file FILE ADDR [-s <SECT> <SECT_ADDR> -s <SECT>
<SECT_ADDR> ...]
ADDR is the starting address of the file's text.
The optional arguments are section-name section-address pairs and
should be specified if the data and bss segments are not contiguous
with the text. SECT is a section name to be loaded at SECT_ADDR.
(gdb) add-symbol-file <...>/scull.ko 0xd0832000 \
-s .bss 0xd0837100 \
-s .data 0xd0836be0
```

We have included a [small script in the sample source \(gdbline.sh\)](#) that can create this command for a given module.

We can now use gdb to examine variables in our loadable module. Here is a quick example taken from a scull debugging session:

```
(gdb) add-symbol-file scull.ko 0xd0832000 \
-s .bss 0xd0837100 \
-s .data 0xd0836be0
add symbol table from file "scull.ko" at
    .text_addr = 0xd0832000
    .bss_addr = 0xd0837100
    .data_addr = 0xd0836be0
(y or n) y
Reading symbols from scull.ko...done.
(gdb) p scull_devices[0]
$1 = {data = 0xcfd66c50,
quantum = 4000,
qset = 1000,
size = 20881,
access_key = 0,
...}
```

Here we see that the first scull device currently holds 20,881 bytes. If we wanted, we could follow the data chain, or look at anything else of interest in the module.

<< See (and try out) the *hello_gdb* kernel module >>

Sample Session:

Ensure that the Makefile includes the '-g' option:

```
EXTRA_CFLAGS += -DDEBUG -g
ccflags-y += -DDEBUG -g
```

make

```
make -C /lib/modules/2.6.17/build SUBDIRS=/mnt/.../hello_gdb modules
make[1]: Entering directory `/mnt/data_5GB/kbuild/linux-2.6.17'
CC [M] /mnt/.../hello_gdb/hello_gdb.o
/mnt/.../hello_gdb/hello_gdb.c: In function `procread':
```

```
/mnt/.../hello_gdb/hello_gdb.c:66: warning: unused variable `p'
Building modules, stage 2.
MODPOST
CC      /mnt/.../hello_gdb/hello_gdb.mod.o
LD [M]  /mnt/.../hello_gdb/hello_gdb.ko
make[1]: Leaving directory `/mnt/data_5GB/kbuild/linux-2.6.17'
#
# insmod hello_gdb.ko
```

Here's a [helper shell script](#) that can construct the appropriate gdb “**add-symbol-file**” command that you'll use within gdb to help debug kernel modules:

```
# cat gdbline.sh
[ ... ]
cd /sys/module/$1/sections
echo "Copy-paste the following lines into GDB"
echo "---snip---"

# Don't issue newlines and the continuation \ ; results in GDB err:
# Unrecognized argument "
# "
[[ -f .text ]] && {
    sudo echo -n "add-symbol-file $2 `/bin/cat .text` "
    #sudo echo " \\"
} || [[ -f .init.text ]] && {
    sudo echo -n "-s .init.text `/bin/cat .init.text` "
}

for section in `[a-z]*`; do
    if [[ ${section} != ".text" || ${section} != ".init.text" ]]; then
        sudo echo -n " -s" ${section} `/bin/cat ${section}`
    fi
done
echo "
---snip---"
$ ./gdbline.sh
Usage: ./gdbline.sh module-name image-filename
       module-name: name of the (already inserted) kernel module (without
the .ko)
       image-filename: pathname to the kernel module.

$ sudo ./gdbline.sh hello_gdb ./hello_gdb.ko
Copy-paste the following lines into GDB
---snip---
add-symbol-file ./hello_gdb.ko 0xfffffffffc386d000 -s .init.text
0xfffffffffc1027000 -s .bss 0xfffffffffc386f380 -s .gnu.linkonce.this_module
0xfffffffffc386f000 -s .init.text 0xfffffffffc1027000 -s .note.gnu.build-id
0xfffffffffc386e000 -s .note.Linux 0xfffffffffc386e024 -s .return_sites
0xfffffffffc386e0b6 -s .rodata.str1.1 0xfffffffffc386e054 -s .strtab
0xfffffffffc10284b0 -s .symtab 0xfffffffffc1028000 -s .text
0xfffffffffc386d000 -s __mcount_loc 0xfffffffffc386e0a6
```

---snip---

Copy the above output from the shell script (the lines between --snip--).

```
hello_gdb # ./gdbline.sh hello_gdb ./hello_gdb.ko
Copy-paste the following lines into GDB
---snip---
add-symbol-file ./hello_gdb.ko 0xffffffffc09f5000 -s .init.text 0xffffffffc0a01000 -s .bss 0xffffffffc09f7400 -s .gnu.
linkonce.this_module 0xffffffffc09f7000 -s .init.text 0xffffffffc0a01000 -s .note.gnu.build-id 0xffffffffc09f6000 -s .n
ote.Linux 0xffffffffc09f6024 -s .return_sites 0xffffffffc09f609e -s .rodata.str1.1 0xffffffffc09f603c -s .strtab 0xffff
ffffc0a02498 -s .symtab 0xffffffffc0a02000 -s .text 0xffffffffc09f5000 -s __mcount_loc 0xffffffffc09f608e
---snip---

hello_gdb # gdb -q -c /proc/kcore /lib/modules/5.10.153-mykernel/build/vmlinux
Reading symbols from /lib/modules/5.10.153-mykernel/build/vmlinux...
[New process 1]
Core was generated by `BOOT_IMAGE=vmlinux-5.10.153-mykernel root=UUID=b67edd03-5aa9-4826-add9-3de240b'.
#0  0x0000000000000000 in fixed_percpu_data ()
(gdb) add-symbol-file ./hello_gdb.ko 0xffffffffc09f5000 -s .init.text 0xffffffffc0a01000 -s .bss 0xffffffffc09f7400 -s
.gnu.linkonce.this_module 0xffffffffc09f7000 -s .init.text 0xffffffffc0a01000 -s .note.gnu.build-id 0xffffffffc09f6000
-s .note.Linux 0xffffffffc09f6024 -s .return_sites 0xffffffffc09f609e -s .rodata.str1.1 0xffffffffc09f603c -s .strtab
0xffffffffc0a02498 -s .symtab 0xffffffffc0a02000 -s .text 0xffffffffc09f5000 -s __mcount_loc 0xffffffffc09f608e
add symbol table from file ".:/hello_gdb.ko" at
    .text_addr = 0xffffffffc09f5000
    .init.text_addr = 0xffffffffc0a01000
    .bss_addr = 0xffffffffc09f7400
    .gnu.linkonce.this_module_addr = 0xffffffffc09f7000
    .init.text_addr = 0xffffffffc0a01000
    .note.gnu.build-id_addr = 0xffffffffc09f6000
    .note.Linux_addr = 0xffffffffc09f6024
    .return_sites_addr = 0xffffffffc09f609e
    .rodata.str1.1_addr = 0xffffffffc09f603c
    .strtab_addr = 0xffffffffc0a02498
    .symtab_addr = 0xffffffffc0a02000
    .text_addr = 0xffffffffc09f5000
    __mcount_loc_addr = 0xffffffffc09f608e
(y or n) y
Reading symbols from ./hello_gdb.ko...
warning: section .init.text not found in /home/osboxes/kaiwanTECH/L5_debug_trg/kernel_debug/hello_gdb/hello_gdb.ko
warning: section .strtab not found in /home/osboxes/kaiwanTECH/L5_debug_trg/kernel_debug/hello_gdb/hello_gdb.ko
warning: section .symtab not found in /home/osboxes/kaiwanTECH/L5_debug_trg/kernel_debug/hello_gdb/hello_gdb.ko
warning: section .text not found in /home/osboxes/kaiwanTECH/L5_debug_trg/kernel_debug/hello_gdb/hello_gdb.ko
(gdb) p mine
$1 = (MYS *) 0xffff894fc5e8ce48
```

<< NOTE: it's IMP to pass along the path to the (debug) vmlinux ! >>

<< Note: Below, the values are different from those above >>

...

(gdb) set print pretty

(gdb) p *mine

```
$4 = {
  tx = 0,
  Corresponds to jiffies value  rx = 0,
  j = 6128429
}
```

(gdb) x/16 0xc37f6a60

```
0xc37f6a60:  0x00000000  0x005d832d  0x5a5a5a5a  0x5a5a5a5a
0xc37f6a70:  0x5a5a5a5a  0x5a5a5a5a  0x5a5a5a5a  0xa55a5a5a
0xc37f6a80:  0x170fc2a5  0xd081e04d  0x170fc2a5  0x00000002
0xc37f6a90:  0x00000001  0xdead4ead  0xffffffff  0xffffffff
```

(gdb) p/x *mine

```
$5 = {
  tx = 0x0,
```

```
    rx = 0x0,  
    j = 0x5d832d  
}  
(gdb) q  
#
```

One other useful trick worth knowing about is this:

(gdb) print *(address)

Here, fill in a hex address for address ; the output is a file and line number for the code corresponding to that address. This technique may be useful, for example, to find out where a function pointer really points.

We still cannot perform typical debugging tasks like setting breakpoints or modifying data; to perform those operations, we need to use a tool like kdb or kgdb.

KGDB

The KGDB project allows one to perform *source-level* debugging of the kernel.

From kernel ver 2.6.26, the kgdb facility is built into the mainline kernel; hence, activation is a lot easier (than before when multiple patches had to be applied).

The section below on “Using QEMU and GDB for source-level kernel debugging” is almost identical to how KGDB is used on an actual machine.

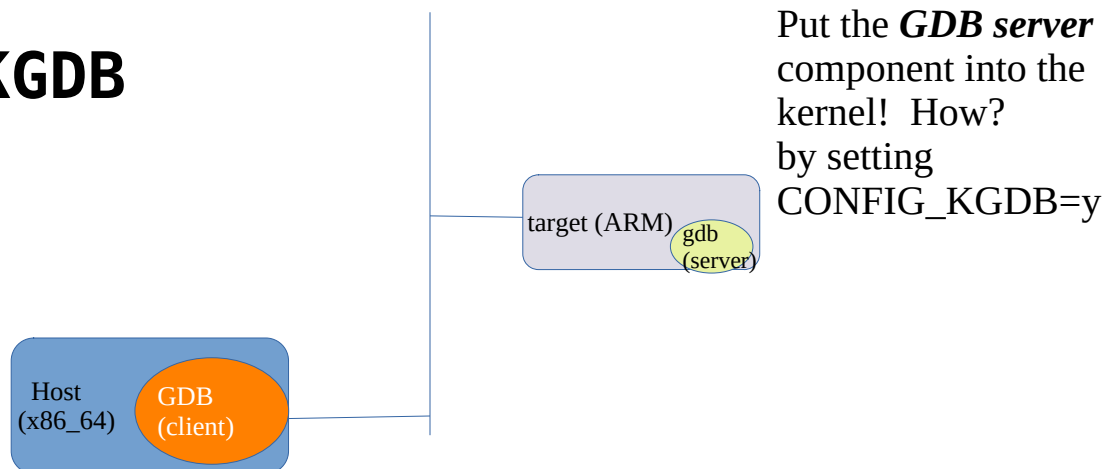
Resources on KGDB:

Kernel docs: [Using kgdb, kdb and the kernel debugger internals](#)

<http://elinux.org/Kgdb>

[A KDB / KGDB SESSION ON THE POPULAR RASPBERRY PI EMBEDDED LINUX BOARD](#)

KGDB



```
(gdb) target remote :1234
```

Using QEMU and GDB for source-level kernel debugging

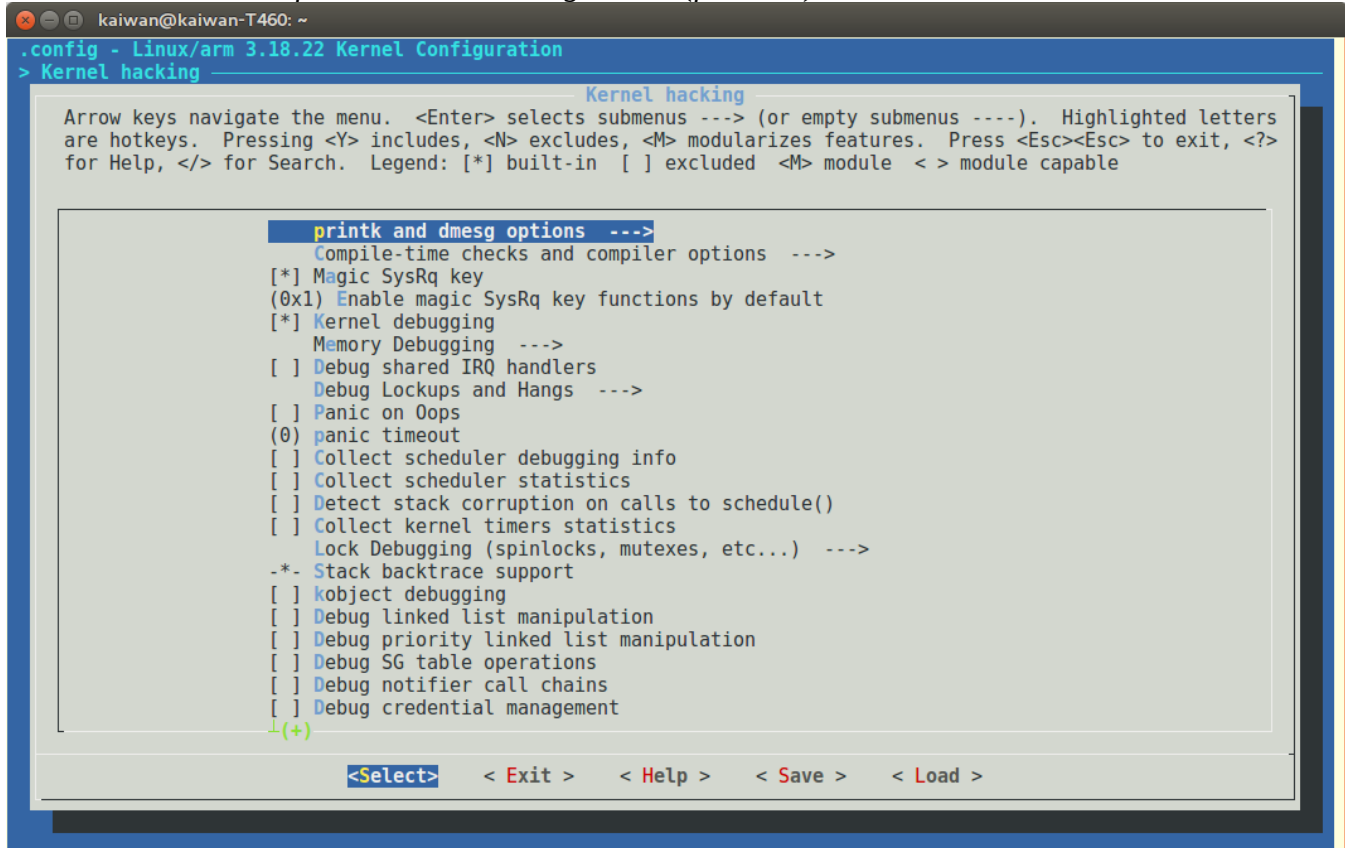
1. Download and install the emulator package, Qemu. <http://wiki.qemu.org/Download>. (We're using the latest one at the time of this writing, i.e., ver 0.14.1)
2. We assume native gdb is present (we're using gdb ver 7.3).

Kernel Only Debugging

3. If the intention is to just perform pure kernel source debugging, then we don't even require a root filesystem.

<<

A screenshot of the “Kernel Hacking” menu (for ARM):



Q. Where are the debug “menus” generated from?

A. The Kconfig files of course.

Specifically:

lib/Kconfig.debug

>>

Build your kernel with debug on and frame pointers on.

CONFIG_DEBUG_KERNEL=y << basic support for kernel debug; turns the
Kernel Hacking menu on >>

CONFIG_DEBUG_INFO=y << build with -g ! includes symbols >>

CONFIG_KGDB=y << KGDB remote debugging support! >>

CONFIG_FRAME_POINTER=y << recommended >>

<< On recent kernels, when configuring with the “menuconfig” UI, goto:

Kernel Hacking / Compile-time checks and compiler options → Compile the kernel with debug info

>>

and, optionally (useful),

CONFIG_DEBUG_SPINLOCK=y
CONFIG_DEBUG_MUTEXES=y
CONFIG_DEBUG_LOCK_ALLOC=y
CONFIG_PROVE_LOCKING=y
CONFIG_DEBUG_BUGVERBOSE=y
CONFIG_EARLY_PRINTK=y

Kernel Command-line Parameters relevant to KGDB

(see Documentation/kernel-parameters.txt)

```
...
ekgdboc=    [X86,KGDB] Allow early kernel console debugging
             ekgdboc=kbd

...

kgdbdbgpc=  [KGDB,HW] kgdb over EHCI usb debug port.
             Format: <Controller#>[,poll interval]
             The controller # is the number of the ehci usb debug
             port as it is probed via PCI. The poll interval is
             optional and is the number seconds in between
             each poll cycle to the debug port in case you need
             the functionality for interrupting the kernel with
             gdb or control-c on the dbgpc connection. When
             not using this parameter you use sysrq-g to break into
             the kernel debugger.

kgdboc=      [KGDB,HW] kgdb over consoles.
             Requires a tty driver that supports console polling,
             or a supported polling keyboard driver (non-usb).
```

```

Serial only format: <serial_device>[,baud]
keyboard only format: kbd
keyboard and serial format: kbd,<serial_device>[,baud]
Optional Kernel mode setting:
kms, kbd format: kms,kbd
kms, kbd and serial format: kms,kbd,<ser_dev>[,baud]

```

```

kgdbwait    [KGDB] Stop kernel execution and enter the
             kernel debugger at the earliest opportunity.

```

...

Using KGDB via a AMD64 Qemu/KVM guest and host

Ref: [Linux Kernel Exploitation 0x0\] Debugging the Kernel with QEMU](#)
 posted by [Keith Makan](#)

A few tips:

- Install all required packages first:

```

sudo apt-get update
sudo apt-get upgrade

```

```

sudo apt-get install git fakeroot build-essential ncurses-dev xz-utils
libssl-dev bc flex libelf-dev bison debootstrap qemu-system-x86

```

- Once the kernel's built, test the Qemu guest VM (leave out the -s -S which makes it wait for gdb to connect); once it works, reboot but this time with the '-S -s' Qemu options:

```
<<
```

```
-S Do not start CPU at startup (you must type 'c' in the monitor).
```

```
-s Shorthand for -gdb tcp::1234, i.e. open a gdbserver on TCP port 1234.
```

```
>>
```

```
$ cat run.sh
```

```
#!/bin/bash
```

```
# ref: http://blog.k3170makan.com/2020/11/linux-kernel-exploitation-0x0-debugging.html
```

```

qemu-system-x86_64 \
  -kernel ../linux-5.10.3/arch/x86/boot/bzImage \
  -append "console=ttyS0 root=/dev/sda earlyprintk=serial nokaslr" \
  -hda ./stretch.img \
  -net user,hostfwd=tcp::10021-:22 -net nic \
  -enable-kvm \
  -nographic \

```



```
-m 2G \
-smp 2 \
-S -s \
-pidfile vm.pid \
2>&1 | tee vm.log
```

It should bootup; you're logged in as root (no passwd).
Shut down and continue...

NOTE

- *You CANNOT run another hypervisor (or Docker container) along with Qemu/KVM with the **-enable-kvm** (accelerator) option !!*
If VirtualBox (or another) is running, and you want to use **-enable-kvm**, first shut it down and then retry...
- Before running GDB on the 'host', cd to the relevant kernel source directory (here it's 5.10.3)... else, you won't see the kernel source (with list, ...)
- Setting hardware breakpoints (with 'hb <func>' is considered a safer bet
- to quit out of Qemu, type Ctrl-a x

Host: run the above script:

```
kdebug_kmake $ ./run.sh
```

Note:

1. First shut down any other hypervisor instance
 2. once run, the guest qemu system will **wait** for GDB to connect from the host:
- On the host, do:

```
$ gdb linux-5.10.3/vmlinux -q
(gdb) target remote :1234
```

```
qemu-system-x86_64 -kernel linux-5.10.3/arch/x86/boot/bzImage -append
console=ttyS0 root=/dev/sda earlyprintk=serial nokaslr -s
```

WARNING: Image format was not specified for 'image/stretch.img' and probing guessed raw.

Automatically detecting the format is dangerous for raw images, write operations on block 0 will be restricted.

Specify the 'raw' format explicitly to remove the restrictions.

```
qemu-system-x86_64: warning: host doesn't support requested feature:
```

```
CPUID.80000001H:ECX.svm [bit 2]
```

```
qemu-system-x86_64: warning: host doesn't support requested feature:
```

```
CPUID.80000001H:ECX.svm [bit 2]
```

```
<< ... kernel GDB server is waiting for GDB client to connect ... >>
```

From another terminal window, connect via GDB client:

```
$ cd <...>/linux-5.10.3
$ gdb -q vmlinux
Reading symbols from vmlinux...
(gdb) target remote :1234
Remote debugging using :1234
0x0000000000000fff0 in gdt_page ()
(gdb)
```

It's worked!

*Try it: set **breakpoints** (on `kmem_cache_alloc()`, `schedule_timeout()`, etc) and see...*

```
(gdb) b kmem_cache_alloc
Breakpoint 3 at 0xffffffff8184a740: file mm/slub.c, line 2903.
(gdb) c
Continuing.
[Switching to Thread 1.2]
Cannot remove breakpoints because program is no longer writable.
Further execution is probably impossible.

Thread 2 hit Breakpoint 3, kmem_cache_alloc (s=0xffff888006217280,
gfpflags=gfpflags@entry=3264) at mm/slub.c:2903
2903 {
(gdb) bt
#0 kmem_cache_alloc (s=0xffff888006217280, gfpflags=gfpflags@entry=3264) at
mm/slub.c:2903
#1 0xffffffff8191b3cd in getname_flags (empty=0x0 <fixed_percpu_data>, flags=1,
filename=0x4ccca3 ".") at fs/namei.c:138
#2 getname_flags (filename=0x4ccca3 ".", flags=1, empty=0x0 <fixed_percpu_data>)
at fs/namei.c:128
#3 0xffffffff8191ec94 in user_path_at_empty (dfd=dfd@entry=-100,
name=name@entry=0x4ccca3 ".", flags=flags@entry=1,
path=path@entry=0xffff888009e07d60, empty=empty@entry=0x0 <fixed_percpu_data>)
at fs/namei.c:2647
#4 0xffffffff818f1ef3 in user_path_at (path=0xffff888009e07d60, flags=1,
name=0x4ccca3 ".", dfd=-100) at ./include/linux/namei.h:59
#5 vfs_statx (dfd=dfd@entry=-100, filename=filename@entry=0x4ccca3 ".",
flags=flags@entry=2048, stat=stat@entry=0xffff888009e07e08,
request_mask=request_mask@entry=2047) at fs/stat.c:185
#6 0xffffffff818f330d in vfs_fstatat (flags=0, stat=0xffff888009e07e08,
filename=0x4ccca3 ".", dfd=-100) at fs/stat.c:207
#7 vfs_stat (stat=0xffff888009e07e08, filename=0x4ccca3 ".") at
./include/linux/fs.h:3121
#8 __do_sys_newstat (filename=0x4ccca3 ".", statbuf=0x7ffc6ca64730) at
fs/stat.c:349
#9 0xffffffff818f33d9 in __se_sys_newstat (statbuf=<optimized out>,
filename=<optimized out>) at fs/stat.c:345
#10 __x64_sys_newstat (regs=0xffff888009e07f58) at fs/stat.c:345
#11 0xffffffff833aefb8 in do_syscall_64 (nr=<optimized out>,
regs=0xffff888009e07f58) at arch/x86/entry/common.c:46
#12 0xffffffff8340008c in entry_SYSCALL_64 () at arch/x86/entry/entry_64.S:120
#13 0x000000000025ee508 in ?? ()
#14 0x000007ffc6ca64730 in ?? ()
```

```
#15 0x0000000000000024 in fixed_percpu_data ()
#16 0x00000000025e9278 in ?? ()
#17 0x00000000025ee508 in ?? ()
#18 0x0000000000000000 in ?? ()
(gdb)
```

Source code's visible and step-able !

```
(gdb) l
2898 {
2899     return slab_alloc_node(s, gfpflags, NUMA_NO_NODE, addr);
2900 }
2901
2902 void *kmem_cache_alloc(struct kmem_cache *s, gfp_t gfpflags)
2903 {
2904     void *ret = slab_alloc(s, gfpflags, _RET_IP_);
2905
2906     trace_kmem_cache_alloc(_RET_IP_, ret, s->object_size,
2907                           s->size, gfpflags);
(gdb)
...
```

Added a breakpoint on schedule_timeout()

```
(gdb) info breakpoints
Num      Type          Disp Enb Address              What
1        breakpoint    keep y   0xffffffff8184a740 in kmem_cache_alloc at mm/slab.c:2903
2        breakpoint    keep y   0xffffffff833d1c70 in schedule_timeout at kernel/time/timer.c:1833
breakpoint already hit 8 times
(gdb) c
Continuing.
Cannot remove breakpoints because program is no longer writable.
Further execution is probably impossible.

Thread 2 hit Breakpoint 2, schedule_timeout (timeout=timeout@entry=125) at kernel/time/timer.c:1833
1833 {
(gdb) bt
#0  schedule_timeout (timeout=timeout@entry=125) at kernel/time/timer.c:1833
#1  0xffffffff81750c17 in freezable_schedule_timeout (timeout=<optimized out>) at ./include/linux/freezer.h:192
#2  kcompactd (p=p@entry=0xffff88807ffb5000) at mm/compaction.c:2818
#3  0xffffffff812152e4 in kthread (_create=0xffff888006b62000) at kernel/kthread.c:292
#4  0xffffffff81006642 in ret_from_fork () at arch/x86/entry/entry_64.S:296
#5  0x0000000000000000 in ?? ()
(gdb) █
...
```

Tip

Also, it's common practise to setup commonly used gdb commands as macros and save them in the gdb initialization file `~/.gdbinit`.

For example:

```
$ cat ~/.gdbinit
set history save on
set history filename ~/.gdb_history
```

```

set output-radix 16
# Careful! disable security protection, auto-load vmlinux-gdb.py
set auto-load safe-path /

#-----
# connect
define connect_serial
    set remotebaud 115200
    target remote /dev/ttyS0
    hb panic
    hb sys_sync
end
define connect_qemu
    target remote :1234
    hb start_kernel
    hb panic
    hb sys_sync
end
...
$

```

SIDEBAR

At times, gdb cannot detect the value of variables and results in:

```

(gdb) p var
$1 = <value optimized out>

```

In these situations, compiling at a high level of optimization might have caused these issues. The kernel is built at '-O2'. Since we cannot realistically rebuild the kernel without optimization, what's to be done?

The only realistic way to debug then is to view the disassembly listing along with the relevant 'C' source and debug. This is easily achieved by entering gdb's TUI mode (notes were provided in the earlier module on gdb usage).

The trick is then to view the local variables and parameters as register values according to the calling convention (and compiler) for that processor architecture.

Also, the 'si' gdb command can be very useful here (when stepping through assembly): step forward a single instruction.

GDB TUI Mode Quick Ref Table

Keyboard key	Action
^X A	Toggle between “regular” GDB and TUI mode
^X 2	Toggle between different TUI layouts: <source-assembly> to <source-registers> to

	<registers-assembly> to <registers-source> to <source-assembly>.
^P	Recall previous command (history) (equivalent to up arrow) at gdb prompt
^N	Recall next command (history) (equivalent to down arrow) at gdb prompt

^ = Ctrl

Note: Linux kernel configurable:

lib/Kconfig.debug

...

config DEBUG_INFO_DWARF4

bool "Generate dwarf4 debuginfo"

depends on DEBUG_INFO

help

Generate dwarf4 debug info. This requires recent versions of gcc and gdb. It makes the debug information larger.

But it significantly improves the success of resolving variables in gdb on optimized code.

...

- If we want to have (or debug) a functioning system with libraries and applications, then obviously we require to build and have QEMU recognize a root filesystem (rfs) image. See the section below which details how to do this.

Sample screenshot of our KGDB session with a Qemu-emulated x86_64 Debian guest (left terminal, white bg) and a x86_64 Ubuntu host (right terminal, dark colour):

```

mm/slab.c
2893         return object;
2894     }
2895
2896     static __always_inline void *slab_alloc(struct kmem_cache *s,
2897         gfp_t gfpflags, unsigned long addr)
2898     {
2899         return slab_alloc_node(s, gfpflags, NUMA_NO_NODE, addr);
2900     }
2901
2902     void *kmem_cache_alloc(struct kmem_cache *s, gfp_t gfpflags)
2903     {
2904         void *ret = slab_alloc(s, gfpflags, _RET_IP_);
2905
2906         trace_kmem_cache_alloc(_RET_IP_, ret, s->object size,
2907             s->size, gfpflags);
2908
2909         return ret;
2910     }
2911     EXPORT_SYMBOL(kmem_cache_alloc);
2912
2913     #ifdef CONFIG_TRACING
2914     void *kmem_cache_alloc_trace(struct kmem_cache *s, gfp_t gfpflags, size_t size)
2915     {
2916         void *ret = slab_alloc(s, gfpflags, _RET_IP_);
2917         trace_kmalloc(_RET_IP_, ret, size, s->size, gfpflags);
2918         ret = kasan_kmalloc(s, ret, size, gfpflags);
2919     }
2920     #endif
2921
remote Thread 1.1 In: kmem_cache_alloc L2906 PC: 0xffffffff8184a7d0
#0 kmem_cache_alloc (s=0xffffffff84d0f240 <boot_kmem_cache>, gfpflags=gfpflags@entry=2304)
   at mm/slab.c:2906
#1 0xffffffff84b878ca in kmem_cache_zalloc (flags=2048, k=<optimized out>)
   at ./include/linux/slab.h:654
#2 bootstrap (static_cache=static_cache@entry=0xffffffff84d0f240 <boot_kmem_cache>)
   at mm/slab.c:4333
#3 0xffffffff84b87a47 in kmem_cache_init () at mm/slab.c:4383
#4 0xffffffff84b11936 in mm_init () at init/main.c:831
#5 0xffffffff84b11af5 in start_kernel () at init/main.c:904
#6 0xffffffff84b10506 in x86_64_start_reservations (
   real_mode_data=real_mode_data@entry=0x13fd0 <exception_stacks+16336> <error: Cannot access m
emory at address 0x13fd0>) at arch/x86/kernel/head64.c:526
#7 0xffffffff84b10584 in x86_64_start_kernel (
   real_mode_data=0x13fd0 <exception_stacks+16336> <error: Cannot access memory at address 0x13
--Type <RET> for more, q to quit, c to continue without paging--])

0.121720] DMA [mem 0x0000000000001000-0x000000000000ffff]
0.122186] DMA32 [mem 0x000000000000100000-0x00000000003ffff]
0.122483] Normal empty
0.122679] Device empty
0.122859] Movable zone start for each node
0.123081] Early memory node ranges
0.123287] node 0: [mem 0x0000000000001000-0x000000000009eff]
0.123776] node 0: [mem 0x0000000000010000-0x00000000003ffff]
0.125028] Zeroed struct page in unavailable ranges: 130 pages
0.125398] Initmem setup node 0 [mem 0x000000000001000-0x00000000003ffff]
0.305808] kasan: KernelAddressSanitizer initialized
0.306994] ACPI: PM-Timer IO Port: 0x608
0.308662] ACPI: LAPIC_NMI (acpi_id[0xff] dfl dfl lint[0x1])
0.309841] IOAPIC[0]: apic_id 0, version 32, address 0xfec00000, GS
0.310467] ACPI: INT_SRC_OVR (bus 0 bus_irq 0 global_irq 2 dfl dfl)
0.311238] ACPI: INT_SRC_OVR (bus 0 bus_irq 5 global_irq 5 high lev
0.311510] ACPI: INT_SRC_OVR (bus 0 bus_irq 9 global_irq 9 high lev
0.311942] ACPI: INT_SRC_OVR (bus 0 bus_irq 10 global_irq 10 high l
0.312198] ACPI: INT_SRC_OVR (bus 0 bus_irq 11 global_irq 11 high l
0.313155] Using ACPI (MADT) for SMP configuration information
0.313628] ACPI: HPET id: 0x8086a201 base: 0xfed00000
0.314852] smpboot: Allowing 2 CPUs, 0 hotplug CPUs
0.317123] PM: hibernation: Registered nosave memory: [mem 0x000000
0.317479] PM: hibernation: Registered nosave memory: [mem 0x0009f0
0.317853] PM: hibernation: Registered nosave memory: [mem 0x000a00
0.318210] PM: hibernation: Registered nosave memory: [mem 0x000f00
0.318681] [mem 0x40000000-0xffffbfff] available for PCI devices
0.319003] Booting paravirtualized kernel on bare hardware
0.319948] clocksource: refined-jiffies: mask: 0xffffffff max_cycle
ns: 7645519600211568 ns
0.321113] setup_percpu: NR_CPUS:8192 nr_cpumask_bits:2 nr_cpu_ids:
0.324042] percpu: Embedded 64 pages/cpu s225280 r8192 d28672 u1048
0.329224] Built 1 zonelists, mobility grouping on. Total pages: 2
0.329564] Policy zone: DMA32
0.330049] Kernel command line: console=ttyS0 root=/dev/sda earlypr
0.333342] Dentry cache hash table entries: 131072 (order: 8, 10485
0.333990] Inode-cache hash table entries: 65536 (order: 7, 524288
0.336950] mem auto-init: stack:off, heap alloc:on, heap free:off
0.337668] AGP: Checking aperture...
0.339403] AGP: No AGP bridge found
0.412702] Memory: 568688K/1048056K available (38921K kernel code,
data, 3272K init, 14252K bss, 479112K re)

```

Sample Session 2 : KGDB on a QEMU ARM System

```
$ qemu-system-arm --help
```

QEMU emulator version 2.5.0 (Debian 1:2.5+dfsg-5ubuntu10.6), Copyright (c) 2003-2008 Fabrice Bellard

usage: qemu-system-arm [options] [disk_image]

```
--snip--
```

```
-S freeze CPU at startup (use 'c' to start execution)
```

```
...
```

```
-gdb dev wait for gdb connection on 'dev'
```

```
-s shorthand for -gdb tcp::1234
```

```
$ cat ./run_kgdb_qemu.sh
```

```
...
```

```
#####
```

```
## UPDATE for your box
```

```
STG=~/.scratchpad/SEALS_staging/
```

```
#####
```

```
ARMPAT=vexpress-a9 ## make sure it's right! ##
```

```
PORT=1235
```

```
qemu-system-arm -m 256 -M ${ARMPLAT} -kernel $1 \
    -drive file=${STG}/images/rfs.img,if=sd,format=raw \
    -append "console=ttyAMA0 root=/dev/mmcblk0 init=/sbin/init" -
nographic \
    -gdb tcp::${PORT} -S
```

...

```
$ ./run_kgdb_qemu.sh </path/to/>zImage
```

REMEMBER this qemu instance is run w/ the **-S QEMU switch: it *waits*** for a gdb client to connect to it...

You are expected to run (in another terminal window):

```
$ arm-none-linux-gnueabi-gdb <path-to-ARM-built-kernel-src-tree>/vmlinux #
<--built w/ -g
```

...

```
# and then have gdb connect to the target kernel using
(gdb) target remote :1235
```

...

<< On the ARM target >>

```
pulseaudio: set_sink_input_volume() failed
pulseaudio: Reason: Invalid argument
pulseaudio: set_sink_input_mute() failed
pulseaudio: Reason: Invalid argument
<< it's waiting for GDB to connect >>
```

In another terminal window:

<< On the x86 host >>

```
$ cd <ARM-kernel-source-tree>
```

```
$ ls -lh vmlinux
```

```
-rwxr-xr-x 1 root root 78M Jan 19 12:46 vmlinux
```

```
$ file vmlinux
```

```
vmlinux: ELF 32-bit LSB executable, ARM, EABI5 version 1 (SYSV), statically
linked, BuildID[sha1]=abc..., not stripped
```

```
$
```

```
$ arm-none-linux-gnueabi-gdb ./vmlinux
```

```
GNU gdb (Sourcery CodeBench Lite 2014.05-29) 7.7.50.20140217-cvs
```

```
Copyright (C) 2014 Free Software Foundation, Inc.
```

```
License GPLv3+: GNU GPL version 3 or later
```

```
<http://gnu.org/licenses/gpl.html>
```

...

```
Type "apropos word" to search for commands related to "word"...
```

```
Reading symbols from ./vmlinux...done.
```

```
(gdb) target remote :1234
```

```
:1234: Connection timed out. << Oops, we gave the wrong port# (see above)
```

```
>>
```

```
(gdb) target remote :1235 << attempt to connect to the (QEMU) ARM
target >>
```


Remote debugging using :1235

0x60000000 in ?? () *<< connected! >>*

(gdb) bt

#0 0x60000000 in ?? ()

(gdb) b vfs_write *<< set a breakpoint >>*

Breakpoint 1 at 0x80139a00: file fs/read_write.c, line 519.

(gdb) c

Continuing.

<< On the ARM target >>

...

[0.000000] Booting Linux on physical CPU 0x0

[0.000000] Initializing cgroup subsys cpuset

[0.000000] Linux version 3.18.22-kgdb (root@kaiwan-T460) (gcc version 4.8.3 20140320 (prerelease) (Sourcery CodeBench Lite 2014.05-29)) #2 SMP Thu Jan 19 12:46:36 IST 2017

[0.000000] CPU: ARMv7 Processor [410fc090] revision 0 (ARMv7), cr=10c5387d

[0.000000] CPU: PIPT / VIPT nonaliasing data cache, VIPT nonaliasing instruction cache

[0.000000] Machine: ARM-Versatile Express

[0.000000] Memory policy: Data cache writeback

[0.000000] CPU: All CPU(s) started in SVC mode.

[0.000227] sched_clock: 32 bits at 24MHz, resolution 41ns, wraps every 178956969942ns

[0.001330] PERCPU: Embedded 11 pages/cpu @8fde9000 s12352 r8192 d24512 u45056

[0.003103] Built 1 zonelists in Zone order, mobility grouping on.

Total pages: 65024

[0.003234] Kernel command line: console=ttyAMA0 root=/dev/mmcblk0 init=/sbin/init

[0.003942] PID hash table entries: 1024 (order: 0, 4096 bytes)

[0.004142] Dentry cache hash table entries: 32768 (order: 5, 131072 bytes)

[0.004510] Inode-cache hash table entries: 16384 (order: 4, 65536 bytes)

[0.008112] Memory: 244816K/262144K available (6163K kernel code, 339K rwddata, 1872K rodata, 436K init, 6169K bss, 17328K reserved)

[0.008229] Virtual kernel memory layout:

[0.008229] vector : 0xfffff000 - 0xfffff1000 (4 kB)

[0.008229] fixmap : 0xffc00000 - 0xffe00000 (2048 kB)

[0.008229] vmalloc : 0x90800000 - 0xff000000 (1768 MB)

[0.008229] lowmem : 0x80000000 - 0x90000000 (256 MB)

[0.008229] modules : 0x7f000000 - 0x80000000 (16 MB)

[0.008229] .text : 0x80008000 - 0x807e0edc (8036 kB)

[0.008229] .init : 0x807e1000 - 0x8084e000 (436 kB)

[0.008229] .data : 0x8084e000 - 0x808a2f38 (340 kB)

[0.008229] .bss : 0x808a2f38 - 0x80ea93f4 (6170 kB)

[0.011639] SLUB: HWalign=64, Order=0-3, MinObjects=0, CPUs=1, Nodes=1

--snip--


```

/devices/mb:kmil/seriol/input/input2
[ 1.927263] EXT3-fs (mmcblk0): error: couldn't mount because of
unsupported optional features (240)
[ 1.935603] EXT2-fs (mmcblk0): error: couldn't mount because of
unsupported optional features (240)
[ 1.968311] EXT4-fs (mmcblk0): mounted filesystem with ordered data
mode. Opts: (null)
[ 1.969544] VFS: Mounted root (ext4 filesystem) readonly on device
179:0.
[ 1.980846] Freeing unused kernel memory: 436K (807e1000 - 8084e000)
[ 2.143933] random: nonblocking pool is initialized
<< breakpoint hit, waiting to be continued >>

```

<< On the x86 host >>

```

Breakpoint 1, vfs_write (file=0x8fbc7780, buf=0xf8f88 "/etc/init.d/rcS
running now ...\n", count=32, pos=0x8f9eff78)
    at fs/read_write.c:519
519  {
(gdb) bt
#0  vfs_write (file=0x8fbc7780, buf=0xf8f88 "/etc/init.d/rcS running
now ...\n", count=32, pos=0x8f9eff78)
    at fs/read_write.c:519
#1  0x8013a00c in SYSC_write (count=<optimized out>, buf=<optimized out>,
fd=<optimized out>) at fs/read_write.c:585
#2  Sys_write (fd=<optimized out>, buf=1019784, count=32) at
fs/read_write.c:577
#3  0x8000f980 in ?? ()
Backtrace stopped: frame did not save the PC
(gdb) l
514  }
515
516  EXPORT_SYMBOL(__kernel_write);
517
518  ssize_t vfs_write(struct file *file, const char __user *buf, size_t
count, loff_t *pos)
519  {
520      ssize_t ret;
521
522      if (!(file->f_mode & FMODE_WRITE))
523          return -EBADF;
(gdb) p file
$1 = (struct file *) 0x8fbc7780
(gdb) set print pretty
(gdb) p *file
$2 = {
  f_u = {
    fu_llist = {
      next = 0x0 <__vectors_start>
    },

```

```
[...]
```

```
f_path = {
    mnt = 0x8f80f0d0,
    dentry = 0x8f61f170
},
f_inode = 0x8f638aa0,
f_op = 0x805ce634 <console_fops>,
```

```
[...]
```

```
---Type <return> to continue, or q <return> to quit---q
```

```
Quit
```

```
(gdb)
```

```
...
```

```
Breakpoint 1, vfs_write (file=0x8fbf9780, buf=0xf9060 "8 4 1 7\n", count=8,
pos=0x8f9eff78) at fs/read_write.c:519
```

```
519 {
```

```
(gdb) info breakpoints
```

```
Num      Type           Disp Enb Address      What
1        breakpoint      keep y   0x80139a00 in vfs_write at
fs/read_write.c:519
```

```
breakpoint already hit 2 times
```

```
(gdb) disable 1
```

```
(gdb) info breakpoints
```

```
Num      Type           Disp Enb Address      What
1        breakpoint      keep n   0x80139a00 in vfs_write at
fs/read_write.c:519
```

```
breakpoint already hit 2 times
```

```
(gdb) c
```

```
Continuing.
```

```
...
```

```
<< Want to setup another breakpoint; so we 'break' into GDB with ^C >>
```

```
^C
```

```
Program received signal SIGINT, Interrupt.
```

```
cpu_v7_do_idle () at arch/arm/mm/proc-v7.S:74
```

```
74      ret    lr
```

```
(gdb) b kmem_cache_free
```

```
Breakpoint 3 at 0x80131e80: file mm/slub.c, line 2679.
```

```
(gdb) c
```

```
Continuing.
```

```
...
```

```
Breakpoint 3, kmem_cache_free (s=0x8f801b80, x=0x8f81e000) at
mm/slub.c:2679 << hit ! >>
```

```
2679 {
```

```
(gdb) bt
```

```
#0  kmem_cache_free (s=0x8f801b80, x=0x8f81e000) at mm/slub.c:2679
#1  0x80148840 in final_putname (name=0x8f81e000) at fs/namei.c:127
#2  0x8013892c in do_sys_open (dfd=-1895797304, filename=<optimized out>,
    flags=-1895797312, mode=<optimized out>)
    at fs/open.c:1007
#3  0x801389d0 in SYSC_openat (mode=<optimized out>, flags=<optimized out>,
    filename=<optimized out>,
    dfd=<optimized out>) at fs/open.c:1025
#4  Sys_openat (dfd=<optimized out>, filename=<optimized out>,
    flags=<optimized out>, mode=<optimized out>)
    at fs/open.c:1019
#5  0x8000f980 in ?? ()
Backtrace stopped: frame did not save the PC
(gdb) l
2674         __slab_free(s, page, x, addr);
2675
2676     }
2677
2678 void kmem_cache_free(struct kmem_cache *s, void *x)
2679 {
2680     s = cache_from_obj(s, x);
2681     if (!s)
2682         return;
2683     slab_free(s, virt_to_head_page(x), x, _RET_IP_);
```

We can see in the above function (`kmem_cache_free`) that the first parameter is a pointer to `struct kmem_cache`. From the current ARM Procedure Call Standard (APCS) document (defined by ARM Ltd [in this PDF document](#)), we know that the first four parameters will be in registers `r0` through `r3`.

Thus the first parameter should be in the **r0** register:

```
(gdb) info registers
r0          0x8f801b80 -1887429760
r1          0x8f81e000 -1887313920
r2          0x0      0

[...]

sp          0x8f043f38 0x8f043f38
lr          0x80148840 -2146138048
pc          0x80131e80 0x80131e80 <kmem_cache_free>
cpsr       0x60000013 1610612755
(gdb) p s
$1 = (struct kmem_cache *) 0x8f801b80    << indeed it is! >>
(gdb) p *s
$2 = {cpu_slab = 0x8084f0d0 <init_thread_union+4304>, flags = 0,
    min_partial = 6, size = 4096, object_size = 4096,
    offset = 0, cpu_partial = 2, oo = {x = 196616}, max = {x = 196616}, min =
    {x = 1}, allocflags = 16384,
```

```
    refcount = 2, ctor = 0x0 <__vectors_start>, inuse = 4096, align = 64,
    reserved = 0,
    name = 0x8f8021c0 "kmalloc-4096", list = {next = 0x8f801c44, prev =
0x8f801b44}, kobj = {
    name = 0x8fab5900 ":t-0004096", entry = {next = 0x8f801c50, prev =
0x8f801b50}, parent = 0x8f9d9e28,
    ...
(gdb) set print pretty
...
```

<< On the ARM target >>

/bin/sh: can't access tty; job control turned off

ARM / \$

ARM / \$ halt

ARM / \$ [7.653539] EXT4-fs (mmcblk0): re-mounted. Opts: (null)

The system is going down NOW!

Sent SIGTERM to all processes

Sent SIGKILL to all processes

Requesting system halt

[9.706334] Flash device refused suspend due to active operation (state 20)

[9.707181] Flash device refused suspend due to active operation (state 20)

[9.708762] reboot: System halted

QEMU: Terminated

\$

<< On the x86 host >>

Remote connection closed

(gdb) q

\$

Debugging a Kernel Module with KGDB

In order to debug a kernel module within KGDB, the only difference from debugging the “regular” in-line kernel code, is that the **kernel module's ELF sections information has to be provided to gdb**, so that gdb can “see” it.

We can achieve this quite easily:

1. (Cross) Compile the kernel module with debug symbols:
In the Makefile, keep the line

```
EXTRA_CFLAGS += -DDEBUG -g -ggdb
or (modern):
ccflags-y += -DDEBUG -g -ggdb
```

(Caution: above, do NOT add -O0 ; the kernel and modules MUST be compiled with optimization on, usually -O2; this is unlike userspace...)

2. Copy the kernel module and the *gdbline.sh* shell script (seen earlier in this module) onto the target
3. With the KGDB session active:
 1. Load the kernel module into RAM
`insmod <mykm.ko>`
 2. Query it's ELF sections using
`/sys/module/<module-name>/sections/.<section>` either directly (or, easier, via the *gdbline.sh* shell script seen earlier).

An Example Session:

Using a custom network driver kernel module and the *gdbline.sh* script

On the target:

```
ARM /myprj/veth_dyn_ops $ ../gdbline.sh veth_dyn_ops ./veth_dyn_ops.ko
add-symbol-file ./veth_dyn_ops.ko 0xbf000000 -s .alt.smp.init 0xbf000f34 -s
.bss 0xbf001a10 -s .data 0xbf001758 -s .exit.text 0xbf000dac
[ ... ] -s __mcount_loc 0xbf00165c
ARM /myprj/veth_dyn_ops $
```

On the host gdb (in this case, it's the x86-to-ARM cross-compiler gdb):

<< cd to the source folder of the kernel module you're going to debug >>

```
(gdb) cd <...>/3_veth_dyn_ops
Working directory <...>/3_veth_dyn_ops
(canonically <.../>/3_veth_dyn_ops).
```

(gdb)*<< Copy-paste the output of the gdbline.sh shell script >>*

```
(gdb) add-symbol-file ./veth_dyn_ops.ko 0xbf000000 -s .alt.smp.init
0xbf000f34 -s .bss 0xbf001a10 -s .data 0xbf001758 [ ... ] -s __mcount_loc
0xbf00165c
```

```
add symbol table from file "./veth_dyn_ops.ko" at
```

```
  .text_addr = 0xbf000000
  .alt.smp.init_addr = 0xbf000f34
  .bss_addr = 0xbf001a10
  .data_addr = 0xbf001758
```

```
...
```

```
  __mcount_loc_addr = 0xbf00165c
```

```
(y or n) y
```

```
Reading symbols from <path/to/>/veth_dyn_ops.ko...warning: section .strtab
not found in <path/to/>/veth_dyn_ops.ko
```

```
warning: section .symtab not found in <path/to/>/veth_dyn_ops.ko
done.
```

```
(gdb) b vnet_ <<tab-tab>>
```

```
vnet_exit          vnet_init          vnet_open          vnet_remove
vnet_stop
vnet_get_stats     vnet_netdev_ops  vnet_probe         vnet_start_xmit
vnet_tx_timeout
```

```
(gdb) b vnet_start_xmit
```

```
Breakpoint 1 at 0xbf0003fc: file <path/to/>/veth_netdrv.c, line 73.
```

```
(gdb) c
```

```
Continuing.
```

<< Have some data sent to the network interface in question, so that the kernel invokes the network driver's 'start_xmit' method.. ... >>

```
Breakpoint 1, vnet_start_xmit (skb=0xc78819c0, dev=0xc61de800)
```

```
  at <path/to/>/veth_netdrv.c:73
```

```
73      if (!skb) { // paranoia!
```

```
(gdb) bt
```

```
#0  vnet_start_xmit (skb=0xc78819c0, dev=0xc61de800)
```

```
  at <path/to/>/veth_netdrv.c:73
```

```
#1  0xc036cd78 in dev_hard_start_xmit (skb=0xc78819c0, dev=0xc61de800,
txq=0xc11e1120) at net/core/dev.c:2222
```

```
#2  0xc0387930 in sch_direct_xmit (skb=0xc78819c0, q=0xc625e200,
dev=0xc61de800, txq=0xc11e1120, root_lock=0xc625e25c) at
```

```
net/sched/sch_generic.c:125
```

```
#3  0xc036d224 in __dev_xmit_skb (skb=0xc78819c0) at net/core/dev.c:2427
```

```
#4  dev_queue_xmit (skb=0xc78819c0) at net/core/dev.c:2502
```

```
#5  0xc0379880 in neigh_resolve_output (neigh=0xc7881900, skb=0xc78819c0)
at net/core/neighbour.c:1274
```

```
#6  0xc039c2d8 in neigh_output (skb=0xc78819c0) at
include/net/neighbour.h:352
```

```
#7  ip_finish_output2 (skb=0xc78819c0) at net/ipv4/ip_output.c:211
```

```
#8  ip_finish_output (skb=0xc78819c0) at net/ipv4/ip_output.c:244
```

```
#9 0xc039c5e8 in NF_HOOK_COND (skb=0xc78819c0) at
include/linux/netfilter.h:233
#10 ip_output (skb=0xc78819c0) at net/ipv4/ip_output.c:317
#11 0xc039b59c in dst_output (skb=0xc78819c0) at include/net/dst.h:430
#12 ip_local_out (skb=0xc78819c0) at net/ipv4/ip_output.c:111
#13 0xc039b5bc in ip_send_skb (skb=0xc78819c0) at net/ipv4/ip_output.c:1371
#14 0xc03be734 in udp_send_skb (skb=0xc78819c0, fl4=0xc11cfd5c) at
net/ipv4/udp.c:753
#15 0xc03c0630 in udp_sendmsg (iocb=<value optimized out>, sk=0xc11d4220,
msg=0xc11cff54, len=16) at net/ipv4/udp.c:966
#16 0xc03c92c4 in inet_sendmsg (iocb=0xc11cfde8, sock=<value optimized
out>, msg=0x4000, size=268503562) at net/ipv4/af_inet.c:744
#17 0xc0357724 in __sock_sendmsg_nosec (sock=0xc7404d80, msg=0xc11cff54,
size=16) at net/socket.c:557
#18 __sock_sendmsg (sock=0xc7404d80, msg=0xc11cff54, size=16) at
net/socket.c:565
#19 sock_sendmsg (sock=0xc7404d80, msg=0xc11cff54, size=16) at
net/socket.c:576
#20 0xc0357f18 in sys_sendto (fd=<value optimized out>, buff=0xc11cfed4,
len=16, flags=0, addr=0xbe845cdc, addr_len=16) at net/socket.c:1708
#21 0xc000ed80 in ?? ()
Cannot access memory at address 0x0
#22 0xc000ed80 in ?? ()
Cannot access memory at address 0x0
Backtrace stopped: previous frame identical to this frame (corrupt stack?)
```

(gdb) l

```
68     struct udphdr *udph;
69     pstVnetIntfCtx pstCtx = netdev_priv (dev);
70     //struct timespec ts1, ts2, tdiff;
71     struct timeval ts1, ts2, tdiff;
72
73     if (!skb) { // paranoia!
74         printk (KERN_ALERT "%s:%s: skb NULL!\n", DRVNAME,
__func__);
75         return -EAGAIN;
76     }
77     // SKB_PEEK(skb);
...
...
97     ip = ip_hdr(skb);
```

(gdb) p ip

```
$3 = (struct iphdr *) 0xc11c3e10
```

(gdb) set print pretty

(gdb) p/x *ip

```
$5 = {
  ihl = 0x5,
  version = 0x4,
  tos = 0x0,
  tot_len = 0x2c00,
  id = 0x0,
  frag_off = 0x40,
  ttl = 0x40,
```

```

    protocol = 0x11,
    check = 0x9924,
    saddr = 0x5010a0a,
    daddr = 0x10010a0a
}
(gdb)

```

<< *Change an existing variable's value!* >>

(gdb) help set variable

Evaluate expression EXP and assign result to variable VAR, using assignment syntax appropriate for the current language (VAR = EXP or VAR := EXP for example). VAR may be a debugger "convenience" variable (names starting with \$), a register (a few standard names starting with \$), or an actual variable in the program being debugged. EXP is any valid expression. This may usually be abbreviated to simply "set".

(gdb) p dev

```
$13 = (struct net_device *) 0xcfa5f800
```

(gdb) p dev->promiscuity

```
$11 = 0x0
```

(gdb) set variable dev->promiscuity = 1

← *here!*

(gdb) p dev->promiscuity

```
$12 = 0x1
```

(gdb)

Note:

From (around) the 3.20 Linux kernel, (and GDB ver >= 7.2 required) we have an interesting Kernel configurable:

lib/Kconfig.debug

```
...
```

```
config GDB_SCRIPTS
```

```
    bool "Provide GDB scripts for kernel debugging"
```

```
    depends on DEBUG_INFO
```

```
    help
```

This creates the required links to GDB helper scripts in the build directory. If you load vmlinux into gdb, the helper scripts will be automatically imported by gdb as well, and additional functions are available to analyze a Linux kernel instance. See Documentation/gdb-kernel-debugging.txt for

further

details.

See [Documentation/gdb-kernel-debugging.txt](#).

...

List of commands and functions

The number of commands and convenience functions may evolve over the time, this is just a snapshot of the initial version:

```
(gdb) apropos lx
function lx_current -- Return current task
function lx_module -- Find module by name and return the module variable
function lx_per_cpu -- Return per-cpu variable
function lx_task_by_pid -- Find Linux task by PID and return the
task_struct variable
function lx_thread_info -- Calculate Linux thread_info from task variable
lx-dmesg -- Print Linux kernel log buffer
lx-lsmod -- List currently loaded modules
lx-symbols -- (Re-)load symbols of Linux kernel and currently loaded
modules
```

Detailed help can be obtained via "help <command-name>" for commands and "help function <function-name>" for convenience functions.

SIDEBAR ::

The MAINTAINERS file


Very useful!

I once had an issue with getting the Python-based GDB scripts running on an ARM-based system. Looked up the MAINTAINERS file, grepping for help; I soon found it!


```
linux-5.10.3 $ grep -i -A5 "gdb " MAINTAINERS
GDB KERNEL DEBUGGING HELPER SCRIPTS
M:      Jan Kiszka <jan.kiszka@siemens.com>
M:      Kieran Bingham <kbingham@kernel.org>
S:      Supported
F:      scripts/gdb/

--
KGDB / KDB /debug_core
M:      Jason Wessel <jason.wessel@windriver.com>
M:      Daniel Thompson <daniel.thompson@linaro.org>
R:      Douglas Anderson <dianders@chromium.org>
L:      kgdb-bugreport@lists.sourceforge.net
S:      Maintained
linux-5.10.3 $
```

Wrote to the maintainers, and had a good reply within an hour!!


Kaiwan N Billimoria
Tue, Jul 18, 2017, 4:00 PM

Hi, Am trying to make use of the GDB helpers 'lx-*' via the CONFIG_GDB_SCRIPTS kernel option. Host sys: Ubuntu 17.04 x86_64 4.10...


Kieran Bingham kieran@bingham.xyz [via ksquared.org.uk](https://ksquared.org.uk)
Tue, Jul 18, 2017, 4:19 PM

to me, jan.kiszka

Hi Kaiwan,


On 18/07/17 11:30, Kaiwan N Billimoria wrote:

- > Hi,
- >
- > Am trying to make use of the GDB helpers 'lx-*' via the
- > CONFIG_GDB_SCRIPTS kernel option.
- > Host sys: Ubuntu 17.04 x86_64 4.10.0-26-generic.
- >
- > Steps I followed are as per the documentation here:
- > <https://www.kernel.org/doc/html/v4.11/dev-tools/gdb-kernel-debugging.html>
- >
- > Yet, it does not seem to work: gdb doesn't recognize any 'lx*' commands.
- > Does it work for ARM-32 ??

Yes, I have used gdb-scripts on both ARM32, and ARM64

...

...


Jan Kiszka
Jul 18, 2017, 5:44 PM

Ask CodeSourcery folks, can't tell. But the Linaro toolchain supports this mode.

Test run on an x86_64 VM:

```
tty0 $ ls -lh vmlinux
-rwxrwxr-x 1 seawolf seawolf 439M Jul 16 11:18 vmlinux* << vmlinux with
debugsym >>
tty0 $ ls -l vmlinux*
```

```
-rwxrwxr-x 1 seawolf seawolf 459318832 Jul 16 11:18 vmlinux*
lrwxrwxrwx 1 seawolf seawolf      58 Jul 16 11:16 vmlinux-gdb.py ->
<...>/linux-4.4.21/scripts/gdb/vmlinux-gdb.py << notice the slink ! >>
-rw-rw-r-- 1 seawolf seawolf 542215936 Jul 16 11:16 vmlinux.o
```

```
tty0 $ cat ~/.gdbinit << setup your ~/.gdbinit to auto-load the python
scripts >>
```

```
# ~/.gdbinit
```

```
# For GDB helper scripts pkg w/ modern kernels (GDB_SCRIPTS)
add-auto-load-safe-path <...>linux-4.4.21/scripts/gdb/vmlinux-gdb.py
tty0 $
```

```
tty0 $ gdb ./vmlinux
```

```
GNU gdb (Ubuntu 7.11.90.20161005-0ubuntu2) 7.11.90.20161005-git
```

```
[...]
```

```
Type "apropos word" to search for commands related to "word"...
```

```
Reading symbols from ./vmlinux...done.
```

```
(gdb) info auto-load
```

```
gdb-scripts: No auto-load scripts.
```

```
libthread-db: No auto-loaded libthread-db.
```

```
local-gdbinit: Local .gdbinit file was not found.
```

```
python-scripts:
```

```
Loaded Script
```

```
Yes /home/seawolf/0tmp/linux-4.4.21/vmlinux-gdb.py
```

```
(gdb) apropos lx
```

```
function lx_current -- Return current task << the lx-* gdb goodies are now
available ! >>
```

```
function lx_module -- Find module by name and return the module variable
```

```
function lx_per_cpu -- Return per-cpu variable
```

```
function lx_task_by_pid -- Find Linux task by PID and return the
```

```
task_struct variable
```

```
function lx_thread_info -- Calculate Linux thread_info from task variable
```

```
lx-dmesg -- Print Linux kernel log buffer
```

```
lx-list-check -- Verify a list consistency
```

```
lx-lsmod -- List currently loaded modules
```

```
lx-ps -- Dump Linux tasks
```

```
lx-symbols -- (Re-)load symbols of Linux kernel and currently loaded
```

```
modules
```

```
(gdb)
```

It may **not** always be the case that GDB supports python scripts. For eg., on a cross-compile GDB that I use at times for embedded Linux work, the support was not present:

```
$ arm-none-linux-gnueabi-gdb ./vmlinux
```

```
GNU gdb (Sourcery CodeBench Lite 2014.05-29) 7.7.50.20140217-cvs << GDB ver
```

```
>= 7.2 yet ... >>
```

```
[...]
```

```
Type "apropos word" to search for commands related to "word"...
```

```
Reading symbols from ./vmlinux...done.
```

```
(gdb) python import sys
```

```
<< check for python scripting support >>
```

```
Python scripting is not supported in this copy of GDB.
```

(gdb)

(Tip: Looks to be an issue with the Code Sourcery (Mentor Graphics) toolchain (this ver only??); Linaro toolchains support python scripting).

[OPTIONAL/FYI]

KDB

<<

Refer “Linux Debugging and Performance Tuning” by Steve Best, Ch 13 “Kernel-Level Debuggers (kgdb and kdb)” section “kdb” page 348 onwards.

>>

Additional Notes for kernel-level debuggers topic: kdb

Author: kaiwan.

Earlier, KDB had to be patched into a Linux kernel via an arch-dependent and an arch-independent patch.

Recent Linux kernels:

The procedure below to apply the two kdb patches is now outdated. That's good news: **kdb is now in mainline (from ver 2.6.26) and is integrated as a “frontend” into the KGDB system.**

[What is different about the old KDB vs the merged KDB?](#)

In April 2009 [KDB 4.4](#) had significant chunks of the code base removed and hooked it up to the same debug core and polled I/O drivers used by kgdb.

At a really high level, the only difference is how you configure and connect. The kdb disassembler was the main casualty of the merger. Architecturally the merged kdb is 99% platform independent and actually gained new features via the debug core API.

...

The kdb command shell differences

- There are fewer commands in the merged kdb & kgdb.
 - The kdump / kexec analysis modules tools were removed in the process of merging the code bases.

- The status command emits a bit less information in that you no longer get the exact output you would from `cat /proc/meminfo`
- The `bt` command does use the kernel's backtracer and not a disassembly engine which shows you function arguments
- In the merged `kdb` & `kgdb`, **presently there is no disassembler** (known as the "id" command)
- In the merged `kdb` & `kgdb`, when using a non-vga style connection you can transition into `kgdb` mode and attach `gdb`, by using the `kgdb` command.
- The ability to reference elements of a structures was removed. This may get implemented a different way in the future.

Architecture Support:

- KDB 4.4 supported only `ia64`, `x86_64` and `i386`
- The merged `kdb` & `kgdb` supports `arm`, `blackfin`, `mips`, `x86_64`, `i386`, `sh`, `powerpc`, and `sparc`.

Kernel Configuration

```
$ make ARCH=<arch> menuconfig
```

```
...
```

Under 'Kernel Hacking / KGDB: kernel debugger':

i.e. select the option labelled:

KGDB_KDB : include kdb frontend for kgdb

, as shown above.

The relevant configuration options are:

```
# grep -i kdb .config
CONFIG_KGDB_KDB=y
CONFIG_KDB_KEYBOARD=y
CONFIG_KDB_CONTINUE_CATASTROPHIC=0
#
```

Save & build. When booted, the kernel now supports debugging via both the KGDB and KDB interfaces.

How do we enter kdb?

There are a few ways (several described below); the “new” technique is conveniently captured by the shell script below:

```
$ cat kdb.sh
#!/bin/sh
# kdb.sh
# Invoke KDB - kernel debugger
# Of course, we have to have KGDB/KDB compiled into the kernel
# Relevant k config options:
# CONFIG_KGDB_KDB=y
# CONFIG_KDB_KEYBOARD=y
```

```
# Enable console for KDB use
# (use tty02 for the pandaboard; ttyAMA0 for QEMU/ARM-vexpress)
CONSOLEDEV=ttyAMA0

echo ${CONSOLEDEV} > /sys/module/kgdboc/parameters/kgdboc || {
    echo "Failed."
    exit 1
}

# Invoke KDB (via Magic-SysRq key 'g' (debug))
echo "+++ Invoking KDB now...(type 'go' @ kdb prompt to exit kdb)"
echo g > /proc/sysrq-trigger
$
```

Run this script on the target and you'll end up at the kdb prompt!

A Quick KDB Session using ARM/Linux on QEMU

```
ARM / # echo g > /proc/sysrq-trigger
```

```
SysRq : HELP : loglevel(0-9) reBoot Crash terminate-all-tasks(E) memory-
full-oom-kill(F) kill-all-tasks(I) thaw-filestems(J) saK show-backtrace-
all-active-cpus(L) show-memory-usage(M) nice-all-RT-tasks(N) powerOff show-
registers(P) show-all-timers(Q) unRaw Sync show-task-states(T) Unmount
show-blocked-tasks(W) dump-ftrace-buffer(Z)
```

```
ARM / #
```

Doesn't work ?? Try this:

```
ARM / # echo ttyAMA0 > /sys/module/kgdboc/parameters/kgdboc
```

```
kgdb: Registered I/O driver kgdboc.
```

```
ARM / #
```

```
ARM / # echo g > /proc/sysrq-trigger << now it works! >>
```

```
SysRq : DEBUG
```

Entering kdb (current=0xcf8a54a0, pid 489) on processor 0 due to Keyboard

Entry

```
[0]kdb> ?
```

Command	Usage	Description
md	<vaddr>	Display Memory Contents, also mdWcN,
e.g. md8c1		
mdr	<vaddr> <bytes>	Display Raw Memory
mdp	<paddr> <bytes>	Display Physical Memory
mds	<vaddr>	Display Memory Symbolically
mm	<vaddr> <contents>	Modify Memory Contents
go	[<vaddr>]	Continue Execution
rd		Display Registers
rm	<reg> <contents>	Modify Registers
ef	<vaddr>	Display exception frame

bt	[<vaddr>]	Stack traceback
btp	<pid>	Display stack for process <pid>
bta	[D R S T C Z E U I M A]	Backtrace all processes matching state
flag		
btc		Backtrace current process on each cpu
btt	<vaddr>	Backtrace process given its struct task
address		
env		Show environment variables
set		Set environment variables
help		Display Help Message
?		Display Help Message
cpu	<cpunum>	Switch to new cpu
kgdb		Enter kgdb mode
ps	[<flags> A]	Display active task list
pid	<pidnum>	Switch to another task
reboot		Reboot the machine immediately
lsmod		List loaded kernel modules
sr	<key>	Magic SysRq key
dmesg	[lines]	Display syslog buffer
defcmd	name "usage" "help"	Define a set of commands, down to
endefcmd		
kill	<-signal> <pid>	Send a signal to a process
summary		Summarize the system
per_cpu	<sym> [<bytes>] [<cpu>]	Display per_cpu variables
grep		Display help on grep
bp	[<vaddr>]	Set/Display breakpoints
bl	[<vaddr>]	Display breakpoints
bc	<bpnum>	Clear Breakpoint
be	<bpnum>	Enable Breakpoint
bd	<bpnum>	Disable Breakpoint
ss		Single Step
dumpcommon		Common kdb debugging
dumpall		First line debugging
dumpcpu		Same as dumpall but only tasks on cpus
ftdump	[skip_#lines] [cpu]	Dump ftrace log
[0]kdb> summary		<< kdb: summ: Summarize the system >>
sysname	Linux	
release	3.18.22-kgdb	
version	#2 SMP Thu Jan 19 12:46:36 IST 2017	
machine	armv7l	
nodename	(none)	
domainname	(none)	
ccversion	CCVERSION	
date	2017-01-19 10:25:03 tz_minuteswest 0	
uptime	00:04	
load avg	0.00 0.00 0.00	
MemTotal:	245252 kB	
MemFree:	233244 kB	
Buffers:	156 kB	

```

[0]kdb> ps                << kdb: ps: Display active task list >>
30 sleeping system daemon (state M) processes suppressed,
use 'ps A' to see all.
Task Addr      Pid    Parent [*] cpu State Thread      Command
0x8f911e00      609        1  1    0   R  0x8f9120e0  *sh

0x8f880000        1        0  0    0   S  0x8f8802e0  init
0x8f911e00      609        1  1    0   R  0x8f9120e0  *sh

[0]kdb> md 0x8f880000     << kdb: md: Display Memory Contents, also mdWcN,
e.g. md8c1 >>
0x8f880000 00000001 8f844000 00000002 00400100 .....@.....@.
0x8f880010 00000000 00000000 00000000 8fbe8a00 .....
0x8f880020 0000021c 00000000 00000000 00000000 .....
0x8f880030 00000078 00000078 00000078 00000000 x...x...x.....
0x8f880040 8059adec 00000000 00000400 00400000 ..Y.....@.
0x8f880050 00000001 00000000 00000000 8f88005c .....\.
0x8f880060 8f88005c 00000000 ff197121 00000000 \.....!q.....
0x8f880070 32e56250 00000000 7e74b0ec 00000000 Pb.2.....t~....

[0]kdb> rd                << kdb: rd: Display Registers >>
r0: 00000067 r1: 00000000 r2: 90800000 r3: 90800730 r4: 80e36b40
r5: 80e36b44 r6: 80856448 r7: 8f20a010 r8: 00000000 r9: 00000008
r10: 00000000 fp: 8f20bebc ip: 8f20bea0 sp: 8f20bea0 lr: 800a62b4
pc: 800a62b4 f0: ?? f1: ?? f2: ?? f3: ?? f4: ?? f5: ?? f6: ??
f7: ??
fps: 00000000 cpsr: a0000013

[0]kdb> btc              << kdb: btc: Backtrace current process on each cpu >>
btc: cpu status: Currently on cpu 0
Available cpus: 0
Stack traceback for pid 609
0x8f911e00      609        1  1    0   R  0x8f9120e0  *sh
[<80017cdc>] (unwind_backtrace) from [<80013c5c>] (show_stack+0x20/0x24)
[<80013c5c>] (show_stack) from [<800b00bc>] (kdb_show_stack+0x54/0x68)
[<800b00bc>] (kdb_show_stack) from [<800b016c>] (kdb_bt1.isra.0+0x9c/0xe8)
[<800b016c>] (kdb_bt1.isra.0) from [<800b0398>] (kdb_bt+0x1e0/0x454)
[<800b0398>] (kdb_bt) from [<800ad99c>] (kdb_parse+0x2cc/0x6c4)
[<800ad99c>] (kdb_parse) from [<800b0450>] (kdb_bt+0x298/0x454)
[<800b0450>] (kdb_bt) from [<800ad99c>] (kdb_parse+0x2cc/0x6c4)
[<800ad99c>] (kdb_parse) from [<800ae57c>] (kdb_main_loop+0x530/0x790)
[<800ae57c>] (kdb_main_loop) from [<800b1174>] (kdb_stub+0x230/0x488)
[<800b1174>] (kdb_stub) from [<800a6e8c>] (kgdb_cpu_enter+0x43c/0x730)
[<800a6e8c>] (kgdb_cpu_enter) from [<800a7464>]
(kgdb_handle_exception+0x194/0x1e0)
[<800a7464>] (kgdb_handle_exception) from [<80017038>]
(kgdb_compiled_brk_fn+0x38/0x40)
[<80017038>] (kgdb_compiled_brk_fn) from [<800083fc>]
(do_undefinstr+0x17c/0x214)
[<800083fc>] (do_undefinstr) from [<80594c10>] (__und_svc_finish+0x0/0x30)
Exception stack(0x8f20be18 to 0x8f20be60)

```



```

be00: 00000000 00000000
be20: 90800000 90800730 80e36b40 80e36b44 80856448 8f20a010 00000000
00000008
be40: 00000000 8f20bebc 8f20bea0 8f20bea0 800a62b4 800a62b4 a0000013
fffffff
[<80594c10>] (__und_svc_finish) from [<800a62b4>]
(kgdb_breakpoint+0x58/0x94)
[<800a62b4>] (kgdb_breakpoint) from [<800a6368>]
(sysrq_handle_dbg+0x4c/0x70)
[<800a6368>] (sysrq_handle_dbg) from [<80349734>]
(__handle_sysrq+0x100/0x1f0)
[<80349734>] (__handle_sysrq) from [<80349c98>]
(write_sysrq_trigger+0x58/0x68)
[<80349c98>] (write_sysrq_trigger) from [<8019083c>]
(proc_reg_write+0x6c/0x94)
[<8019083c>] (proc_reg_write) from [<80139ab4>] (vfs_write+0xb4/0x1c0)
[<80139ab4>] (vfs_write) from [<8013a00c>] (SyS_write+0x4c/0xa0)
[<8013a00c>] (SyS_write) from [<8000f980>] (ret_fast_syscall+0x0/0x50)

```

```

[0]kdb> cpu      << kdb: cpu: <cpunum>           Switch to new cpu >>
Currently on cpu 0
Available cpus: 0

```

```

[0]kdb> lsmod
Module                               Size  modstruct  Used by

```

```
[0]kdb> dumpall      << kdb: dumpall: First line debugging (a kind of wrapper cmd) >>
```

```
[dumpall]kdb>      pid R
```

KDB current process is **sh(pid=609)**

```
[dumpall]kdb>      -dumpcommon      << commands shown like this =>
internally executed by
the wrapper cmd 'dumpall' >>
```

```
[dumpcommon]kdb>    set BTAPROMPT 0
```

```
[dumpcommon]kdb>    set LINES 10000
```

```
[dumpcommon]kdb>    -summary
```

```
sysname      Linux
release      3.18.22-kgdb
version      #2 SMP Thu Jan 19 12:46:36 IST 2017
machine      armv7l
nodename     (none)
domainname   (none)
ccversion    CCVERSION
date         2017-01-19 10:25:03 tz_minuteswest 0
uptime       00:12
load avg     0.00 0.00 0.00
```

```
MemTotal:    245252 kB
MemFree:     233244 kB
Buffers:     156 kB
```

```
[dumpcommon]kdb>    -cpu
```

Currently on cpu 0

Available cpus: 0

```
[dumpcommon]kdb>    -ps
```

30 sleeping system daemon (state M) processes suppressed,
use 'ps A' to see all.

Task Addr	Pid	Parent	[*]	cpu	State	Thread	Command
0x8f911e00	609	1	1	0	R	0x8f9120e0	*sh
0x8f880000	1	0	0	0	S	0x8f8802e0	init
0x8f911e00	609	1	1	0	R	0x8f9120e0	*sh

```
[dumpcommon]kdb>    -dmesg 600
```

buffer only contains 164 lines, first 164 lines printed

```
<6>[ 0.000000] Booting Linux on physical CPU 0x0
<6>[ 0.000000] Initializing cgroup subsys cpuset
<5>[ 0.000000] Linux version 3.18.22-kgdb (root@kaiwan-T460) (gcc
version 4.8.3 20140320 (prerelease) (Sourcery CodeBench Lite 2014.05-29) )
#2 SMP Thu Jan 19 12:46:36 IST 2017
<6>[ 0.000000] CPU: ARMv7 Processor [410fc090] revision 0 (ARMv7),
cr=10c5387d
```

```
<6>[ 0.000000] CPU: PIPT / VIPT nonaliasing data cache, VIPT nonaliasing
instruction cache
<4>[ 0.000000] Machine: ARM-Versatile Express
```

```
--snip--
```

```
<6>[ 2.148985] VFS: Mounted root (ext4 filesystem) readonly on device
179:0.
<6>[ 2.161095] Freeing unused kernel memory: 436K (807e1000 - 8084e000)
<5>[ 2.434529] random: nonblocking pool is initialized
<6>[ 3.174955] EXT4-fs (mmcblk0): re-mounted. Opts: data=ordered
<6>[ 3.241284] smsc911x smsc911x eth0: SMSC911x/921x identified at
0x90a40000, IRQ: 47
```

```
<6>[ 71.141098] kgdb: Registered I/O driver kgdboc.
```

```
<6>[ 83.832502] SysRq : DEBUG
```

```
[dumpcommon]kdb> -bt
```

```
Stack traceback for pid 609
```

```
0x8f911e00 609 1 1 0 R 0x8f9120e0 *sh
[<80017cdc>] (unwind_backtrace) from [<80013c5c>] (show_stack+0x20/0x24)
[<80013c5c>] (show_stack) from [<800b00bc>] (kdb_show_stack+0x54/0x68)
[<800b00bc>] (kdb_show_stack) from [<800b016c>] (kdb_bt1.isra.0+0x9c/0xe8)
[<800b016c>] (kdb_bt1.isra.0) from [<800b0508>] (kdb_bt+0x350/0x454)
[<800b0508>] (kdb_bt) from [<800ad99c>] (kdb_parse+0x2cc/0x6c4)
[<800ad99c>] (kdb_parse) from [<800ade3c>] (kdb_exec_defcmd+0xa8/0xf0)
[<800ade3c>] (kdb_exec_defcmd) from [<800ad99c>] (kdb_parse+0x2cc/0x6c4)
[<800ad99c>] (kdb_parse) from [<800ade3c>] (kdb_exec_defcmd+0xa8/0xf0)
[<800ade3c>] (kdb_exec_defcmd) from [<800ad99c>] (kdb_parse+0x2cc/0x6c4)
[<800ad99c>] (kdb_parse) from [<800ae57c>] (kdb_main_loop+0x530/0x790)
[<800ae57c>] (kdb_main_loop) from [<800b1174>] (kdb_stub+0x230/0x488)
[<800b1174>] (kdb_stub) from [<800a6e8c>] (kgdb_cpu_enter+0x43c/0x730)
[<800a6e8c>] (kgdb_cpu_enter) from [<800a7464>]
(kgdb_handle_exception+0x194/0x1e0)
[<800a7464>] (kgdb_handle_exception) from [<80017038>]
(kgdb_compiled_brk_fn+0x38/0x40)
[<80017038>] (kgdb_compiled_brk_fn) from [<800083fc>]
(do_undefinstr+0x17c/0x214)
[<800083fc>] (do_undefinstr) from [<80594c10>] (__und_svc_finish+0x0/0x30)
Exception stack(0x8f20be18 to 0x8f20be60)
be00: 00000067
00000000
be20: 90800000 90800730 80e36b40 80e36b44 80856448 8f20a010 00000000
00000008
be40: 00000000 8f20bebc 8f20bea0 8f20bea0 800a62b4 800a62b4 a0000013
ffffffff
[<80594c10>] (__und_svc_finish) from [<800a62b4>]
(kgdb_breakpoint+0x58/0x94)
[<800a62b4>] (kgdb_breakpoint) from [<800a6368>]
(sysrq_handle_dbg+0x4c/0x70)
[<800a6368>] (sysrq_handle_dbg) from [<80349734>]
(__handle_sysrq+0x100/0x1f0)
```

```
[<80349734>] (__handle_sysrq) from [<80349c98>]
(write_sysrq_trigger+0x58/0x68)
[<80349c98>] (write_sysrq_trigger) from [<8019083c>]
(proc_reg_write+0x6c/0x94)
[<8019083c>] (proc_reg_write) from [<80139ab4>] (vfs_write+0xb4/0x1c0)
[<80139ab4>] (vfs_write) from [<8013a00c>] (Sys_write+0x4c/0xa0)
[<8013a00c>] (Sys_write) from [<8000f980>] (ret_fast_syscall+0x0/0x50)
```

```
[dumpall]kdb> -bta
```

30 sleeping system daemon (state M) processes suppressed,
use 'ps A' to see all.

Stack traceback for pid 609

```
0x8f911e00      609      1 1      0   R  0x8f9120e0 *sh
[<80017cdc>] (unwind_backtrace) from [<80013c5c>] (show_stack+0x20/0x24)
[<80013c5c>] (show_stack) from [<800b00bc>] (kdb_show_stack+0x54/0x68)
[<800b00bc>] (kdb_show_stack) from [<800b016c>] (kdb_bt1.isra.0+0x9c/0xe8)
[<800b016c>] (kdb_bt1.isra.0) from [<800b0260>] (kdb_bt+0xa8/0x454)
[<800b0260>] (kdb_bt) from [<800ad99c>] (kdb_parse+0x2cc/0x6c4)
[<800ad99c>] (kdb_parse) from [<800ade3c>] (kdb_exec_defcmd+0xa8/0xf0)
[<800ade3c>] (kdb_exec_defcmd) from [<800ad99c>] (kdb_parse+0x2cc/0x6c4)
[<800ad99c>] (kdb_parse) from [<800ae57c>] (kdb_main_loop+0x530/0x790)
[<800ae57c>] (kdb_main_loop) from [<800b1174>] (kdb_stub+0x230/0x488)
[<800b1174>] (kdb_stub) from [<800a6e8c>] (kgdb_cpu_enter+0x43c/0x730)
[<800a6e8c>] (kgdb_cpu_enter) from [<800a7464>]
(kgdb_handle_exception+0x194/0x1e0)
[<800a7464>] (kgdb_handle_exception) from [<80017038>]
(kgdb_compiled_brk_fn+0x38/0x40)
[<80017038>] (kgdb_compiled_brk_fn) from [<800083fc>]
(do_undefinstr+0x17c/0x214)
[<800083fc>] (do_undefinstr) from [<80594c10>] (__und_svc_finish+0x0/0x30)
Exception stack(0x8f20be18 to 0x8f20be60)
```

```
be00:                                00000067
00000000
be20: 90800000 90800730 80e36b40 80e36b44 80856448 8f20a010 00000000
00000008
be40: 00000000 8f20bebc 8f20bea0 8f20bea0 800a62b4 800a62b4 a0000013
ffffffff
```

```
[<80594c10>] (__und_svc_finish) from [<800a62b4>]
(kgdb_breakpoint+0x58/0x94)
[<800a62b4>] (kgdb_breakpoint) from [<800a6368>]
(sysrq_handle_dbg+0x4c/0x70)
[<800a6368>] (sysrq_handle_dbg) from [<80349734>]
(__handle_sysrq+0x100/0x1f0)
[<80349734>] (__handle_sysrq) from [<80349c98>]
(write_sysrq_trigger+0x58/0x68)
[<80349c98>] (write_sysrq_trigger) from [<8019083c>]
(proc_reg_write+0x6c/0x94)
[<8019083c>] (proc_reg_write) from [<80139ab4>] (vfs_write+0xb4/0x1c0)
[<80139ab4>] (vfs_write) from [<8013a00c>] (Sys_write+0x4c/0xa0)
[<8013a00c>] (Sys_write) from [<8000f980>] (ret_fast_syscall+0x0/0x50)
```

Stack traceback for pid 1

```
0x8f880000      1      0 0      0   S  0x8f8802e0 init
```

```
[<8058e5b0>] (__schedule) from [<8058e9c4>] (schedule+0x40/0x8c)
[<8058e9c4>] (schedule) from [<8002b340>] (do_wait+0x220/0x28c)
[<8002b340>] (do_wait) from [<8002b7b4>] (Sys_wait4+0x74/0xd8)
[<8002b7b4>] (Sys_wait4) from [<8000f980>] (ret_fast_syscall+0x0/0x50)
[0]kdb>
[0]kdb> go    << kdb: go: [<vaddr>]          Continue Execution >>
ARM / #
```

When kdb is enabled, a kernel Oops causes the debugger to be invoked, and the keyboard LEDs blink. Once the kdb prompt is displayed, you can enter kdb commands.

If the system will be run in graphical mode with kdb enabled, it is recommend that kdb be set up to use a serial console so that the kdb prompt can be seen.

kdb Activation

kdb can be activated by configuring it at kernel build time. If kdb off by default (CONFIG_KDB_OFF) was not selected during kernel configuration, kdb is active by default.

[DEPRECATED : below]

The runtime options are as follows:

- Activate it by passing the `kdb=on` flag to the kernel during boot.

Turn it on or off through the `/proc` file system entry

- To enable kdb through the `/proc`, use this command:

```
# echo "1" > /proc/sys/kernel/kdb
```

To disable kdb through the `/proc`, use this command:

```
# echo "0" > /proc/sys/kernel/kdb
```

If kdb is needed during boot, specify the `kdb=early` flag to the kernel during boot. This allows kdb to be activated during the boot process.

kdb is invoked in the following ways:

- Pressing the Pause key on the keyboard manually invokes kdb.
- Whenever there is an Oops or kernel panic.
- If kdb is set up through a serial port between two machines, pressing Ctrl-A invokes kdb from the serial console. Several programs can be used to communicate with kdb through a serial port, such as minicom and kermit.

Note-The key sequence Ctrl-A has been changed in version 4.4 of kdb. It is now Esc-KDB.

- Also see: [Entering kdb from qemu](#)

<<

FYI, an interesting blog post (from your author):

[“A KDB / KGDB session on the popular Raspberry Pi embedded Linux board”](#)

>>

Another Sample Session

We'll use the same Oops-generating kernel module – *hello_oops.ko* – that we used earlier.

Important:

When trying this out, first **switch to console mode** (on a PC, Ctrl-Alt-F1 should do it); as kdb will show up only here of course (not in the GUI X Windows mode).

<< On console >>

```
# uname -r
2.6.22-kdb
# insmod hello_oops.ko
# dmesg
--snip--
[4295047.638000] hello_oops 1.0 initialize
# cat /proc/oops-test
...
[ 3739.122922] *pte = 00000000
[ 3739.122975] Oops: 0000 [#9]
[ 3739.123023] PREEMPT
[ 3739.123109] Modules linked in: ipv6 autofs4 hello_oops nls_ascii vfat
fat
sd_mod usb_storage scsi_mod dm_mod button battery ac uhci_hcd shpchp
i2c_i801
i2c_core floppy
[ 3739.123889] CPU:      0
[ 3739.123891] EIP:      0060:[<d081e031>]      Not tainted VLI
[ 3739.123893] EFLAGS: 00010286   (2.6.22-kdb #1)
[ 3739.124046] EIP is at procread+0x31/0x5f [hello_oops]
[ 3739.124100] eax: 00000039   ebx: 00000039   ecx: 00000000   edx:
d081e0de
[ 3739.124156] esi: 00001000   edi: c9c52000   ebp: c9d23f30   esp:
c9d23f1c
[ 3739.124212] ds: 007b   es: 007b   fs: 0000   gs: 0033   ss: 0068
[ 3739.124267] Process cat (pid: 5807, ti=c9d23000 task=ca8e0b00
task.ti=c9d23000)
[ 3739.124322] Stack:c9c52000 d081e0a7 000016af 00000000 cd261bc0 c9d23f70
c0183b88 00000c00
```

```

[ 3739.124717] c9d23f5c 00000000 00000c00 cd261bc0 00000000 ce29ca84
00001000 0804d858
[ 3739.125112] 00000000 00000000 ca4539c0 0804d858 c030b1a0 c9d23f90
c01561f2 c9d23f9c
[ 3739.125506] Call Trace:
[ 3739.125598] [<c01040f5>] show_trace_log_lvl+0x19/0x2e
[ 3739.125698] [<c01041b7>] show_stack_log_lvl+0x99/0xa1
[ 3739.125789] [<c01043b9>] show_registers+0x1b8/0x2cb
[ 3739.124212] ds: 007b es: 007b fs: 0000 gs: 0033 ss: 0068
[ 3739.124267] Process cat (pid: 5807, ti=c9d23000 task=ca8e0b00
task.ti=c9d23000)
[ 3739.124322] Stack:c9c52000 d081e0a7 000016af 00000000 cd261bc0 c9d23f70
c0183b88 00000c00
[ 3739.124717] c9d23f5c 00000000 00000c00 cd261bc0 00000000 ce29ca84
00001000 0804d858
[ 3739.125112] 00000000 00000000 ca4539c0 0804d858 c030b1a0 c9d23f90
c01561f2 c9d23f9c
[ 3739.125506] Call Trace:
[ 3739.125598] [<c01040f5>] show_trace_log_lvl+0x19/0x2e
[ 3739.125698] [<c01041b7>] show_stack_log_lvl+0x99/0xa1
[ 3739.125789] [<c01043b9>] show_registers+0x1b8/0x2cb
[ 3739.125879] [<c0104628>] die+0x100/0x1dc
[ 3739.125968] [<c01142be>] do_page_fault+0x44d/0x528
[ 3739.126059] [<c02fdc7a>] error_code+0x6a/0x70
[ 3739.126152] [<c0183b88>] proc_file_read+0x116/0x229
[ 3739.126247] [<c01561f2>] vfs_read+0x8a/0x109
[ 3739.126341] [<c01564b3>] sys_read+0x3d/0x61
[ 3739.126431] [<c0103c86>] sysenter_past_esp+0x5f/0x85
[ 3739.126522] =====
[ 3739.126571] Code: 10 c7 44 24 0c 00 00 00 00 8b 15 00 40 3d c0 8b 92 a4
00
00 00 c7 44 24 04 a7 e0 81 d0 89 04 24 89 54 24 08 e8 49 71 9b ef 89 c3
<a1>
00 00 00 00 c7 44 24 04 e0 e0 81 d0 c7 04 24 eb e0 81 d0 89
[ 3739.129038] EIP: [<d081e031>] procread+0x31/0x5f [hello_oops] SS:ESP
0068:c9d

```

<< The system traps into kdb now >>

Entering kdb (current=0xca8e0b00, pid 5807) Oops: Oops

due to oops @ 0xd081e031

```

eax = 0x00000039 ebx = 0x00000039 ecx = 0x00000000 edx = 0xd081e0de
esi = 0x00001000 edi = 0xc9c52000 esp = 0xc9d23f1c eip = 0xd081e031
ebp = 0xc9d23f30 xss = 0xc0290068 xcs = 0x00000060 eflags = 0x00010286
xds = 0x0000007b xes = 0xc9c5007b origeax = 0xffffffff &regs = 0xc9d23ee4
kdb> ps

```

22 sleeping system daemon (state M) processes suppressed

Task Addr	Pid	Parent	[*]	cpu	State	Thread	Command
0xca8e0b00	5807	5806	1	0	R	0xca8e0cb0	*cat
0xc1241450	1	0	0	0	S	0xc1241600	init


```

0xcdfd0070      1210      1  0  0  S  0xcdfd0220  udevd
0xcff70030      3472      1  0  0  R  0xcff701e0  klogd
0xcff0b490      3501      1  0  0  S  0xcff0b640  portmap
0xcff48070      3521      1  0  0  S  0xcff48220  rpc.statd
0xcff6e170      4370      1  0  0  S  0xcff6e320  apmd
0xc83be0b0      4460      1  0  0  S  0xc83be260  smartd
0xc7c7cb80      4470      1  0  0  S  0xc7c7cd30  acpid

```

...

...

more> q

kdb> bt

Stack traceback for pid 5807

```

0xca8e0b00      5807      5806  1  0  R  0xca8e0cb0  *cat

```

```

esp      eip      Function (args)

```

kdb_bb: address 0xffffffff not recognised

Using old style backtrace, unreliable with no arguments

```

esp      eip      Function (args)

```

```

0xc9d23f10 0xd081e031 [hello_oops]procread+0x31      <--- function+offset
causing the Oops

```

```

0xc9d23f34 0xc0183b88 proc_file_read+0x116

```

```

0xc9d23f74 0xc01561f2 vfs_read+0x8a

```

kdb> id procread <--- **[DEPRECATED] disassemble the procread function**

```

0xd081e000 procread:      push    %ebp

```

```

0xd081e001 procread+0x1:  mov     %esp,%ebp

```

```

0xd081e003 procread+0x3:  push    %ebx

```

...

```

0xd081e02a procread+0x2a:  call    0xc01d5178 sprintf

```

```

0xd081e02f procread+0x2f:  mov     %eax,%ebx

```

```

0xd081e031 procread+0x31:  mov     0x0,%eax

```

```

0xd081e036 procread+0x36:  movl    $0xd081e0e0,0x4(%esp)

```

```

0xd081e03e procread+0x3e:  movl    $0xd081e0eb, (%esp)

```

```

0xd081e045 procread+0x45:  mov     %eax,0x8(%esp)

```

kdb> id procread+0x45

```

0xd081e045 procread+0x45:  mov     %eax,0x8(%esp)

```

```

0xd081e049 procread+0x49:  call    0xc0119a43 printk

```

```

0xd081e04e procread+0x4e:  mov     0xc(%ebp),%eax

```

```

0xd081e051 procread+0x51:  movl    $0x1, (%eax)

```

```

0xd081e057 procread+0x57:  add     $0x10,%esp

```

```

0xd081e05a procread+0x5a:  mov     %ebx,%eax

```

```

0xd081e05c procread+0x5c:  pop     %ebx

```

```

0xd081e05d procread+0x5d:  pop     %ebp

```

```

0xd081e05e procread+0x5e:  ret

```

```

0xd081e05f hello_oops_cleanup_module:  push    %ebp

```

```

0xd081e060 hello_oops_cleanup_module+0x1:  xor     %edx,%edx

```

```

0xd081e062 hello_oops_cleanup_module+0x3:  mov     %esp,%ebp

```

```

0xd081e064 hello_oops_cleanup_module+0x5:  mov     $0xd081e0f6,%eax

```

```

0xd081e069 hello_oops_cleanup_module+0xa:  sub     $0xc,%esp

```

```

0xd081e06c hello_oops_cleanup_module+0xd:  call    0xc018453e

```

```

remove_proc_entry

```



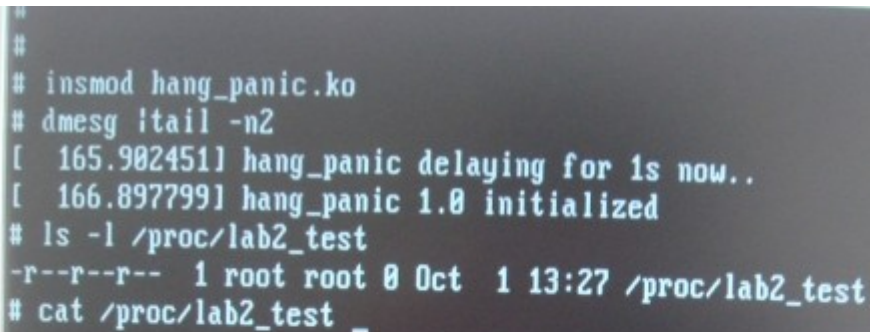
```
0xd081e071 hello_oops_cleanup_module+0x12:    movl    $0xd081e136,0x8(%esp)
kdb> go
```

```
Catastrophic error detected
kdb_continue_catastrophic=0, type go a second time if you really want to
continue
kdb> go
```

```
Catastrophic error detected
kdb_continue_catastrophic=0, attempting to continue
Segmentation fault
#
```

<<

*Instructor Note: photos taken of the screen showing the **hang_panic.ko** kernel module debug steps with kdb >>*



```
#
#
# insmod hang_panic.ko
# dmesg |tail -n2
[ 165.982451] hang_panic delaying for 1s now..
[ 166.897799] hang_panic 1.0 initialized
# ls -l /proc/lab2_test
-r--r--r-- 1 root root 0 Oct  1 13:27 /proc/lab2_test
# cat /proc/lab2_test _
```

1. With (buggy) kernel module hang_panic.ko : About to trigger the Oops by reading the proc entry. Of course, we're running on the console. (Photo above).

```
[ 332.945412] c012705c 00000011 c045b000 0000000a c0436ff8 c011d503 c04f
3f30 c0403000
[ 332.945000] 00000046 c0403f4c c0105004
[ 332.946005] Call Trace:
[ 332.946097] [] show_trace_log_lvl+0x19/0x2e
[ 332.946192] [] show_stack_log_lvl+0x99/0xa1
[ 332.946202] [] show_registers+0x1b8/0x2cb
[ 332.946372] [] die+0x100/0x1dc
[ 332.946461] [] do_page_fault+0x44d/0x520
[ 332.946552] [] error_code+0x6a/0x70
[ 332.946644] [] __do_softirq+0x38/0x78
[ 332.946734] [] do_softirq+0x44/0xab
[ 332.946824] =====
[ 332.946873] Code: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00 00 00 00 00 00 00 00 00 00 00 ed 1e af de ff ff ff ff ff ff ff ff <18> 0a 4f
c0 6c 0d 4a c0 16 81 38 c0 00 00 00 00 00 00 00 00 00 00
[ 332.949335] EIP: [] tasklist_lock+0x18/0x00 SS:ESP 0068:c0436fb4

Entering kdb (current=0xc03d0bc0, pid 0) Oops: Oops
due to oops @ 0xc0400a10
eax = 0xce785f30 ebx = 0xce785f30 ecx = 0x00000001 edx = 0x00000000
esi = 0xc045bb00 edi = 0x00000100 esp = 0xc0436fb4 eip = 0xc0400a10
ebp = 0xc0436fe4 xss = 0xc0290068 xcs = 0x00000060 eflags = 0x00010206
xds = 0x0000007b xes = 0xc045007b origeax = 0xffffffff &regs = 0xc0436f7c
kdb> _
```

2. [Enter] pressed, the Oops occurs (sure enough); and, there, kdb kicks in!

```

<6>[ 332.940593] hang_panic setting up timer (1s) now..
<1>[ 332.943809] BUG: unable to handle kernel NULL pointer dereference at virt
al address 00000000
<1>[ 332.943900] printing eip:
<4>[ 332.943942] c0400a18
<1>[ 332.943945] *pde = 0dbd2067
<1>[ 332.943986] *pte = 00000000
<0>[ 332.944031] Oops: 0002 [#1]
<0>[ 332.944072] PREEMPT
<4>[ 332.944150] Modules linked in: hang_panic dm_mod button battery ac uhci_hc
d shpchp i2c_i801 i2c_core floppy
<0>[ 332.944608] CPU: 0
<0>[ 332.944609] EIP: 0060:[c0400a18] Not tainted VLI
<0>[ 332.944611] EFLAGS: 00010286 (2.6.22-kdb #1)
<0>[ 332.944748] EIP is at tasklist_lock+0x18/0x80
<0>[ 332.944794] eax: ce785f30 ebx: ce785f30 ecx: 00000001 edx: 00000000
<0>[ 332.944851] esi: c045bb00 edi: 00000100 ebp: c0436fe4 esp: c0436fb4
<0>[ 332.944907] ds: 007b es: 007b fs: 0000 gs: 0000 ss: 0068
<0>[ 332.944962] Process swapper (pid: 0, ti=c0436000 task=c03d0bc0 task.ti=c04
03000)
<0>[ 332.945017] Stack: c01203f2 c03d07e0 c0436fd0 c0127029 00000000 c0400a18 c
0436fcc c0436fcc
<0>[ 332.945412] c012705c 00000011 c045b000 0000000a c0436ff8 c011d503 c
0403f30 c0403000
more>

```

3. Output of kdb's 'dmesg' command (1 of 2)


```
<0>[ 332.944907] ds: 007b es: 007b fs: 0000 gs: 0000 ss: 0068  
<0>[ 332.944962] Process swapper (pid: 0, ti=c0436000 task=c03d0bc0 task.ti=c0  
03000)  
<0>[ 332.945017] Stack: c01203f2 c03d87e0 c0436fd0 c0127029 00000000 c0400a18  
0436fcc c0436fcc  
<0>[ 332.945412] c012705c 00000011 c045b888 0000000a c0436ff8 c011d503  
0403f30 c0403000  
more>  
  
<0>[ 332.945888] 00000046 c0403f4c c0105804  
<0>[ 332.946005] Call Trace:  
<0>[ 332.946097] [] show_trace_log_lvl+0x19/0x2e  
<0>[ 332.946192] [] show_stack_log_lvl+0x99/0xa1  
<0>[ 332.946282] [] show_registers+0x1b8/0x2cb  
<0>[ 332.946372] [] die+0x100/0x1dc  
<0>[ 332.946461] [] do_page_fault+0x44d/0x528  
<0>[ 332.946552] [] error_code+0x6a/0x70  
<0>[ 332.946644] [] __do_softirq+0x38/0x78  
<0>[ 332.946734] [] do_softirq+0x44/0xab  
<0>[ 332.946824] =====  
<0>[ 332.946873] Code: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00  
00 00 00 00 00 00 00 00 00 00 ed 1e af de ff ff ff ff ff ff ff ff <18> 0e  
48 c0 6c 0d 4a c0 16 81 38 c0 00 00 00 00 00 00 00 00 00 00  
<0>[ 332.949335] EIP: [] tasklist_lock+0x18/0x80 SS:ESP 0068:c0436fbf  
kdb>
```

4. (Continued) Output of kdb's 'dmesg' command (2 of 2)

```

[ 332.949335] EIP: [c0400a18] tasklist_lock+0x18/0x80 SS:ESP 0068:c0
Entering kdb (current=0xc03d0bc0, pid 0) Oops: Oops
due to oops @ 0xc0400a18
eax = 0xce785f30 ebx = 0xce785f30 ecx = 0x00000001 edx = 0x00000000
esi = 0xc045bb00 edi = 0x00000100 esp = 0xc0436fb4 eip = 0xc0400a18
ebp = 0xc0436fe4 xss = 0xc0290060 xcs = 0x00000060 eflags = 0x00010286
xds = 0x0000007b xes = 0xc045007b origeax = 0xffffffff &regs = 0xc0436
kdb> ps

22 sleeping system daemon (state M) processes suppressed
Task Addr      Pid  Parent [*] cpu State Thread      Command
0xc03d0bc0      0      0  1  0  R  0xc03d0d70 *swapper

0xc1241450      1      0  0  0  S  0xc1241600  init
0xce071490     704      1  0  0  S  0xce071640  minilogd
0xc1264b00    1200      1  0  0  S  0xc1264cb0  udevd
0xcd8580f0    2134      1  0  0  S  0xcd8582a0  init
0xcd816ac0    2139    2134  0  0  S  0xcd816c70  sh
kdb>

```

5. Output of kdb's 'ps' command; the process marked with the asterisk '*' was the one running when the Oops occurred.

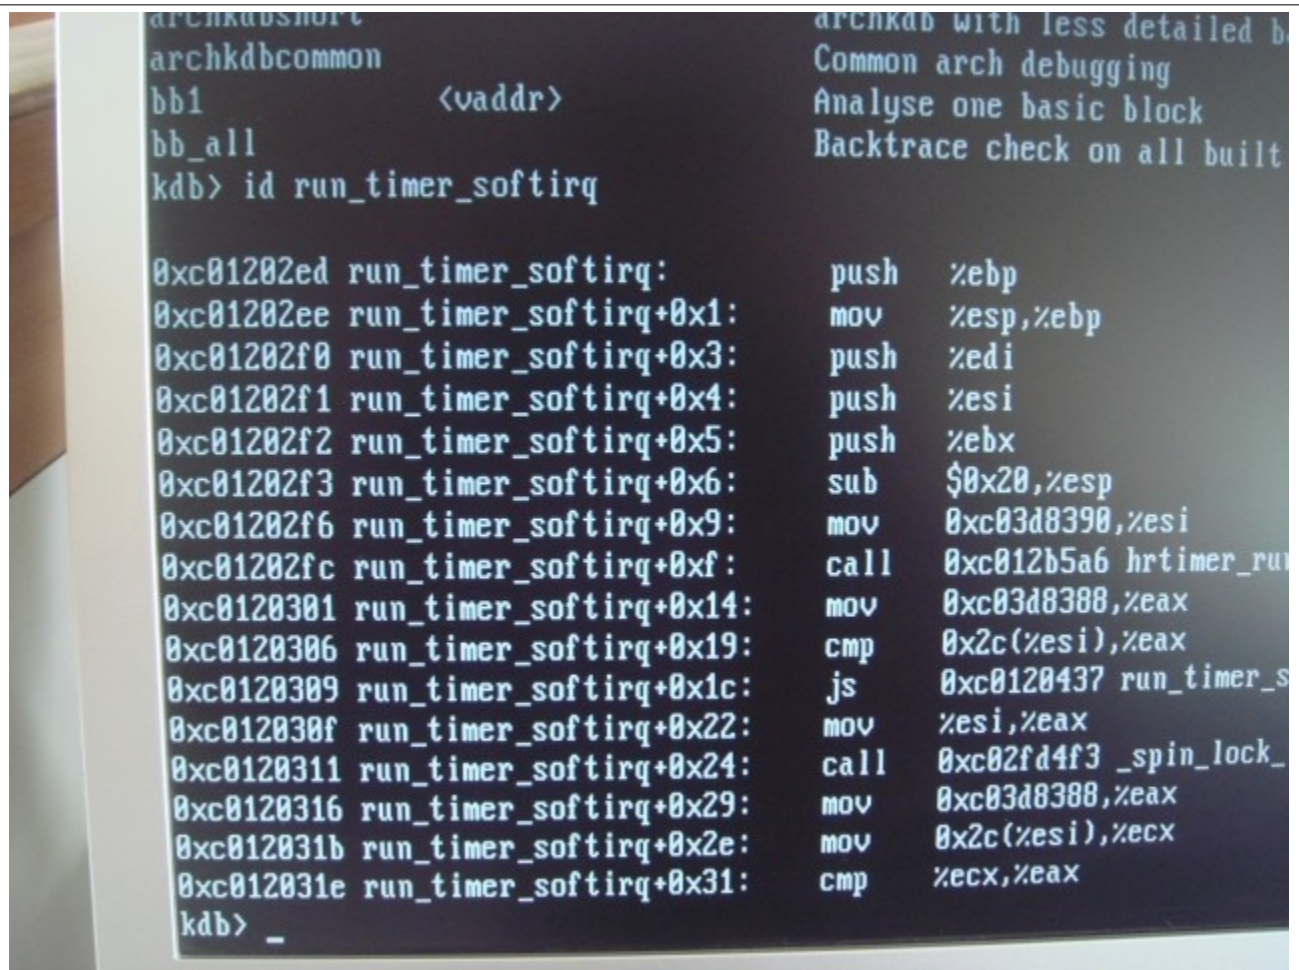
```

<0>[ 332.949335] EIP: [<c0400a18>] tasklist_lock+0x18/0x80 SS:ESP 0068:c0436fb4
kdb> bt

Stack traceback for pid 0
0xc03d0bc0      0      0 1 0 R 0xc03d0d70 *swapper
esp      eip      Function (args)
===== <softirq_ctx>
kdb_bb: address 0xffffffff not recognised
Using old style backtrace, unreliable with no arguments
esp      eip      Function (args)
===== <softirq_ctx>
0xc0436fb4 0xc01203f2 run_timer_softirq+0x105
0xc0436fc0 0xc0127029 __rcu_process_callbacks+0x8c
0xc0436fd4 0xc012705c rcu_process_callbacks+0x21
0xc0436fe8 0xc011d503 __do_softirq+0x38
0xc0436ffc 0xc0105804 do_softirq+0x44
===== <normal>
0xc0403f30 0xc0105804 do_softirq+0x44
0xc0403f48 0xc0139e81 handle_level_irq
0xc0403f50 0xc011d582 irq_exit+0x2c
0xc0403f58 0xc0105703 do_IRQ+0xb8
0xc0403f80 0xc0103e97 common_interrupt+0x23
0xc0403fa0 0xc011007b acpi_copy_wakeup_routine+0x1b
0xc0403fb0 0xc010204e default_idle+0x2a
more> _

```

6. Output of kdb's 'bt' command; note that the last recognizable function that ran is 'run_timer_softirq' in (soft)irq context. This clearly tells us that some timer code was running when the crash happened. Also, note that the code offset is 0x105 bytes.



```

archkdb with less detailed b
Common arch debugging
Analyse one basic block
Backtrace check on all built

kdb> id run_timer_softirq

0xc01202ed run_timer_softirq:      push    %ebp
0xc01202ee run_timer_softirq+0x1:  mov     %esp,%ebp
0xc01202f0 run_timer_softirq+0x3:  push    %edi
0xc01202f1 run_timer_softirq+0x4:  push    %esi
0xc01202f2 run_timer_softirq+0x5:  push    %ebx
0xc01202f3 run_timer_softirq+0x6:  sub     $0x20,%esp
0xc01202f6 run_timer_softirq+0x9:  mov     0xc03d8390,%esi
0xc01202fc run_timer_softirq+0xf:  call    0xc012b5a6 hrtimer_run
0xc0120301 run_timer_softirq+0x14: mov     0xc03d8388,%eax
0xc0120306 run_timer_softirq+0x19: cmp     0x2c(%esi),%eax
0xc0120309 run_timer_softirq+0x1c: js      0xc0120437 run_timer_s
0xc012030f run_timer_softirq+0x22: mov     %esi,%eax
0xc0120311 run_timer_softirq+0x24: call    0xc02fd4f3 _spin_lock_
0xc0120316 run_timer_softirq+0x29: mov     0xc03d8388,%eax
0xc012031b run_timer_softirq+0x2e: mov     0x2c(%esi),%ecx
0xc012031e run_timer_softirq+0x31: cmp     %ecx,%eax
kdb> _

```

7. Investigate further by viewing the output of kdb's 'id' (disassemble) command on the relevant function (1 of 2).

```

0xc01203e5 run_timer_softirq+0xf8:  and    $0xffffffff000,%eax
0xc01203ea run_timer_softirq+0xfd:  mov     0x14(%eax),%edi
0xc01203ed run_timer_softirq+0x100:  mov     %ebx,%eax
0xc01203ef run_timer_softirq+0x102:  call    *0xffffffffe4(%ebp)
0xc01203f2 run_timer_softirq+0x105:  mov     %esp,%eax
kdb>

0xc01203f4 run_timer_softirq+0x107:  and     $0xffffffff000,%eax
0xc01203f9 run_timer_softirq+0x10c:  mov     0x14(%eax),%eax
0xc01203fc run_timer_softirq+0x10f:  cmp     %eax,%edi
0xc01203fe run_timer_softirq+0x111:  je      0xc012041f run_timer_softirq+0x11f
0xc0120400 run_timer_softirq+0x113:  mov     %eax,0xc(%esp)
0xc0120404 run_timer_softirq+0x117:  mov     0xffffffffe4(%ebp),%eax
0xc0120407 run_timer_softirq+0x11a:  mov     %edi,0x8(%esp)
0xc012040b run_timer_softirq+0x11e:  movl    $0xc0388a5b,(%esp)
0xc0120412 run_timer_softirq+0x125:  mov     %eax,0x4(%esp)
0xc0120416 run_timer_softirq+0x129:  call    0xc0119a43 printk
0xc012041b run_timer_softirq+0x12e:  ud2a
0xc012041d run_timer_softirq+0x130:  jmp     0xc012041d run_timer_softirq+0x130
0xc012041f run_timer_softirq+0x132:  mov     %esi,%eax
0xc0120421 run_timer_softirq+0x134:  lea     0xffffffffe8(%ebp),%ebx
0xc0120424 run_timer_softirq+0x137:  call    0xc02fd4f3 _spin_lock_irq
0xc0120429 run_timer_softirq+0x13c:  mov     0xffffffffe8(%ebp),%ecx
kdb> _

```

8. Investigate further by viewing the output of kdb's 'id' (disassemble) command on the relevant function; this time we can see offset 0x105 (2 of 2).

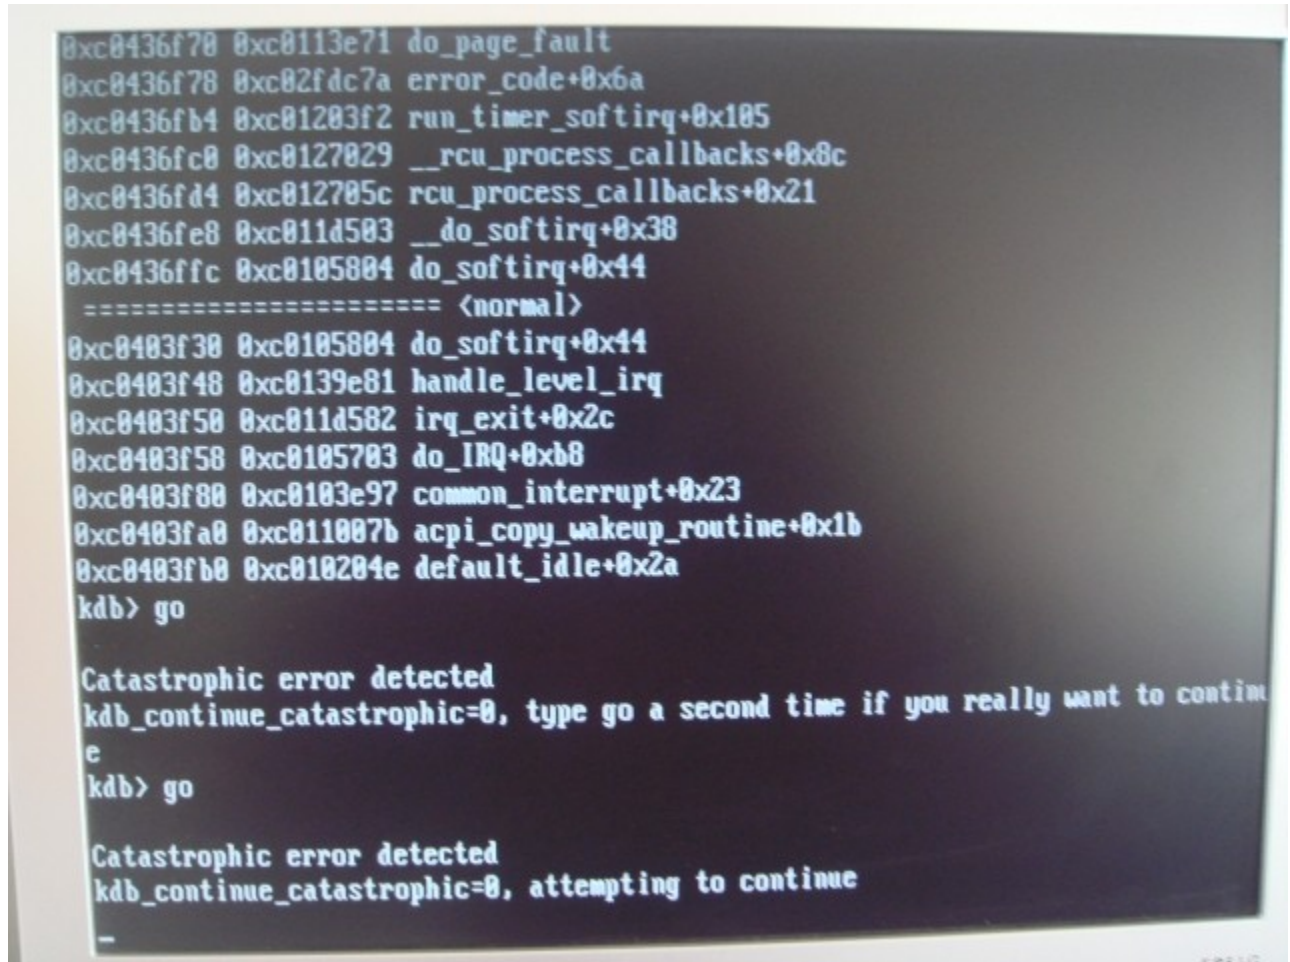
Of course, this does not necessarily mean that *this* code is buggy: rather it should point us to carefully check how *we invoked* the timer. Careful study will point us to the bug, viz., the timer data structure was declared *locally* and is therefore out of scope (yielding some arbitrary junk values) when the timer fires and the kernel code attempts to access the timer_struct contents.


```
ssb                                     Single step to branch/call
pt_regs                               Format struct pt_regs
more> q
kdb> lsmod

Module                               Size  modstruct    Used by
hang_panic                           2560  0xd081e780    0 (Live) [ ]
dm_mod                               48616  0xd0863700    0 (Live) [ ]
button                               6672  0xd0848780    0 (Live) [ ]
battery                              8964  0xd0845080    0 (Live) [ ]
ac                                    4356  0xd0840e80    0 (Live) [ ]
uhci_hcd                             21788  0xd083e280    0 (Live) [ ]
shpchp                               27960  0xd0837a80    0 (Live) [ ]
i2c_i801                             8080  0xd0805d00    0 (Live) [ ]
i2c_core                             19984  0xd0856b80    1 (Live) [ i2c_i801 ]
floppy                               54212  0xd087bb80    0 (Live) [ ]
kdb> go

Catastrophic error detected
kdb_continue_catastrophic=0, type go a second time if you really want to continue
kdb> go_
```

9. Output of kdb's 'lsmod' and 'go' commands (1 of 2).



```
0xc0436f70 0xc0113e71 do_page_fault
0xc0436f78 0xc02fdc7a error_code+0x6a
0xc0436fb4 0xc01203f2 run_timer_softirq+0x105
0xc0436fc0 0xc0127029 __rcu_process_callbacks+0x8c
0xc0436fd4 0xc012705c rcu_process_callbacks+0x21
0xc0436fe8 0xc011d503 __do_softirq+0x38
0xc0436ffc 0xc0105804 do_softirq+0x44
===== <normal>
0xc0403f30 0xc0105804 do_softirq+0x44
0xc0403f48 0xc0139e81 handle_level_irq
0xc0403f50 0xc011d582 irq_exit+0x2c
0xc0403f58 0xc0105703 do_IRQ+0xb8
0xc0403f80 0xc0103e97 common_interrupt+0x23
0xc0403fa0 0xc011007b acpi_copy_wakeup_routine+0x1b
0xc0403fb0 0xc010204e default_idle+0x2a
kdb> go

Catastrophic error detected
kdb_continue_catastrophic=0, type go a second time if you really want to continue
kdb> go

Catastrophic error detected
kdb_continue_catastrophic=0, attempting to continue
_
```

10. Output of kdb's second 'go' command (2 of 2).

[FYI / OPTIONAL]**Useful Tips****Resource**[KernelDebuggingTricks](#)

...

gdb on vmlinux

One can disassemble a built kernel using gdb on the vmlinux image. This is useful when one gets a kernel Oops message and a stack dump - one can then disassemble the object code and see where the Oops is occurring. For example:

gdb debian/build/build-generic/vmlinux**(gdb) disassemble printk**

Dump of assembler code for function printk:

```

0xffffffff8023dce0 <printk+0>:  sub    $0xd8,%rsp
0xffffffff8023dce7 <printk+7>:  lea    0xe0(%rsp),%rax
0xffffffff8023dcef <printk+15>: mov    %rsi,0x28(%rsp)
0xffffffff8023dcf4 <printk+20>: mov    %rsp,%rsi
0xffffffff8023dcf7 <printk+23>: mov    %rdx,0x30(%rsp)
0xffffffff8023dcfc <printk+28>: mov    %rcx,0x38(%rsp)
0xffffffff8023dd01 <printk+33>: mov    %rax,0x8(%rsp)
0xffffffff8023dd06 <printk+38>: lea    0x20(%rsp),%rax
0xffffffff8023dd0b <printk+43>: mov    %r8,0x40(%rsp)
0xffffffff8023dd10 <printk+48>: mov    %r9,0x48(%rsp)
0xffffffff8023dd15 <printk+53>: movl   $0x8, (%rsp)
0xffffffff8023dd1c <printk+60>: movl   $0x30,0x4(%rsp)
0xffffffff8023dd24 <printk+68>: mov    %rax,0x10(%rsp)
0xffffffff8023dd29 <printk+73>: callq  0xffffffff8023d980 <vprintk>
0xffffffff8023dd2e <printk+78>: add    $0xd8,%rsp
0xffffffff8023dd35 <printk+85>: retq
End of assembler dump.

```

Objdump

If one has the built object code at hand, one can disassemble the object using objdump as follows:

```
objdump -SdCg debian/build/build-generic/fs/dcache.o
```

Using GDB to find the location where your kernel panicked or oopsed.

A quick and easy way to find the line of code where your kernel panicked or oopsed is to use GDB list command. You can do this as follows.

Lets assume your panic/oops message says something like:

```
[ 174.507084] Stack:
```

```

[ 174.507163] ce0bd8ac 00000008 00000000 ce4a7e90 c039ce30 ce0bd8ac c0718b04
c07185a0
[ 174.507380] ce4a7ea0 c0398f22 ce0bd8ac c0718b04 ce4a7eb0 c037deee ce0bd8e0
ce0bd8ac
[ 174.507597] ce4a7ec0 c037dfe0 c07185a0 ce0bd8ac ce4a7ed4 c037d353 ce0bd8ac
ce0bd8ac
[ 174.507888] Call Trace:
[ 174.508125] [<c039ce30>] ? sd_remove+0x20/0x70
[ 174.508235] [<c0398f22>] ? scsi_bus_remove+0x32/0x40
[ 174.508326] [<c037deee>] ? __device_release_driver+0x3e/0x70
[ 174.508421] [<c037dfe0>] ? device_release_driver+0x20/0x40
[ 174.508514] [<c037d353>] ? bus_remove_device+0x73/0x90
[ 174.508606] [<c037bccf>] ? device_del+0xef/0x150
[ 174.508693] [<c0399207>] ? __scsi_remove_device+0x47/0x80
[ 174.508786] [<c0399262>] ? scsi_remove_device+0x22/0x40
[ 174.508877] [<c0399324>] ? __scsi_remove_target+0x94/0xd0
[ 174.508969] [<c03993c0>] ? __remove_child+0x0/0x20
[ 174.509060] [<c03993d7>] ? __remove_child+0x17/0x20
[ 174.509148] [<c037b868>] ? device_for_each_child+0x38/0x60
[ 174.509241] [<c039938f>] ? scsi_remove_target+0x2f/0x60
[ 174.509393] [<d0c38907>] ? __iscsi_unbind_session+0x77/0xa0
[scsi_transport_iscsi]
[ 174.509699] [<c015272e>] ? run_workqueue+0x6e/0x140
[ 174.509801] [<d0c38890>] ? __iscsi_unbind_session+0x0/0xa0
[scsi_transport_iscsi]
[ 174.509977] [<c0152888>] ? worker_thread+0x88/0xe0
[ 174.510047] [<c01566a0>] ? autoremove_wake_function+0x0/0x40

```

Lets say you want to know what line of code represents `sd_remove+0x20/0x70`. cd to the ubuntu debian/build/build-generic directory in your kernel tree and run gdb on the ".o" file which has the function `sd_remove()` in this case in `sd.o`, and use the gdb "list" command, (gdb) list *(function+0xoffset), in this case function is `sd_remove()` and offset is 0x20, and gdb should tell you the line number where you hit the panic or oops. This has worked for me very reliably for most cases.

manjo@hungry:~/devel/ubuntu/kernel/ubuntu-karmic-397906/debian/build/build-generic/drivers/scsi\$ gdb sd.o

GNU gdb (GDB) 6.8.50.20090628-cvs-debian

Copyright (C) 2009 Free Software Foundation, Inc.

License GPLv3+: GNU GPL version 3 or later

<<http://gnu.org/licenses/gpl.html>>

This is free software: you are free to change and redistribute it.

There is NO WARRANTY, to the extent permitted by law. Type "show copying" and "show warranty" for details.

This GDB was configured as "x86_64-linux-gnu".

For bug reporting instructions, please see:

<<http://www.gnu.org/software/gdb/bugs/>>...

(gdb) list *(sd_remove+0x20)

0x1650 is in sd_remove (/home/manjo/devel/ubuntu/kernel/ubuntu-karmic-397906/drivers/scsi/sd.c:2125).

```
2120     static int sd_remove(struct device *dev)
```

```
2121     {
```

```
2122         struct scsi_disk *sdkp;
2123
2124         async_synchronize_full();
2125         sdkp = dev_get_drvdata(dev);
2126         blk_queue_prep_rq(sdkp->device->request_queue,
scsi_prep_fn);
2127         device_del(&sdkp->dev);
2128         del_gendisk(sdkp->disk);
2129         sd_shutdown(dev);
(gdb)
```

objdump

Use the objdump utility to glean useful information from the object file.

```
$ man objdump
```

```
--snip--
```

```
-g
```

```
--debugging
```

```
    Display debugging information. This attempts to parse
debugging
information stored in the file and print it out using a C like syntax.
Only certain types of debugging information have been implemented. Some
other types are supported by readelf -w.
```

```
-e
```

```
--debugging-tags
```

```
    Like -g, but the information is generated in a format compatible
with ctags tool.
```

```
-d
```

```
--disassemble
```

```
    Display the assembler mnemonics for the machine instructions
from objfile.
This option only disassembles those sections which are expected to contain
instructions.
```

```
-D
```

```
--disassemble-all
```

```
    Like -d, but disassemble the contents of all sections, not just
those expected to contain instructions.
```

```
--snip--
```

```
-S
```

```
--source
```

```
    Display source code intermixed with disassembly, if possible.
Implies -d.
```

```
--snip--
```

```
--start-address=address
```

```
    Start displaying data at the specified address. This affects
the output of the -d, -r and -s options.
```

```
--stop-address=address
```

```
    Stop displaying data at the specified address. This affects the
output of the -d, -r and -s options.
```

```
--snip--
```

```
$
```


A very useful piece of information one obtains from a backtrace ('bt'), either from gdb directly, or indirectly, via KGDB's or KDB's backtrace command, is the exact location of where the processor was in a function when the Oops occurred. We can easily search for and display the disassembled code for this offset (and the code around it) using objdump.

For example, we earlier saw a demo of generating an Oops using the hello_oops.ko kernel module, which of course did this by dereferencing a NULL pointer.

Here's a snippet of that Oops dump:

```
...
CPU:      0
EIP:      0060:[<d085d041>]   Tainted: G      U VLI
EFLAGS: 00010286   (2.6.16.21-0.8-default #1)
EIP is at procread+0x1d/0x48 [hello_oops]
eax: 00000039   ebx: d085d024   ecx: c9ec9f40   edx: c9ec9f40
esi: 00000400   edi: 00000400   ebp: c3e55000   esp: c9ec9f38
ds: 007b   es: 007b   ss: 0068
Process cat (pid: 6154, threadinfo=c9ec8000 task=c197aab0)
Stack: <0>c3e55000 d085d132 0000180a 00000000 d085d024 c0174370 00000400 c9ec9f70
       00000000 00000400 08051000 00000000 cfc11e40 00000000 00000000 c87c3740
       c0174270 08051000 00000400 c014b351 c9ec9fa4 c87c3740 ffffffff 00000400
Call Trace:
[<d085d024>] procread+0x0/0x48 [hello_oops]
[<c0174370>] proc_file_read+0x100/0x234
[<c0174270>] proc_file_read+0x0/0x234
[<c014b351>] vfs_read+0xa8/0x14d
[<c014b6bd>] sys_read+0x3c/0x63
[<c010299b>] sysenter_past_esp+0x54/0x79
Code: 68 9c d0 85 d0 e8 9d ac 8b ef 83 c4 0c c3 53 ba 00 e0 ff ff 6a 00 21 e2 8b
12 ff b2 b0 00 00 00 68 32 d1 85 d0 50 e8 8f 81 94 ef <ff> 35 00 00 00 00 89 c3 68
91 d0 85 d0 68 6b d1 85 d0 e8 65 ac
```

You can see above (highlighted in bold) that the processor **was at the procread function at an offset of 0x1d bytes from the start of the function (whose length is approximately 0x48 bytes) when the Oops occurred.**

So what's at offset 0x1d? Use objdump to easily discover the same:

```
$ objdump -g -d hello_oops.o
```

```
hello_oops.o:      file format elf32-i386
```

```
Disassembly of section .text:
```

```
00000000 <procread>:
   0:   53                      push    %ebx
   1:   ba 00 f0 ff ff         mov     $0xffffffff00,%edx
   6:   21 e2                  and     %esp,%edx
   8:   8b 0a                  mov     (%edx),%ecx
  a:   ff 72 10              pushl   0x10(%edx)
  d:   ff b1 94 00 00 00     pushl   0x94(%ecx)
 13:   68 00 00 00 00         push    $0x0
 18:   50                      push    %eax
 19:   e8 fc ff ff ff         call    1a <procread+0x1a>
 1e:   ff 35 00 00 00 00     pushl   0x0
```

```

24:  89 c3          mov    %eax,%ebx
26:  68 39 00 00 00  push    $0x39
2b:  68 44 00 00 00  push    $0x44
...
...
$

```

Another example, Mixing source and assembler:

```
$ objdump -g -S hang_panic.o
```

```
hang_panic.o:      file format elf32-i386
```

```
Disassembly of section .text:
```

```

00000000 <ding>:
#define procname      "lab2_test"

static int ding(unsigned long dope)
{
    printk(KERN_INFO
0:    50                      push    %eax
1:    68 00 00 00 00          push    $0x0
6:    68 0b 00 00 00          push    $0xb
b:    e8 fc ff ff ff          call    c <ding+0xc>
    "%s: In timeout function \"%ding\" now...dope=%ld\n",
    MODULE_NAME, dope);
/*
info: PID %d, interrupt-context: %d, processor # %d\n",
    MODULE_NAME, dope, current->pid,
    in_interrupt() ? 1:0, smp_processor_id());
*/
    return 0;
10:   83 c4 0c                add     $0xc,%esp
}
13:   31 c0                  xor     %eax,%eax

/* read callback */
static int my_procread(char *buf, char **start, off_t offset,
    int count, int *eof, void *data)
{
    int n;
    struct timer_list timer;

    n = sprintf(buf, "In proc read callback.\n\
        current->pid, smp_processor_id());

    /* Setup a timeout */

```



```

:objdump: hang_panic.o: no recognized debugging information
    printk(KERN_INFO "%s setting up timer (1s) now..\n",
           MODULE_NAME);
    init_timer (&timer);
    timer.expires = HZ;      /* 1 sec */
    timer.function = ding;
    timer.data = jiffies;
    add_timer (&timer);

    *eof=1;
    return n;
}

/*
 * function hang_panic_init
 */
static int __init hang_panic_init_module(void)
{
    if (!create_proc_read_entry (procname, 0444, /* mode */
                                NULL,          /* parent dir is /proc */
                                my_procread,    /* read callback */
                                NULL )) {      /* client data or tag */
        printk(KERN_ALERT "Proc entry creation failure,
aborting..\n");
    15:      c3                                ret

00000016 <my_procread>:
    16:  53                                push    %ebx
    17:  ba 00 f0 ff ff                    mov     $0xfffff000,%edx
    1c:  83 ec 30                          sub     $0x30,%esp
    1f:  21 e2                              and     %esp,%edx
    21:  8b 0a                              mov     (%edx),%ecx
    23:  ff 72 10                          pushl   0x10(%edx)
    26:  ff b1 94 00 00 00                pushl   0x94(%ecx)
    2c:  68 3d 00 00 00                    push    $0x3d
    31:  50                                push    %eax
    32:  e8 fc ff ff ff                    call    33 <my_procread+0x1d>
    37:  68 00 00 00 00                    push    $0x0
    3c:  89 c3                              mov     %eax,%ebx
    3e:  68 76 00 00 00                    push    $0x76
    43:  e8 fc ff ff ff                    call    44 <my_procread+0x2e>
    48:  83 c4 18                          add     $0x18,%esp
    4b:  ba 01 00 00 00                    mov     $0x1,%edx
    50:  89 e0                              mov     %esp,%eax
    52:  89 54 24 0c                        mov     %edx,0xc(%esp)
    56:  8b 15 00 00 00 00                mov     0x0,%edx
    5c:  b9 ad 4e ad de                    mov     $0xdead4ead,%ecx
    61:  c7 44 24 20 00 00 00            movl    $0x0,0x20(%esp)
    68:  00
    69:  89 54 24 1c                        mov     %edx,0x1c(%esp)
    6d:  ba e8 03 00 00                    mov     $0x3e8,%edx

```

```

72:  c7 44 24 14 6e ad 87    movl    $0x4b87ad6e,0x14(%esp)
79:  4b
7a:  89 4c 24 10            mov     %ecx,0x10(%esp)
7e:  c7 44 24 08 e8 03 00    movl    $0x3e8,0x8(%esp)
85:  00
86:  c7 44 24 18 00 00 00    movl    $0x0,0x18(%esp)
8d:  00
8e:  e8 fc ff ff ff        call    8f <my_procread+0x79>
93:  8b 44 24 3c            mov     0x3c(%esp),%eax
                                return -ENOMEM;
    }

    /* Delay for 1 second; just for the heck of it */
    printk(KERN_INFO "%s delaying for 1s now..\n", MODULE_NAME);
    set_current_state (TASK_INTERRUPTIBLE);
...
...
79:  c3                    ret
7a:  68 2d 01 00 00        push    $0x12d
7f:  68 00 00 00 00        push    $0x0
84:  68 31 01 00 00        push    $0x131
89:  e8 fc ff ff ff        call    8a <init_module+0x8a>
8e:  31 c0                xor     %eax,%eax
90:  83 c4 0c            add     $0xc,%esp
93:  c3                    ret
$

```

IMP: How exactly can one disassemble the kernel image file for a given function (by name)?

The following shell script lets you achieve this, using the extremely versatile tool `objdump`; it relies on your providing the path to the *uncompressed* kernel image (`vmlinux`) and the function name you'd like to disassemble. Using the '-S' switch with `objdump` even attempts to intermix C source with assembly!

```

$ cat disasfunc.sh
#!/bin/sh
# Source: http://kgdb.linsyssoft.com/downloads/miscmacros
#
# Eg.
# ./disasfunc.sh <path/to/>vmlinux do_fork
#
if [ $# != 2 ]
then
    echo disasfun objectfile functionname
    exit 1
fi
OBJFILE=$1
FUNNAME=$2

```

```

ADDRSZ=`objdump -t $OBJFILE | gawk -- '{
    if (\\$3 == \\\"F\\\" && \\$6 == \\\"$FUNNAME\\\") {
        printf(\\\"%s %s\\\", \\$1, \\$5)
    }
}'`
ADDR=`echo $ADDRSZ | gawk '{ printf(\\\"%s\\\", \\$1)}'`
SIZE=`echo $ADDRSZ | gawk '{ printf(\\\"%s\\\", \\$2)}'`
if [ -z \"$ADDR\" -o -z \"$SIZE\" ]
then
    echo Cannot find address or size of function $FUNNAME
    exit 2
fi
objdump -S $OBJFILE --start-address=0x$ADDR --stop-address=$((0x$ADDR +
0x$SIZE))
$

```

Eg.

```
$ ./disasfunc.sh <path>/<to>/vmlinux strlen
```

```
<path>/<to>/vmlinux:      file format elf32-i386
```

Disassembly of section .text:

```

08078ac0 <strlen>:
EXPORT_SYMBOL(memscan);
#endif

```

```

#ifdef __HAVE_ARCH_STRNLEN
size_t strlen(const char *s, size_t count)
{
    8078ac0: 55                push    %ebp
    8078ac1: 89 e5            mov     %esp,%ebp
        int d0;
        int res;
        asm volatile("movl %2,%0\n\t"
    8078ac3: 8b 55 0c         mov     0xc(%ebp),%edx
    8078ac6: 8b 4d 08         mov     0x8(%ebp),%ecx
    8078ac9: 89 c8            mov     %ecx,%eax
    8078acb: eb 06            jmp     8078ad3 <strlen+0x13>
    8078acd: 80 38 00         cmpb    $0x0,(%eax)
    8078ad0: 74 07            je      8078ad9 <strlen+0x19>
    8078ad2: 40              inc     %eax
    8078ad3: 4a              dec     %edx
    8078ad4: 83 fa ff         cmp     $0xffffffff,%edx
    8078ad7: 75 f4            jne     8078acd <strlen+0xd>
    8078ad9: 29 c8            sub     %ecx,%eax
        "3:\tsubl %2,%0"
        : "=a" (res), "=&d" (d0)
        : "c" (s), "1" (count)
        : "memory");
}
#endif

```

```

    return res;
}
8078adb: 5d          pop    %ebp
8078adc: c3          ret
$

```

Setting gdb debugging flags:

```
EXTRA_CFLAGS += -O0 -DDEBUG -g -ggdb -g-dwarf4
```

```
$ man gcc
```

```
...
...
```

```
-g
```

Produce debugging information in the operating system's native format (stabs, COFF, XCOFF, or DWARF). GDB can work with this debugging information.

On most systems that use stabs format, `-g` enables use of extra debugging information that only GDB can use; this extra information makes debugging work better in GDB but will probably make other debuggers crash or refuse to read the program. If you want to control for certain whether to generate the extra information, use `-gstabs+`, `-gstabs`, `-gxcoff+`, `-gxcoff`, or `-gvms` (see below).

Unlike most other C compilers, GCC allows you to use `-g` with `-O`. The shortcuts taken by optimized code may occasionally produce surprising results: some variables you declared may not exist at all; flow of control may briefly move where you did not expect it; some statements may not be executed because they compute constant results or their values were already at hand; some statements may execute in different places because they were moved out of loops.

Nevertheless it proves possible to debug optimized output. This makes it reasonable to use the optimizer for programs that might have bugs.

The following options are useful when GCC is generated with the capability for more than one debugging format.

```
-ggdb
```

Produce debugging information for use by GDB. This means to use the most expressive format available (DWARF, stabs, or the native format if neither of those are supported), including GDB extensions if at all possible.

```
...
```

```
-gdwarf
```

```
-gdwarf-version
```

Produce debugging information in DWARF format (if that is supported). The value of version may be either 2, 3, 4 or 5; the default version for most targets is 4. DWARF Version 5 is only experimental.

```
...
```

```
$
```

Enabling these flags in the Makefile is pretty common practise to enable symbolic debugging information.

More gcc / as flags:

```
EXTRA_CFLAGS += -DDEBUG -g -ggdb -gdwarf-4 -Wa, -a, -ad, -as
```

These gcc flags cause the **assembler** ('as') to emit a wealth of information: essentially, one is able to **view a mix of the source code** (this includes the source of all '#include 'd files), the **resulting assembly and machine language** generated by the compiler/assembler, all in a printer-friendly (paginated) ASCII text format.

Specifically, the meanings of the flags are:

```
$ man gcc
```

```
...
...
```

```
    Passing Options to the Assembler
```

```
    You can pass options to the assembler.
```

```
    -Wa,option
```

```
        Pass option as an option to the assembler.  If option contains
commas, it is split into multiple options at the commas.
```

```
...
```

```
$
```

```
$ man as
```

```
...
...
```

```
    -a[cdhlms]
```

```
        Turn on listings, in any of a variety of ways:
```

```
        -ac omit false conditionals
```

```
        -ad omit debugging directives
```

```
        -ah include high-level source
```

```
        -al include assembly
```

```
        -am include macro expansion
```

```
        -an omit forms processing
```

```
        -as include symbols
```

```
    =file
```

```
        set the name of the listing file
```

You may combine these options; for example, use -aln for assembly listing without forms processing. The =file option, if used, must be the last one. By itself, -a defaults to -ahls.

```
...
```

```
$
```

When make is run, we get output similar to the following:

\$ make

make -C /lib/modules/2.6.9-34.ELsmp/build SUBDIRS=/mnt/<...>/lab2 modules

make[1]: Entering directory `/usr/src/kernels/2.6.9-34.EL-smp-i686'

CC [M] /mnt/<...>/lab2/hang_panic.o

GAS LISTING page 1

```

1          .file "hang_panic.c"
9          .Ltext0:
10         .section .rodata.str1.1,"aMS",@progbits,1
11         .LC0:
12 0000 68616E67      .string "hang_panic"
12      5F70616E
12      696300
13         .LC1:
14 000b 3C363E25      .string "<6>%s: In timeout function \"ding\"
now...dope=%ld\n"
14      733A2049
14      6E207469
14      6D656F75
14      74206675
15         .text
17         ding:
18         .LFB463:
19         .file 1 "/mnt/<...>/lab2/hang_panic.c"
1:/mnt/<...>/lab2/hang_panic.c **** /*
2:/mnt/<...>/lab2/hang_panic.c ****  * hang_panic.c v1.0 <date>
3:/mnt/<...>/lab2/hang_panic.c ****  *
4:/mnt/<...>/lab2/hang_panic.c ****  *
5:/mnt/<...>/lab2/hang_panic.c ****  * Simple kernel module that delays
for a second during init,
6:/mnt/<...>/lab2/hang_panic.c ****  * creates a proc file; when the
proc file is read, a timeout
7:/mnt/<...>/lab2/hang_panic.c ****  * is setup (for 1 second). Once
timeout occurs, the function
8:/mnt/<...>/lab2/hang_panic.c ****  * 'ding' is called.
9:/mnt/<...>/lab2/hang_panic.c ****  *
10:/mnt/<...>/lab2/hang_panic.c ****  * But, of course, it has a bug (or
bugs). Read, run, test,
11:/mnt/<...>/lab2/hang_panic.c ****  * debug, correct, verify, repeat.
12:/mnt/<...>/lab2/hang_panic.c ****  *
13:/mnt/<...>/lab2/hang_panic.c ****  * (c) kaiwan.
14:/mnt/<...>/lab2/hang_panic.c ****  *
15:/mnt/<...>/lab2/hang_panic.c ****  * This program is free software;
you can redistribute it
16:/mnt/<...>/lab2/hang_panic.c ****  * and/or modify it under the terms
of the GNU Library
[...
31:/mnt/<...>/lab2/hang_panic.c **** */
32:/mnt/<...>/lab2/hang_panic.c ****

```

```

33:/mnt/<...>/lab2/hang_panic.c **** #include <linux/module.h>
34:/mnt/<...>/lab2/hang_panic.c **** #include <linux/kernel.h>
35:/mnt/<...>/lab2/hang_panic.c **** #include <linux/init.h>
36:/mnt/<...>/lab2/hang_panic.c **** #include <linux/sched.h>
37:/mnt/<...>/lab2/hang_panic.c **** #include <linux/proc_fs.h>
38:/mnt/<...>/lab2/hang_panic.c **** #include <linux/timer.h>
39:/mnt/<...>/lab2/hang_panic.c **** #include <linux/interrupt.h>
40:/mnt/<...>/lab2/hang_panic.c ****

```

GAS LISTING

page 2

```

41:/mnt/<...>/lab2/hang_panic.c **** #define MODULE_VER    "1.0"
42:/mnt/<...>/lab2/hang_panic.c **** #define MODULE_NAME    "hang_panic"
43:/mnt/<...>/lab2/hang_panic.c **** #define procname      "lab2_test"
44:/mnt/<...>/lab2/hang_panic.c ****
45:/mnt/<...>/lab2/hang_panic.c **** static int ding(unsigned long dope)
46:/mnt/<...>/lab2/hang_panic.c **** {
20                .loc 1 46 0
21                .LVL0:
47:/mnt/<...>/lab2/hang_panic.c ****      printk(KERN_INFO
22                .loc 1 47 0
23 0000 50          pushl %eax
24                .LCFI0:
25 0001 68000000     pushl $.LC0
25                00
26                .LCFI1:
27 0006 680B0000     pushl $.LC1
27                00
28                .LCFI2:
29 000b E8FCFFFF     call printk
29                FF
30                .LVL1:
48:/mnt/<...>/lab2/hang_panic.c ****      "%s: In timeout
function \"ding\" now...dope=%ld\\n",
49:/mnt/<...>/lab2/hang_panic.c ****      MODULE_NAME, dope);
50:/mnt/<...>/lab2/hang_panic.c **** /*
51:/mnt/<...>/lab2/hang_panic.c **** info: PID %d, interrupt-context: %d,
processor # %d\\n",
52:/mnt/<...>/lab2/hang_panic.c ****      MODULE_NAME, dope, current-
>pid,
53:/mnt/<...>/lab2/hang_panic.c ****      in_interrupt() ? 1:0,
smp_processor_id());
54:/mnt/<...>/lab2/hang_panic.c **** */
55:/mnt/<...>/lab2/hang_panic.c ****      return 0;
31                .loc 1 55 0
32 0010 83C40C      addl $12,%esp
33                .LCFI3:
56:/mnt/<...>/lab2/hang_panic.c **** }
34                .loc 1 56 0
35 0013 31C0        xorl %eax,%eax
36 0015 C3          ret

```

```

37          .LFE463:
39          .section .rodata.str1.1
40          .LC2:

...
--snip--
...

63:/mnt/<...>/lab2/hang_panic.c ****      int n;
64:/mnt/<...>/lab2/hang_panic.c ****      struct timer_list timer;
65:/mnt/<...>/lab2/hang_panic.c ****
66:/mnt/<...>/lab2/hang_panic.c ****      n = sprintf(buf, "In proc read
callback.\n\
67:/mnt/<...>/lab2/hang_panic.c **** PID %d running on processor # %d\n",
68:/mnt/<...>/lab2/hang_panic.c ****      current->pid,
smp_processor_id());
69:/mnt/<...>/lab2/hang_panic.c ****
70:/mnt/<...>/lab2/hang_panic.c ****      /* Setup a timeout */
71:/mnt/<...>/lab2/hang_panic.c ****      printk(KERN_INFO "%s setting up
timer (1s) now..\n",
90          .loc 1 71 0
91 0037 68000000      pushl $.LC0
91          00
92          .LCFI10:
93          .loc 3 9 0
94 003c 89C3      movl %eax,%ebx
95          .LVL6:

...
...

323 00bb C3      ret

...
333          __mod_author124:
334 0000 61757468      .string "author=My Name"
334          6F723D4D
334          79204E61
334          6D6500
335 000f 00000000      .align 32
...
...
jiffies
__mod_timer
create_proc_entry
schedule_timeout
Building modules, stage 2.
MODPOST
CC      /mnt/<...>/lab2/hang_panic.mod.o
LD [M]   /mnt/<...>/lab2/hang_panic.ko
make[1]: Leaving directory `/usr/src/kernels/2.6.9-34.EL-smp-i686'
$

```

Misc:

See the book “Linux Debugging and Performance Tuning”, by Steve Best for excellent “real-world” kdb Debug sessions; page 354 onwards.

Kinda interesting:

[Determining cause of Linux kernel panic – on stackexchange](#)

- Spotted this one:

<http://lkml.org/lkml/2010/6/1/138>

which the kernel bugzilla reported [here](#).

kaiwanTECH Linux OS Corporate Training Programs

Please do check out our current offering of world-class, seriously-valuable, high on returns, technical Linux OS corporate training programs here: <http://bit.ly/ktcorp>