



# ***BUILDING THE MODERN LINUX KERNEL***

## Important Notice : Courseware - Legal

This courseware is both the product of the author and of freely available opensource and/or public domain materials. Wherever external material has been shown, it's source and ownership have been clearly attributed. We acknowledge all copyrights and trademarks of the respective owners.

The contents of the **courseware PDFs are considered proprietary** and thus cannot be copied or reproduced in any form whatsoever without the explicit written consent of the author.

Only the programs - **source code** and binaries (where applicable) - that form part of this courseware, and that are made available to the participant, are released under the terms of the [permissive MIT license](#).

Under the terms of the MIT License, you can certainly use the source code provided here; you must just attribute the original source (author of this courseware and/or other copyright/trademark holders).

**VERY IMPORTANT ::** Before using this source(s) in your project(s), you **\*MUST\*** check with your organization's legal staff that it is appropriate to do so.

The courseware PDFs are **\*not\*** under the MIT License, they are to be kept confidential, non-distributable without consent, for your private internal use only.

The duration, contents, content matter, programs, etc. contained in this courseware and companion participant VM are subject to change at any point in time without prior notice to individual participants.

Care has been taken in the preparation of this material, but there is no warranty, expressed or implied of any kind, and we can assume no responsibility for any errors or omissions. No liability is assumed for incidental or consequential damages in connection with or arising out of the use of the information or programs contained herein.

2000-2024 Kaiwan N Billimoria  
kaiwanTECH, Bangalore, India.

### **kaiwanTECH Linux OS Corporate Training Programs**

*Please do check out our current offering of world-class, seriously-valuable, high on returns, technical Linux OS corporate training programs here: <http://bit.ly/ktcorp>*

## Building the 2.6 (onwards) Linux Kernel - Quick Step Summary, biased to x86

0. Install all required packages (see below)

1. Download and Extract the kernel source tree

a) Download the new kernel source; f.e. for the 5.4.1 kernel source tree:

wget

<https://mirrors.edge.kernel.org/pub/linux/kernel/v5.x/linux-5.4.1.tar.xz>

b) Verify it with gpg [see procedure below]

c) Extract it into some location under your home directory

`tar xf linux-5.4.1.tar.xz`

(Alternatively, one can always use git(1) to download a particular version)

2. Configuration : select kernel support options as required for the new kernel  
(*make [x|g|menu]config*);

`make [ARCH=<arch>] menuconfig`

is recommended. ARCH determines the architecture (cpu) the kernel is being configured and built for, the default is x86; others are:

alpha	arm	c6x	h8300	ia64	m68k	mips	nios2
parisc	riscv	sh	um	x86	arc	arm64	csky
hexagon		Kconfig	microblaze		nds32	openrisc	powerpc s390
sparc	unicore32		xtensa				

3. Build the kernel and loadable modules:

`make -j[n]`

Builds the compressed kernel image (*arch/<arch>/boot/[b|z|u]image*), uncompressed kernel image (*./vmlinux*), *System.map* and kernel modules.

4. Install the just-built kernel modules with

`sudo make [INSTALL_MOD_PATH=<path/to/modules/dir>] modules_install`  
Installs the kernel modules under */lib/modules/`uname -r`*, or, if defined, under *INSTALL\_MOD\_PATH*

5. Set up your boot options as required (LILO / GRUB)

*For x86:*

`sudo make install`

- a) Creates and installs the initrd image under /boot
- b) Updates the bootloader configuration file to boot the new kernel (first entry)

Step 1b. Verifying the kernel source

Ref: <https://www.kernel.org/signature.html>

For the 6.1.25 kernel:

1. Download the kernel source; we download this file:  
<https://mirrors.edge.kernel.org/pub/linux/kernel/v6.x/linux-6.1.25.tar.xz>
2. We require the public key(s); retrieve it from the keyserver (let's get Linus's and Greg-KH's public keys):  
  

```
$ gpg2 --locate-keys torvalds@kernel.org gregkh@kernel.org
gpg: key 38DBBDC86092693E: public key "Greg Kroah-Hartman
<gregkh@kernel.org>" imported
gpg: Total number processed: 1
gpg:             imported: 1
gpg: key 79BE3E4300411886: public key "Linus Torvalds
<torvalds@kernel.org>" imported
gpg: Total number processed: 1
...
```
3. Now let's try to verify the tarred-and-compressed source tree:  

```
$ xzcat ../linux-6.1.25.tar.xz | gpg2 --verify ./linux-
6.1.25.tar.sign -
gpg: Signature made Thursday 20 April 2023 04:06:06 PM IST
gpg:             using RSA key
647F28654894E3BD457199BE38DBBDC86092693E
gpg: Good signature from "Greg Kroah-Hartman
<gregkh@kernel.org>" [unknown]
gpg: WARNING: This key is not certified with a trusted
signature!
gpg:             There is no indication that the signature belongs
to the owner.
...
```
4. As mentioned [here](#):  
To make the "WARNING" message go away you can indicate that you choose to trust that key using TOFU - "Trust on First Use" :

```
$ xzcat ../linux-6.1.25.tar.xz | gpg2 --trust-model tofu --
verify ../linux-6.1.25.tar.sign -
gpg: Signature made Thursday 20 April 2023 04:06:06 PM IST
gpg:                using RSA key
647F28654894E3BD457199BE38DBBDC86092693E
gpg: Good signature from "Greg Kroah-Hartman
<gregkh@kernel.org>" [full]
gpg: gregkh@kernel.org: Verified 1 signatures in the past 0
seconds. Encrypted
    0 messages.
$
There, it's verified!
```

5. If it shows “BAD SIGNATURE” then there’s an issue; don’t use this kernel source! else, all well, untar it and build it!

A script to do precisely this [is available here!](#)

**TIP:** Software packages to install for kernel build and kernel dev (below, on an Ubuntu 20.04 LTS system):

(Not all packages listed below are required for the kernel build itself)

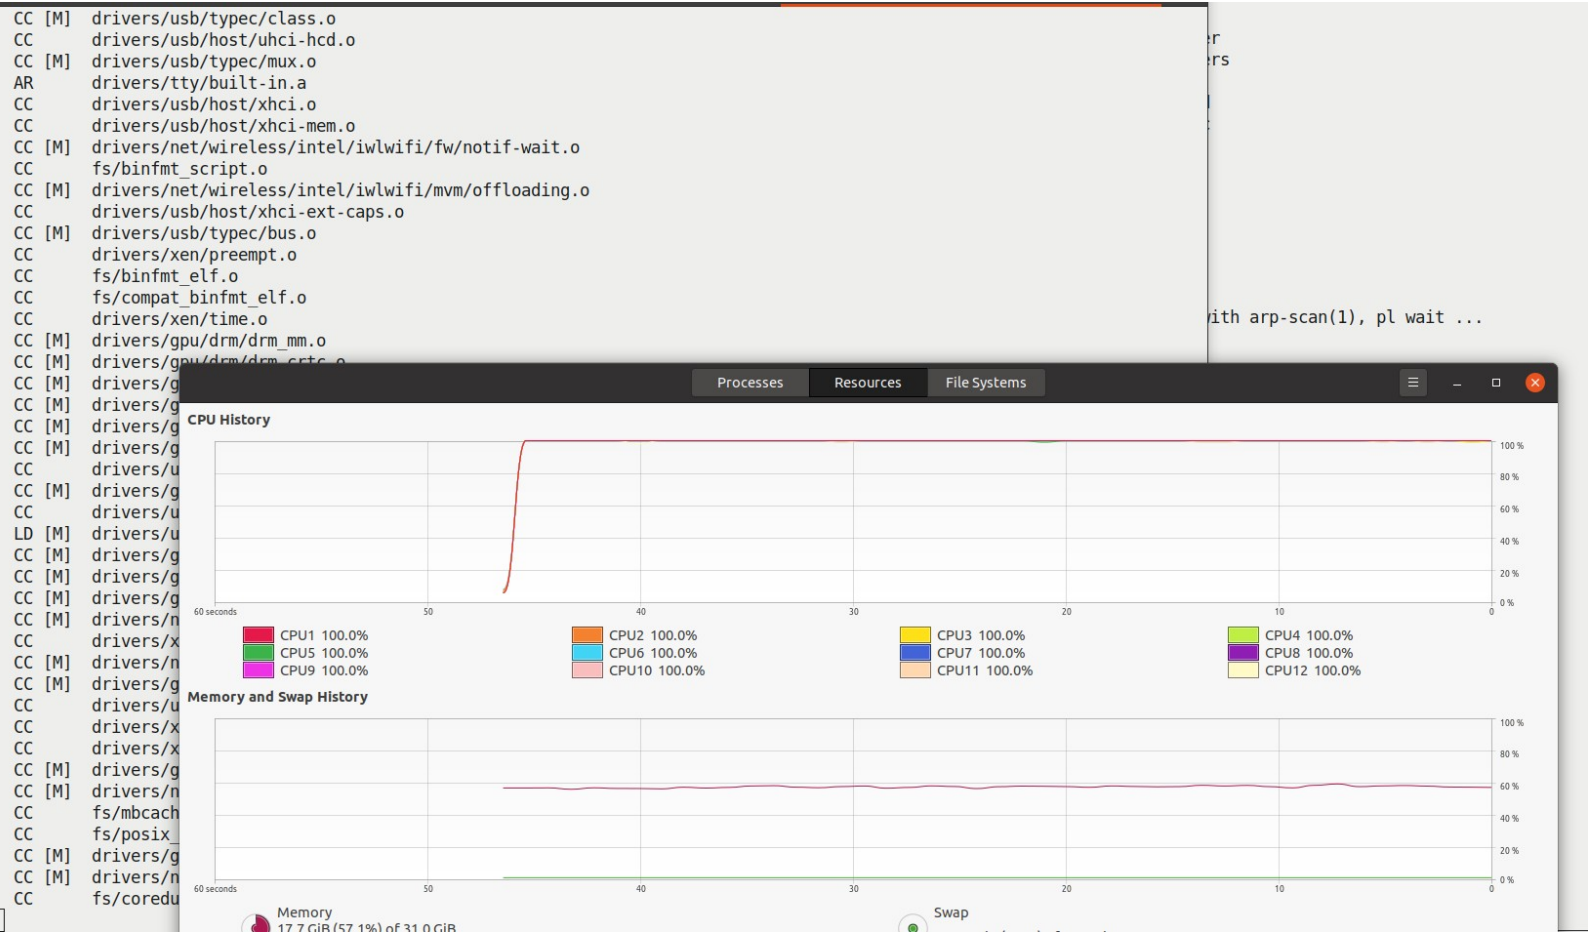
```
sudo apt update
sudo apt install -y gcc make perl

sudo apt install -y \
    bison build-essential flex \
    libelf-dev libssl-dev ncurses-dev \
    pahole tar util-linux xz-utils zstd
```

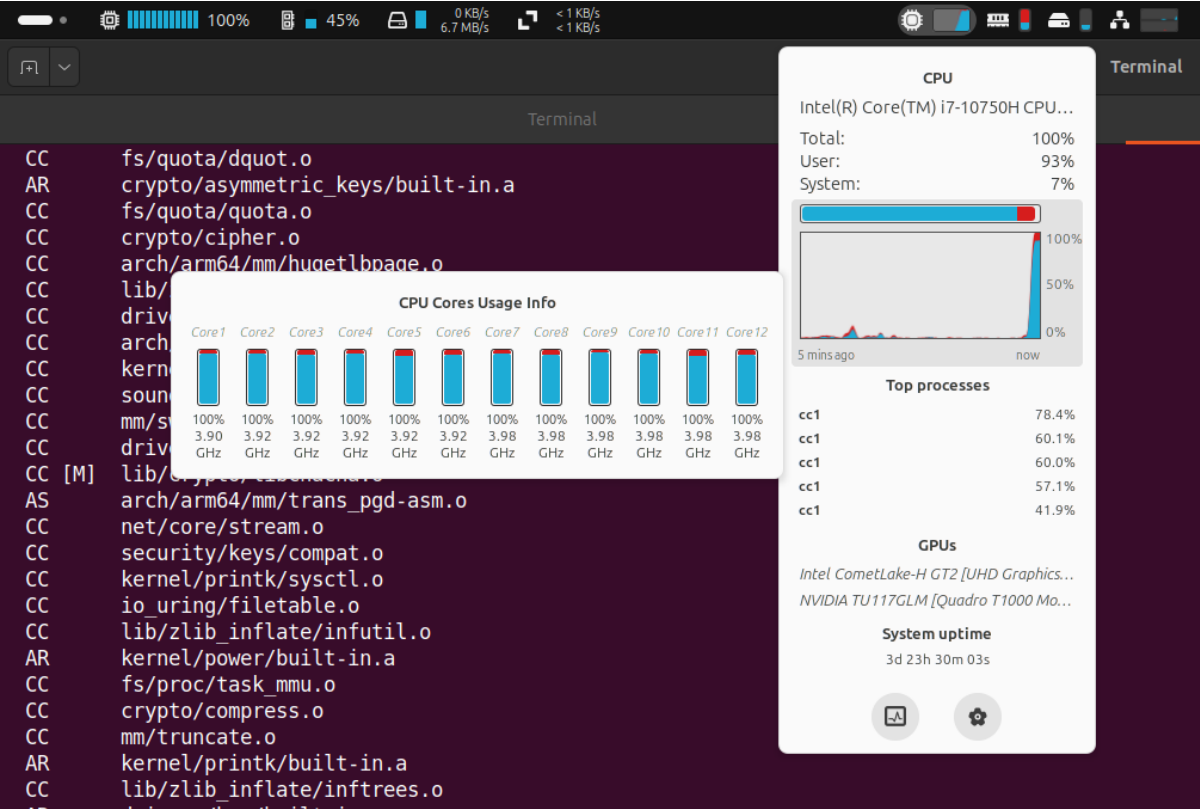
As well, see this convenience script:

[https://github.com/PacktPublishing/Linux-Kernel-Programming\\_2E/blob/main/ch1/pkg\\_install4ubuntu\\_lkp.sh](https://github.com/PacktPublishing/Linux-Kernel-Programming_2E/blob/main/ch1/pkg_install4ubuntu_lkp.sh)

*Partial screenshot while building the Linux kernel on a host with 12 CPU cores (with make -j24 )!*



Another, along similar lines... (the applets on the top bar, right side are Astra Monitor!; on the left are the TopHat applets):




*Resources:*

- ✓ [How to Build Linux Kernel From Scratch {Step-By-Step Guide}](#), Nov 2020 [builds the recent 6.0.7 kernel!]
- ✓ [How to compile and install Linux Kernel 5.6.9 from source code](#) Author: Vivek Gite Last updated: May 2, 2020
- ✓ [How to Compile Linux Kernel from Source to Build Custom Kernel](#)
- ✓ Useful! Which kernel configurables should we turn On minimally? Hard to answer, depends, but a good summary available here: [systemd README](#) : see the “Requirements” section – kernel configurables
- ✓ [How to Configure the GRUB2 Boot Loader’s Settings](#)  
[in brief: edit /etc/default/grub ; sudo update-grub ]
- ✓ Kbuild:  
[Kbuild: the Linux Kernel Build System](#), LJ, Dec 2012  
[How does kbuild actually work? \[SO\]](#)



## Kernel Development Model and Releases

[Source](#)



### IMPORTANT

*In production, strive to use an updated – the latest stable – LTS kernel.*

Ref: [What Stable Kernel Should I Use, Greg Kroah-Hartman, Aug 2018](#)



### More stability for the kernel source tree

- ▶ Issue: bug and security fixes only released for most recent stable kernel versions.
- ▶ Some people need to have a recent kernel, but with long term support for security updates.
- ▶ You could get long term support from a commercial embedded Linux provider.
- ▶ You could reuse sources for the kernel used in Ubuntu Long Term Support releases (5 years of free security updates).
- ▶ The <http://kernel.org> front page shows which versions will be supported for some time (up to 2 or 3 years), and which ones won't be supported any more ("EOL: End Of Life")

mainline:	3.14-rc8	2014-03-25
stable:	3.13.7	2014-03-24
stable:	3.11.10 [EOL]	2013-11-29
longterm:	3.12.15	2014-03-26
longterm:	3.10.34	2014-03-24
longterm:	3.4.84	2014-03-24
longterm:	3.2.55	2014-02-15
longterm:	2.6.34.15 [EOL]	2014-02-10
longterm:	2.6.32.61	2013-06-10
linux-next:	next-20140327	2014-03-27

Kernel versions**Which kernel version to use?**

*Linux OS : Security and Hardening – An Overview*

## Modern OS Hardening Countermeasures

Common Hardening Countermeasures include

- 1) Using Managed Programming Language
- 2) Compiler Protections
- 3) Library Protection
- 4) Executable Space Protection
- 5) [K]ASLR (address space randomization)
- 6) Better Testing

**“If you are not using a stable / longterm kernel, your machine is insecure”**

**- Greg Kroah-Hartman**

**All this is good, great, but... the Most Important Thing:**

**“If you are not using the latest kernel, you don't have the most recently added security defenses, which, in the face of newly exploited bugs, may render your machine less secure than it could have been”**

← **Kees Cook**, Google (Pixel Security), KSPF lead dev

**Who will provide this (very) Long Term kernel support?**

- LTS (Long Term Stable) kernels
- SLTS (Super LTS) kernels too!

from the **Civil Infrastructure Platform (CIP)** group [\[link\]](#)

A **Linux Foundation (LF)** project

Src: <https://wiki.linuxfoundation.org/civilinfrastructureplatform/start>

As of Aug 2023:

### CIP activities

#### Kernel maintainership

The first action taken by the CIP project is to select and maintain Linux kernels for a very long time (+10 years). To achieve this goal we have a group of experts. All the actions and decisions related to the maintenance of the Linux kernels selected by the CIP project are available at [CIP Kernel maintenance](#).

The current released CIP kernels are as follows.

Version	Maintainer(s)	First Release	Projected EOL	Target Releases/Month*
SLTS v6.1	Nobuhiro Iwamatsu & Pavel Machek	2023-07-14	2033-08	2
SLTS v6.1-rt	Pavel Machek	2023-07-16	2033-08	1
SLTS v5.10	Nobuhiro Iwamatsu & Pavel Machek	2021-12-05	2031-01	1
SLTS v5.10-rt	Pavel Machek	2021-12-08	2031-01	0.5
SLTS v4.19	Nobuhiro Iwamatsu & Pavel Machek	2019-01-11	2029-01	1
SLTS v4.19-rt	Pavel Machek	2019-01-11	2029-01	0.5
SLTS v4.4	Ulrich Hecht	2017-01-17	2027-01	1
SLTS v4.4-rt	Pavel Machek	2017-11-16	2027-01	0.5

\*) Critical issues can trigger additional releases.

4.4 SLTS kernel support until at least 2026, possibly 2036!

4.19 SLTS kernel support including ARM64

More on which kernel to use can be seen here:  
[What Stable Kernel Should I Use, Greg Kroah-Hartman, Aug 2018](#)

...

Here's the my short list of what kernel you should use, ranked from best to worst options. I'll go into the details of all of these below, but if you just want the summary of all of this, here it is:

Hierarchy of what kernel to use, from best solution to worst:

- Supported kernel from your favorite Linux distribution
- Latest stable release
- Latest LTS release
- Older LTS release that is still being maintained

What kernel to never use:

- Unmaintained kernel release

...

### Older LTS Release

...

I will note that some manufacturers are already doing this today. Sony is one great example of this, updating to the latest 4.4.y release on many of their new phones for their quarterly security release. Another good example is the small company Essential which has been tracking the 4.4.y releases faster than anyone that I know of.

There is **one huge caveat** when using a kernel like this. The number of security fixes that get backported are not as great as with the latest LTS release, because the traditional model of the devices that use these older LTS kernels is a much more reduced user model. These kernels are not to be used in any type of “general computing” model where you have untrusted users or virtual machines, as the ability to do some of the recent Spectre-type fixes for older releases is greatly reduced, if present at all in some branches.

So again, only use older LTS releases in a device that you fully control, or lock down with a very strong security model (like Android enforces using SELinux and application isolation). Never use these releases on a server with untrusted users, programs, or virtual machines.

...

So, here’s a short list of different types of devices, and what I would recommend for their kernels:

- Laptop / Desktop: Latest stable release
- Server: Latest stable release or latest LTS release
- Embedded device: Latest LTS release or older LTS release if the security model used is very strong and tight.

---

### Quick Tips

#### ----- TIP -----

#### CONFIGURING the kernel:

Do **not skip** this step even if you are only upgrading one minor version. New configuration options are added in each release, and odd problems will turn up if the configuration files are not set up as expected. **If you want to carry your existing configuration to a new version with minimal work, use "make oldconfig", which will only ask you for the answers to new questions.**

```
cd <kernel-src-tree>
```

**make help**

'''

Configuration targets:

```

config      - Update current config utilising a line-oriented program
nconfig     - Update current config utilising a ncurses menu based
program
menuconfig  - Update current config utilising a menu based program
xconfig     - Update current config utilising a Qt based front-end
gconfig     - Update current config utilising a GTK+ based front-end
oldconfig   - Update current config utilising a provided .config
as base

```

[...]

&lt;&lt;

FYI

**Android** (AOSP) kernel configuration defaults:<https://android.googlesource.com/kernel/configs/>

Also, **clang** (LLVM) is used to build the AOSP (for all- kernel, modules and the root filesystem); to build the Android kernel:

```
make LLVM=1 LLVM_IAS=1 all
```

&gt;&gt;

**TIP****USEFUL! Common case:**

Build a kernel with appropriate configurables **for the (Linux) system you are currently running on** (or even another, for that matter):

**Approach 1**

Simple approach: simply copy in the existing kernel config, of the kernel you're currently running and use it as a starting point:

```
cp /boot/config-$(uname -r) .config
```

**Tip:**

Check which kernel configs exist under /boot and copy the one closest to the kernel version you're now going to build:

```
$ ls /boot/config-*
```

```
/boot/config-5.10.140-lkd-kernel  /boot/config-6.1.0-lkp2e-01+
/boot/config-6.1.11-lkp-kernel   /boot/config-5.15.0-43-generic
/boot/config-6.1.0-lkp2e-01+.old
```

So, if we're building, say, 6.1.23 kernel, copy the `/boot/config-6.1.11-lkp-kernel` to `.config` and then proceed!

## Approach 2

From [Documentation/admin-guide/README.rst](#)

```
...
"make localmodconfig" Create a config based on current config and loaded
modules (lsmod). Disables any module option that is not needed for the
loaded modules.
```

To create a localmodconfig for another machine, store the `lsmod` of that machine into a file and pass it in as a `LSMOD` parameter.

```
target$ lsmod > /tmp/mylsmod

host$ make LSMOD=/tmp/mylsmod localmodconfig
[...]
```

```
host$ make ...
```

The above also works when cross compiling.

```
...
```

## TIP

With kernel ver  $\geq 5.10$  kernel builds FAIL with:

```
make
...
AS      certs/system_certificates.o
make[1]: *** No rule to make target 'debian/canonical-revoked-
certs.pem', needed by 'certs/x509_revocation_list'.  Stop.
make[1]: *** Waiting for unfinished jobs....
CC      certs/blacklist.o
...
```

## Tip:

<https://askubuntu.com/questions/1362455/i-am-installing-kernel-in-my-ubuntu-but-getting-an-error>

Do this

```
scripts/config --disable SYSTEM_REVOCATION_KEYS
scripts/config --disable SYSTEM_TRUSTED_KEYS
```

Once these configs are disabled, it builds perfectly.

On Fedora, same thing, plus ensure that the *openssl* package is installed.

----- **TIP** -----

*>= 5.10 (or so) kernel build FAILS with:*  
**pahole required ...**

This is due to the kernel config CONFIG\_DEBUG\_INFO\_BTf being turned On. To fix this, **turn it Off**.

Kernel Hacking

Compile-time checks and compiler options  
Generate BTF typeinfo

NOTE- pahole (ver >= 1.16) being installed prevents this issue (but Ubuntu 20.04 and lower don't have a recent enough ver of pahole).

----- **TIP** -----

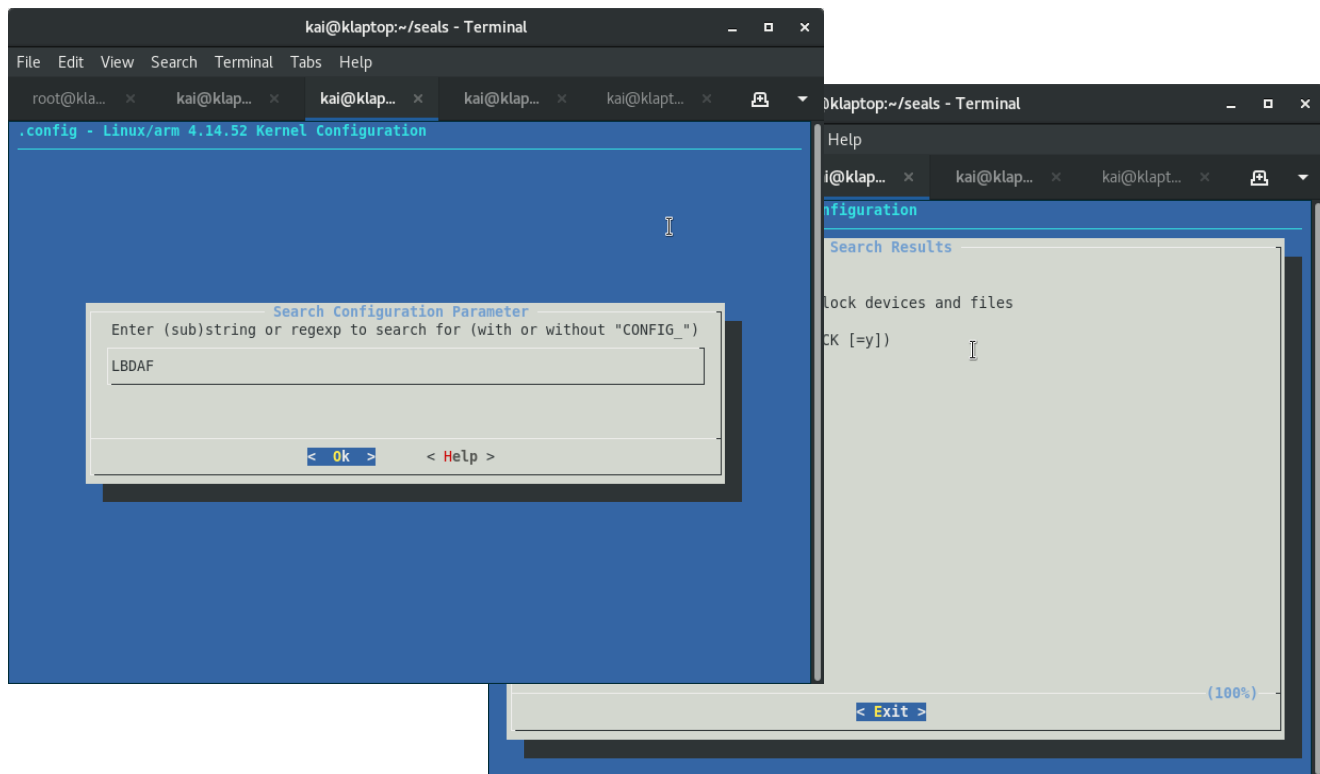
*Configure the kernel carefully!*  
On an ARM system:

```
ARM # mount -o remount,rw /
EXT4-fs (mmcblk0): Filesystem with huge files cannot be mounted RDWR
without CONFIG_LBDAF
EXT4-fs (mmcblk0): re-mounted. Opts: data=ordered
ARM #
```


*In order to mount it as 'rw', the ext4 filesystem requires the CONFIG\_LBDAF option to be set.*

----- **TIP** -----

So, where in the *make menuconfig* menu is this CONFIG\_LBDAF kernel config option?  
Search for it with the "usual" "/" operator!



Ah, we can now see: its under the menu item “*Enable the block layer*”.  
So navigate there:



Turn on the highlighted option above by toggling it with the spacebar, save and build.



----- **TIP** -----

### *Configuring the kernel non-interactively*

Use the *scripts/config* script !

The help screen shows the options:

#### **\$ scripts/config**

Manipulate options in a .config file from the command line.

Usage:

config options command ...

commands:

```
--enable|-e option    Enable option
--disable|-d option   Disable option
--module|-m option    Turn option into a module
--set-str option string
                      Set option to "string"
--set-val option value
                      Set option to value
--undefine|-u option  Undefine option
--state|-s option     Print state of option (n,y,m,undef)

--enable-after|-E beforeopt option
                      Enable option directly after other option
--disable-after|-D beforeopt option
                      Disable option directly after other option
--module-after|-M beforeopt option
                      Turn option into module directly after other
option
```

commands can be repeated multiple times

options:

```
--file config-file    .config file to change (default .config)
--keep-case|-k        Keep next symbols' case (dont' upper-case it)
```

config doesn't check the validity of the .config file. This is done at next make time.

By default, config will upper-case the given symbol. Use --keep-case to keep the case of all following symbols unchanged.

config uses 'CONFIG\_' as the default symbol prefix. Set the environment variable CONFIG\_ to the prefix to use. Eg.: CONFIG\_="FOO\_" config ...

\$

**TIP**

- With the GRUB2 bootloader, edit `/etc/default/grub` as root and add your custom kernel command-line parameters:

For eg.

```
'''
# GRUB_CMDLINE_LINUX_DEFAULT="quiet console=tty0 console=ttyS0,9600"
GRUB_CMDLINE_LINUX_DEFAULT="debug initcall_debug nolapic_timer 3"
'''
$ sudo update-grub
...
```

**TIP**

- When booting a kernel in a VM (virtual machine), it's often useful to turn off the local APIC timer; use the below option in the target kernel's command line (via the GRUB2 bootloader menu system):  
`nolapic_timer [X86-32,APIC] Do not use the local APIC timer.`

**TIP**

The kernel build fails with:

```
[...] fatal error: openssl/opensslv.h: No such file or directory
```

See: [OpenSSL missing during ./configure. How to fix?](#)

[...]

The OpenSSL library is usually already installed, but you have to install the header files. Depending on your Linux distribution, you'll need these packages:

- Red Hat, Fedora, CentOS - `openssl-devel`
- Debian, Ubuntu - `libssl-dev`
- Arch - `openssl`

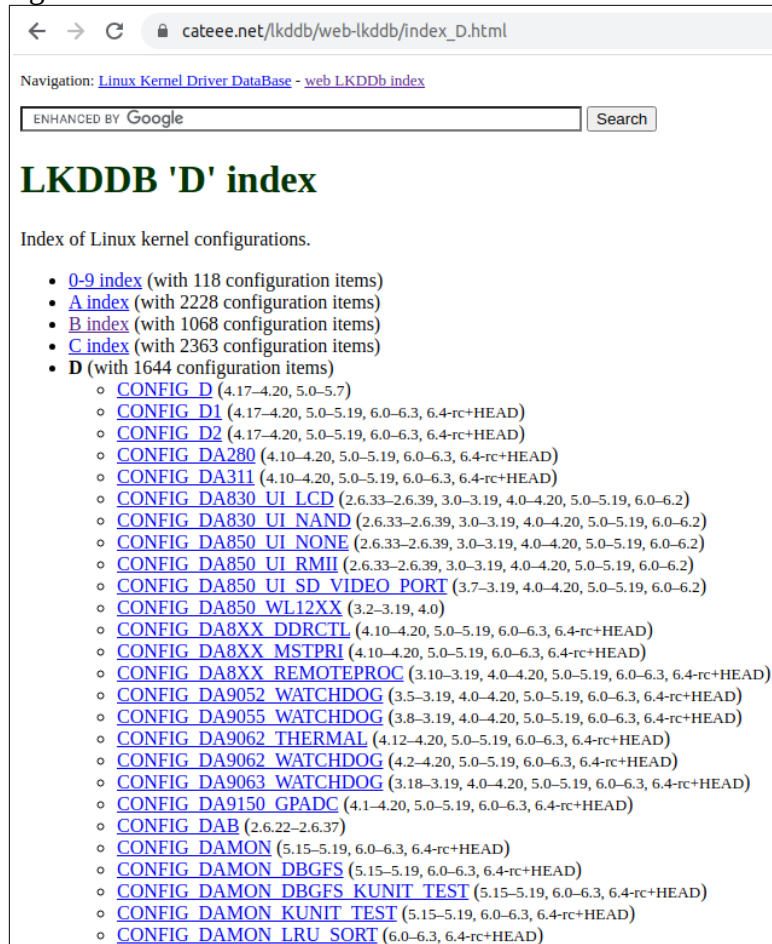
**TIP**

**Q. Are all kernel configurables (CONFIG\_FOO) shown in one place, along with which kernel versions they're supported on?**

**A. Yes!**

<https://cateee.net/lkddb/web-lkddb/>

Eg. screenshot:



### TIP

**Very useful! “make help”**

Output below from a recent Linux kernel (4.10.0-rc2):

**\$ make help**

Cleaning targets:

- clean - Remove most generated files but keep the config and enough build support to build external modules
- mrproper - Remove all generated files + config + various backup files
- distclean - mrproper + remove editor backup and patch files

Configuration targets:

- config - Update current config utilising a line-oriented program
- nconfig - Update current config utilising a ncurses menu based

```

        program
menuconfig - Update current config utilising a menu based program
xconfig   - Update current config utilising a Qt based front-end
gconfig   - Update current config utilising a GTK+ based front-end
oldconfig - Update current config utilising a provided .config as base
localmodconfig - Update current config disabling modules not loaded
localyesconfig - Update current config converting local mods to core
silentoldconfig - Same as oldconfig, but quietly, additionally update deps
defconfig  - New config with default from ARCH supplied defconfig
savedefconfig - Save current config as ./defconfig (minimal config)
allnoconfig - New config where all options are answered with no
allyesconfig - New config where all options are accepted with yes
allmodconfig - New config selecting modules when possible
alldefconfig - New config with all symbols set to default
randconfig  - New config with random answer to all options
listnewconfig - List new options
olddefconfig - Same as silentoldconfig but sets new symbols to their
               default value
kvmconfig   - Enable additional options for kvm guest kernel support
xenconfig   - Enable additional options for xen dom0 and guest kernel support
tinyconfig  - Configure the tiniest possible kernel

```

#### Other generic targets:

```

all - Build all targets marked with [*]
* vmlinux - Build the bare kernel
* modules - Build all modules
modules_install - Install all modules to INSTALL_MOD_PATH (default: /)
firmware_install - Install all firmware to INSTALL_FW_PATH
                  (default: $(INSTALL_MOD_PATH)/lib/firmware)
dir/ - Build all files in dir and below
dir/file.[ois] - Build specified target only
dir/file.lst - Build specified mixed source/assembly target only
              (requires a recent binutils and recent build (System.map))
dir/file.ko - Build module including final link
modules_prepare - Set up for building external modules
tags/TAGS - Generate tags file for editors
cscope - Generate cscope index
gtags - Generate GNU GLOBAL index
kernelrelease - Output the release version string (use with make -s)
kernelversion - Output the version stored in Makefile (use with make -s)
image_name - Output the image name (use with make -s)
headers_install - Install sanitised kernel headers to INSTALL_HDR_PATH
                 (default: ./usr)

```

#### Static analysers

```

checkstack - Generate a list of stack hogs
namespacecheck - Name space analysis on compiled kernel
versioncheck - Sanity check on version.h usage
includecheck - Check for duplicate included header files
export_report - List the usages of all exported symbols
headers_check - Sanity check on exported headers
headerdep - Detect inclusion cycles in headers
coccicheck - Check with Coccinelle.

```

## Kernel selftest

- kseltest - Build and run kernel selftest (run as root)  
Build, install, and boot kernel before  
running kseltest on it
- kseltest-clean - Remove all generated kseltest files
- kseltest-merge - Merge all the config dependencies of kseltest to existed  
.config.

## Kernel packaging:

- rpm-pkg - Build both source and binary RPM kernel packages
- binrpm-pkg - Build only the binary kernel RPM package
- deb-pkg - Build both source and binary deb kernel packages
- bindeb-pkg - Build only the binary kernel deb package
- tar-pkg - Build the kernel as an uncompressed tarball
- targz-pkg - Build the kernel as a gzip compressed tarball
- tarbz2-pkg - Build the kernel as a bzip2 compressed tarball
- tarxz-pkg - Build the kernel as a xz compressed tarball
- perf-tar-src-pkg - Build perf-4.10.0-rc2.tar source tarball
- perf-targz-src-pkg - Build perf-4.10.0-rc2.tar.gz source tarball
- perf-tarbz2-src-pkg - Build perf-4.10.0-rc2.tar.bz2 source tarball
- perf-tarxz-src-pkg - Build perf-4.10.0-rc2.tar.xz source tarball

## Documentation targets:

Linux kernel internal documentation in different formats (Sphinx):

- htmldocs - HTML
- latexdocs - LaTeX
- pdfdocs - PDF
- epubdocs - EPUB
- xmldocs - XML
- cleandocs - clean all generated files

make SPHINXDIRS="s1 s2" [target] Generate only docs of folder s1, s2  
valid values for SPHINXDIRS are: media core-api security admin-guide gpu process dev-  
tools driver-api doc-guide

make SPHINX\_CONF={conf-file} [target] use \*additional\* sphinx-build  
configuration. This is e.g. useful to build with nit-picking config.

Linux kernel internal documentation in different formats (DocBook):

- htmldocs - HTML
- pdfdocs - PDF
- psdocs - Postscript
- xmldocs - XML DocBook
- mandocs - man pages
- installmandocs - install man pages generated by mandocs
- cleandocs - clean all generated DocBook files

make DOCB00KS="s1.xml s2.xml" [target] Generate only docs s1.xml s2.xml  
valid values for DOCB00KS are: z8530book.xml kernel-hacking.xml kernel-locking.xml  
deviceiobook.xml writing\_usb\_driver.xml networking.xml kernel-api.xml filesystems.xml  
lsm.xml kgdb.xml gadget.xml libata.xml mtdnand.xml libs.xml rapidio.xml genericirq.xml

```
s390-drivers.xml uio-howto.xml scsi.xml sh.xml regulator.xml wl.xml
writing_musb_glue_layer.xml iio.xml
```

```
make DOCB00KS="" [target] Don't generate docs from Docbook
    This is useful to generate only the ReST docs (Sphinx)
```

Architecture specific targets (x86):

```
* bzImage      - Compressed kernel image (arch/x86/boot/bzImage)
install        - Install kernel using
                  (your) ~/bin/installkernel or
                  (distribution) /sbin/installkernel or
                  install to $(INSTALL_PATH) and run lilo
fdimage        - Create 1.4MB boot floppy image (arch/x86/boot/fdimage)
fdimage144     - Create 1.4MB boot floppy image (arch/x86/boot/fdimage)
fdimage288     - Create 2.8MB boot floppy image (arch/x86/boot/fdimage)
isoimage       - Create a boot CD-ROM image (arch/x86/boot/image.iso)
                  bzdisk/fdimage*/isoimage also accept:
                  FDARGS="..." arguments for the booted kernel
                  FDINITRD=file initrd for the booted kernel
```

```
i386_defconfig      - Build for i386
x86_64_defconfig    - Build for x86_64
```

```
make V=0|1 [targets] 0 => quiet build (default), 1 => verbose build
make V=2 [targets] 2 => give reason for rebuild of target
make O=dir [targets] Locate all output files in "dir", including .config
make C=1 [targets] Check all c source with $CHECK (sparse by default)
make C=2 [targets] Force check of all c source with $CHECK
make RECORDMCOUNT_WARN=1 [targets] Warn about ignored mcount sections
make W=n [targets] Enable extra gcc checks, n=1,2,3 where
    1: warnings which may be relevant and do not occur too often
    2: warnings which occur quite often but may still be relevant
    3: more obscure warnings, can most likely be ignored
    Multiple levels can be combined with W=12 or W=123
```

Execute "make" or "make all" to build all targets marked with [\*]  
 For further info see the ./README file  
 \$

### Additional TIP

The *make help* is **actually arch-specific**; an example for ARM – notice how the available architecture targets are shown:

**\$ make ARCH=arm help**

Cleaning targets:

```
clean      - Remove most generated files but keep the config and
              enough build support to build external modules
mrproper   - Remove all generated files + config + various backup files
distclean  - mrproper + remove editor backup and patch files
```

## Configuration targets:

- config - Update current config utilising a line-oriented program
- nconfig - Update current config utilising a ncurses menu based program
- menuconfig - Update current config utilising a menu based program

[...]

## Devicetree:

- \* dtbs - Build device tree blobs for enabled boards
- dtbs\_install - Install dtbs to /boot/dtbs/5.4.0
- dt\_binding\_check - Validate device tree binding documents
- dtbs\_check - Validate device tree source files

[...]

## Architecture specific targets (arm):

- \* zImage - Compressed kernel image (arch/arm/boot/zImage)
- Image - Uncompressed kernel image (arch/arm/boot/Image)
- \* xipImage - XIP kernel image, if configured (arch/arm/boot/xipImage)
- uImage - U-Boot wrapped zImage
- bootpImage - Combined zImage and initial RAM disk  
(supply initrd image via make variable INITRD=<path>)

[...]

- am200epdkit\_defconfig - Build for am200epdkit
- aspeed\_g4\_defconfig - Build for aspeed\_g4
- aspeed\_g5\_defconfig - Build for aspeed\_g5
- assabet\_defconfig - Build for assabet
- at91\_dt\_defconfig - Build for at91\_dt

[...]

- versatile\_defconfig - Build for versatile
- vexpress\_defconfig - Build for vexpress
- vf610m4\_defconfig - Build for vf610m4
- viper\_defconfig - Build for viper
- vt8500\_v6\_v7\_defconfig - Build for vt8500\_v6\_v7
- xcep\_defconfig - Build for xcep
- zeus\_defconfig - Build for zeus
- zx\_defconfig - Build for zx

make V=0|1 [targets] 0 => quiet build (default), 1 => verbose build

make V=2 [targets] 2 => give reason for rebuild of target

make O=dir [targets] Locate all output files in "dir", including .config

[...]

\$

**TIP***Generating the kernel in a package format*`make help`

...

Kernel packaging:

<code>rpm-pkg</code>	- Build both source and binary RPM kernel
<code>packages</code>	
<code>binrpm-pkg</code>	- Build only the binary kernel RPM package
<code>deb-pkg</code>	- Build both source and binary deb kernel
<code>packages</code>	
<code>bindeb-pkg</code>	- Build only the binary kernel deb package
<code>snap-pkg</code>	- Build only the binary kernel snap package (will connect to external hosts)

...

Say, you want the RPM package format for the source and binary kernel packages:

Requires *rpmbuild* :`sudo apt install rpm`

Generate them:

**`make [-jn] rpm-pkg`**

...

(For the Debian package format:

`make [-jn] deb-pkg`

...

).

Alternately, one can always supply a custom script that builds a custom package.

**TIP**

Attempting to compile an LKM (Loadable Kernel Module), requires the

`/lib/modules/$(uname -r)/build`

symlink to be correctly setup – pointing to the kernel source tree. Even if that's okay, the LKM build may fail with this message:

[...]

```
ERROR: Kernel configuration is invalid.
       include/generated/autoconf.h or include/config/auto.conf are missing.
       Run 'make oldconfig && make prepare' on kernel src to fix it.
```

[...]

Follow the instructions as shown above and proceed.



----- **TIP** -----

Accurately Measuring / Benchmarking kernel build times:

See this detailed post by Ingo Molnar, May 2018:

<https://lkml.org/lkml/2018/5/2/74>

----- **TIP** -----

**Configuring the Linux kernel as RTL – Real Time Linux – i.e. Linux as an RTOS!**

1. Download and extract the kernel source tree (here, as an example, we're using the 5.10.152 kernel):  
`wget`  
<https://mirrors.edge.kernel.org/pub/linux/kernel/v5.x/linux-5.10.152.tar.xz>
2. Download the kernel RTL patch for this source tree from here:  
<https://mirrors.edge.kernel.org/pub/linux/kernel/projects/rt/>

F.e.:

`wget https://mirrors.edge.kernel.org/pub/linux/kernel/projects/rt/5.10/patch-5.10.152-rt75.patch.xz`

**NOTE!**


- *The particular RTL kernel patch version may differ from what's shown here*
- *you'll need to exactly match the vanilla kernel source version and the RTL kernel patch version*

3. Extract it (in the dir above the 5.10.106 source)

4. Apply the patch...

```
cd linux-5.10.106
patch -p1 < ../patch-5.10.106-rt64.patch
(Verify it applies cleanly; can use the --dry-run switch to patch first to check)
```

5. Run `make menuconfig`; goto the *General Setup* → *Preemption Model* menu: select the last one – *Fully preemptible kernel (Real-Time)* : **CONFIG\_PREEMPT\_RT=y**



Make any other required mods to the kernel config, save and exit.

Build and boot from the kernel in the usual way; ***you're now running Linux as an RTOS!***

#### **TIP**

- ***A simple convenience script for kernel build (lightly tested... YMMV!):***

[https://github.com/PacktPublishing/Linux-Kernel-Programming\\_2E/blob/main/ch2/kbuild.sh](https://github.com/PacktPublishing/Linux-Kernel-Programming_2E/blob/main/ch2/kbuild.sh)

#### **TIP**

*How can I check which distro kernels are installed?*

Deb/Ubuntu:

```
dpkg --get-architecture | grep linux-image
```

RedHat/Fedora/...

```
rpm -qa kernel
```

or

```
dnf list installed kernel
```

Ref-

<https://www.cyberciti.biz/faq/howto-display-all-installed-linux-kernel-version/>

<< These are package-based kernels- they won't include 'manually' built ones of course! >>

---

**TIP**

---

**Change the default boot kernel (Ubuntu)**

<https://askubuntu.com/questions/216398/set-older-kernel-as-default-grub-entry>

GRUB bootloader (basic) customizations:  
(Materials below from my LKP 2E book):

## Customizing GRUB – the basics

Customizing GRUB is quite easy to do; we can always, as root, edit its config file: `/etc/default/grub`. Do note the following:

- The following steps are to be carried out on the “target” system itself (not on the host) – in our case, within the x86\_64 Ubuntu 22.04 guest VM. Of course, if working on Linux natively, you can go ahead on the same system.
- This procedure has been tested and verified on our x86\_64 Ubuntu 22.04 LTS guest system only.

So what exactly are we setting out to do here? We want GRUB to show us its menu at boot, before running our favorite OS, allowing us to customize it further. Here’s a quick series of steps to do this:

1. First, let’s be safe and keep a backup copy of the GRUB bootloader config file:

```
sudo cp /etc/default/grub /etc/default/grub.orig
```

2. Edit it. You can use `vi` or your editor of choice:

```
sudo vi /etc/default/grub
```

3. To always show the GRUB prompt at boot, insert this line:

```
GRUB_HIDDEN_TIMEOUT_QUIET=false
```

On some Linux distros, you might instead have the `GRUB_TIMEOUT_STYLE=hidden` directive; simply change it to `GRUB_TIMEOUT_STYLE=menu` to achieve the same effect. Always showing the bootloader menu at boot is good during development and testing and is typically turned off in production for both speed and security.



Talking about security, always ensure that access to the firmware (BIOS/UEFI) and bootloader is password-protected.

4. Set the timeout to boot the default OS (in seconds) as required. The default is 10 seconds; here, we set it to 3 seconds:

```
GRUB_TIMEOUT=3
```

Setting the preceding timeout value to the following values will produce the following outcomes:

- 0: Boot the system immediately without displaying the menu.
- -1: Wait indefinitely.

Furthermore, if a `GRUB_HIDDEN_TIMEOUT` directive is present in the config file, just comment it out:

```
#GRUB_HIDDEN_TIMEOUT=1
```

5. Finally, run the `update-grub` program as *root* to have your changes take effect:

```
sudo update-grub
```

The preceding command will typically cause the `initramfs` image to be refreshed (regenerated). Once done, you're ready to reboot the system. Hang on a second, though! The following section shows you how you can modify GRUB's configuration to boot by default into a kernel of your choice.

...

This may not be what we want; as a real example, on our x86 Ubuntu 22.04 LTS guest VM, we can set it to the default Ubuntu *distro kernel* by, as earlier, editing the `/etc/default/grub` file (as root, of course), like so:

```
GRUB_DEFAULT="Advanced options for Ubuntu>Ubuntu, with Linux 5.19.0-43-generic"
```



Of course, this implies that if your distro is updated or upgraded, you must again manually change the preceding line to reflect the new distro kernel that you wish to boot into by default, and then run `sudo update-grub`.

Right, our freshly edited GRUB configuration file is shown as follows:

```
$ cat /etc/default/grub
[...]
#GRUB_DEFAULT=0
GRUB_DEFAULT="Advanced options for Ubuntu>Ubuntu, with Linux 5.19.0-43-generic"
#GRUB_TIMEOUT_STYLE=hidden
GRUB_HIDDEN_TIMEOUT_QUIET=false
GRUB_TIMEOUT=3
GRUB_DISTRIBUTOR=`lsb_release -i -s 2> /dev/null || echo Debian`
GRUB_CMDLINE_LINUX_DEFAULT="quiet splash"
GRUB_CMDLINE_LINUX="quiet splash"
[...]
```

As in the previous section, don't forget: if you make any changes here, run the `sudo update-grub` command to have your changes take effect.

...

To see the GRUB menu at boot:

- Depending on your UEFI (or legacy BIOS), the boot order can differ; so press F12 (or perhaps F1/F2) to see all available options
- Fedora: When Linux is selected to boot, keep the Shift key(s) pressed; only then does the GRUB menu seem to appear (in spite of the changes done above!)

*Fedora:*

No `update-grub` command; the equivalent is this:

```
# grub2-mkconfig --output=/boot/grub2/grub.cfg
```

----- **TIP** -----

**Security** for the Linux kernel:

[Ref: How To Secure the Linux Kernel, B Day, Jan 2022](#)

----- **TIP** -----

- Finally, remember, an (almost) guaranteed way to succeed :-)

When you get those build / boot errors, etc... that you cannot fix: copy the exact error message into the clipboard, go to Google, type “linux kernel build fails with: <paste-the-error-message>” !

*One might be surprised at how often this helps ;-)*

If not, post your (well thought-out) question.

---

***Actual option switches, flags used by gcc when building the Linux kernel for an ARMv7 (Cortex-A9) ARM Versatile Express platform:***

*Toolchain is from Linaro for ARM (Aarch32):*

```
$ arm-linux-gnueabi-gcc --version
arm-linux-gnueabi-gcc (Linaro GCC 7.3-2018.05) 7.3.1 20180425 [linaro-7.3-
2018.05 revision d29120a424ecfbc167ef90065c0eeb7f91977701]
Copyright (C) 2017 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

$
```

```
$ make V=1 -j8 ARCH=arm CROSS_COMPILE=arm-linux-gnueabi- all
```

```
[...]
```

<< *Note:*

*a) this is not a 'debug' build*

*b) We've inserted newlines into the output stream below to make it more human-readable*

>>

```
arm-linux-gnueabi-gcc
-Wp,-MD,arch/arm/kernel/.sys_arm.o.d -nostdinc
```

```

-isystem
<...>/gcc-linaro-7.3.1-2018.05-x86_64_arm-linux-gnueabihf/bin/./lib/gcc/arm-
linux-gnueabihf/7.3.1/include
-I./arch/arm/include -Iarch/arm/include/generated/uapi
-Iarch/arm/include/generated -Iinclude -I./arch/arm/include/uapi
-Iarch/arm/include/generated/uapi -I./include/uapi
-Iinclude/generated/uapi
-include ./include/linux/kconfig.h
-D __KERNEL__
-mlittle-endian
-Wall -Wundef -Wstrict-prototypes -Wno-trigraphs
-fno-strict-aliasing -fno-common -Werror-implicit-function-declaration
-Wno-format-security -std=gnu89 -fno-dwarf2-cfi-asm -fno-ipa-sra
-mabi=aapcs-linux -mno-thumb-interwork -mfpv=vfp -funwind-tables
-marm -D__LINUX_ARM_ARCH__=7 -march=armv7-a -msoft-float -Uarm
-fno-delete-null-pointer-checks -fno-PIE
-02
-param=allow-store-data-races=0 -DCC_HAVE_ASM_GOTO
-Wframe-larger-than=1024 -fno-stack-protector -Wno-unused-but-set-variable
-Wno-unused-const-variable -fomit-frame-pointer
-fno-var-tracking-assignments
-g -Wdeclaration-after-statement -Wno-pointer-sign -fno-strict-overflow
-fconserve-stack -Werror=implicit-int -Werror=strict-prototypes
-Werror=date-time -D"KBUILD_STR(s)=#s"
-D"KBUILD_BASENAME=KBUILD_STR(sys_arm)"
-D"KBUILD_MODNAME=KBUILD_STR(sys_arm)"
-c -o arch/arm/kernel/sys_arm.o arch/arm/kernel/sys_arm.c

[...]
```

### *Below is a 'debug' build*

```

<...>/gcc-linaro-7.3.1-2018.05-x86_64_arm-linux-gnueabihf/bin/./libexec/gcc/
arm-linux-gnueabihf/7.3.1/cc1
-quiet
-nostdinc
-I ./arch/arm/include
-I ./arch/arm/include/generated
-I ./include
-I ./arch/arm/include/uapi
-I ./arch/arm/include/generated/uapi
-I ./include/uapi
-I ./include/generated/uapi
-imultilib .
-multiarch arm-linux-gnueabihf
-iprefix
<...>/gcc-linaro-7.3.1-2018.05-x86_64_arm-linux-gnueabihf/bin/./lib/gcc/arm-
linux-gnueabihf/7.3.1/
```

```

-isysroot
<...>/gcc-linaro-7.3.1-2018.05-x86_64_arm-linux-gnueabihf/bin/./arm-linux-
gnueabihf/libc
-D __KERNEL__
-D __LINUX_ARM_ARCH__=7
-U arm
-D CC_HAVE_ASM_GOTO
-D KBUILD_BASENAME="ioctl"
-D KBUILD_MODNAME="ioctl"
-isystem
<...>/gcc-linaro-7.3.1-2018.05-x86_64_arm-linux-gnueabihf/bin/./lib/gcc/arm-
linux-gnueabihf/7.3.1/include
-include ./include/linux/kconfig.h
-MD block/.ioctl.o.d block/ioctl.c
-quiet
-dumpbase ioctl.c
-mlittle-endian -mapcs -mno-sched-prolog
-mabi=aapcs-linux -mno-thumb-interwork -mfpv=vfp
-marm -march=armv7-a -mfloat-abi=soft -mtune=cortex-a9
-mtls-dialect=gnu
-auxbase-strip block/ioctl.o
-g -gdwarf-4 -O2 -Wall
-Wundef -Wstrict-prototypes -Wno-trigraphs
-Werror=implicit-function-declaration -Wno-format-security -Wno-frame-address
-Wformat-truncation=0 -Wformat-overflow=0 -Wno-int-in-bool-context
-Wframe-larger-than=1024 -Wno-unused-but-set-variable
-Wunused-const-variable=0 -Wdeclaration-after-statement -Wno-pointer-sign
-Werror=implicit-int -Werror=strict-prototypes -Werror=date-time
-Werror=incompatible-pointer-types -Werror=designated-init
-std=gnu90 -p -fno-strict-aliasing -fno-common -fshort-wchar -fno-PIE
-fno-dwarf2-cfi-asm -fno-ipa-sra -funwind-tables
-fno-delete-null-pointer-checks -fno-reorder-blocks -fno-ipa-cp-clone
-fno-partial-inlining -fstack-protector -fno-omit-frame-pointer
-fno-optimize-sibling-calls -fno-var-tracking-assignments -fno-strict-
overflow
-fno-merge-all-constants -fmerge-constants -fstack-check=no -fconserve-stack
--param allow-store-data-races=0 -o /tmp/cc0Ujck.s

```

[...]

<< Another build run ... >>

### arm-linux-gnueabihf-gcc

```

-Wp,-MD,arch/arm/kernel/.process.o.d -nostdinc
-isystem
<...>/gcc-linaro-7.2.1-2017.11-x86_64_arm-linux-gnueabihf/bin/./lib/gcc/arm-
linux-gnueabihf/7.2.1/include
-I./arch/arm/include -I./arch/arm/include/generated
-I./include -I./arch/arm/include/uapi -I./arch/arm/include/generated/uapi
-I./include/uapi -I./include/generated/uapi

```



```
-include ./include/linux/kconfig.h
-D__KERNEL__
-mlittle-endian -Wall -Wundef -Wstrict-prototypes -Wno-trigraphs
-fno-strict-aliasing -fno-common -fshort-wchar
-Werror-implicit-function-declaration -Wno-format-security -std=gnu89
-fno-PIE -fno-dwarf2-cfi-asm -fno-omit-frame-pointer -mapcs -mno-sched-prolog

-fno-ipa-sra -mabi=aapcs-linux -mno-thumb-interwork -mfpu=vfp -funwind-tables
-marm -D__LINUX_ARM_ARCH__=7 -march=armv7-a -msoft-float -Uarm
-fno-delete-null-pointer-checks -Wno-frame-address -Wno-format-truncation
-Wno-format-overflow -Wno-int-in-bool-context
-O2 --param=allow-store-data-races=0 -DCC_HAVE_ASM_GOTO -fno-reorder-blocks
-fno-ipa-cp-clone -fno-partial-inlining -Wframe-larger-than=1024
-fstack-protector -Wno-unused-but-set-variable -Wno-unused-const-variable
-fno-omit-frame-pointer -fno-optimize-sibling-calls
-fno-var-tracking-assignments
-g -gdwarf-4 -pg -Wdeclaration-after-statement -Wno-pointer-sign
-fno-strict-overflow -fno-merge-all-constants -fmerge-constants
-fno-stack-check -fconserve-stack -Werror=implicit-int
-Werror=strict-prototypes -Werror=date-time -Werror=incompatible-pointer-
types
-Werror=designated-init -DKBUILD_BASENAME='"process"'
-DKBUILD_MODNAME='"process"' -c -o arch/arm/kernel/process.o
arch/arm/kernel/process.c
...
```

---

## Kernel custom configuration -OR- Setting up your own menu entries in kernel build

The Kconfig file, within each source folder, is the relevant one . Each source folder has a *Kconfig*:

*init/Kconfig* - defines the 'General Setup' menu items!

*Lets explain with an example:*

We'd like to add a kernel menu option under the "General Setup" menu:

```
"[ ] My Amazing Kernel Feature"
```

By default it should be OFF.

When turned ON, the effect will be that we compile the kernel with some additional gcc switches (details below).

How is this setup?

1. Edit the *init/Kconfig* file

(why this one? -because the "General Setup" menu items are defined here. As another eg, the "Kernel Hacking" menu is defined in *lib/Kconfig.debug* ).

1.1 Add the following paragraph to set things up:

```
config AMAZING
    bool "My Amazing Kernel Feature"
    default n
    help
        Turns on the hook that will cause this kernel to ...
        blah blah blah
    ...
```

1.2 Save & exit

## 2. Run 'make menuconfig'

### 2.1 Goto the General Setup menu

### 2.2 You should now see a new entry - the one just created, like this:

```
[ ] My Amazing Kernel Feature
```

### 2.3 Explore, change it, see the help..

### 2.4 Once done, save & exit (the config).

## 3. The resulting .config file will reflect whether the user selected the new entry or not:

```
$ grep CONFIG_AMAZING .config
```

If turned ON, the result of the grep above will be:

```
CONFIG_AMAZING=y
```

If turned OFF (the default in our example), the result of the grep above will be:

```
# CONFIG_AMAZING is not set
```

Additionally, the corresponding "CONFIG\_FOO" define is auto-generated in the *include/generated/autoconf.h* header: (example below shows when it was selected (ON)):

```
<...>/include/generated/autoconf.h:49:#define CONFIG_AMAZING 1
```

## 4. Edit the (toplevel or other) Makefile to figure action based on our new directive:

... <here we inserted these lines into the toplevel Makefile>

```
ifdef CONFIG_AMAZING
KBUILD_CFLAGS += -finstrument-functions -g
endif
```

...

So, if the "My Amazing Kernel Feature" option is used, all C source files will now be compiled with the additional "-finstrument-functions -g" gcc switches. If left unselected (during the kernel config step), these switches will not be used during compilation.

Cool!

Proof that it works: :-)

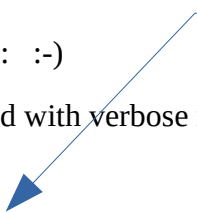
See the kernel build with verbose mode on (using `make V=1 zImage ...`):

...

```
arm-buildroot-linux-uclibcgnueabi-gcc -Wp,-MD,block/.elevator.o.d -nostdinc
```

[...]

```
-Wframe-larger-than=1024 -fno-stack-protector -Wno-unused-but-set-variable
-fno-omit-frame-pointer -fno-optimize-sibling-calls -g -finstrument-
functions -g -Wdeclaration-after-statement -Wno-pointer-sign -fno-strict-
overflow -fconserve-stack -DCC_HAVE_ASM_GOTO -D"KBUILD_STR(s)=#s" -
D"KBUILD_BASENAME=KBUILD_STR(elevator)" -
D"KBUILD_MODNAME=KBUILD_STR(elevator)" -c -o block/.tmp_elevator.o
block/elevator.c
```



Using a cross-compiler here

## TIP

The kernel config is sometimes available within the proc filesystem, as the pseudofile `/proc/config.gz` – a gzip'ed config file.

If it doesn't show up, try this:

```
# modprobe configs
# ls -l /proc/config.gz
-r--r--r-- 1 root root 44216 Jun 23 10:13 /proc/config.gz
# zcat /proc/config.gz |grep INITRD
...
```

&lt;&lt;

For working with Git in general, (and with the Linux kernel upstream development in particular), please refer to the **‘Git – The Basics’** PDF tutorial.  
>>

### ***The “kbuild Test Robot” - employed by the kernel community***

From:

0-DAY kernel test infrastructure      Open Source Technology Center  
<https://lists.01.org/>      Intel Corporation

An example:

## **[kernel-hardening] [PATCH 4/6] Protectable Memory**

---

**kbuild test robot** <lkp@intel.com>

Fri, Feb 2, 2018 at 11:11 AM

To: Igor Stoppa <igor.stoppa@huawei.com>

Cc: kbuild-all@01.org, jglisse@redhat.com, keescook@chromium.org, mhocko@kernel.org, labbott@redhat.com, hch@infradead.org, willy@infradead.org, cl@linux.com, linux-security-module@vger.kernel.org, linux-mm@kvack.org, linux-kernel@vger.kernel.org, kernel-hardening@lists.openwall.com, Igor Stoppa <igor.stoppa@huawei.com>

Hi Igor,

Thank you for the patch! Perhaps something to improve:

[auto build test WARNING on linus/master]

[also build test WARNING on v4.15]

[cannot apply to next-20180201]

[if your patch is applied to the wrong git tree, please drop us a note to help improve the system]

url: <https://github.com/0day-ci/>

config: i386-randconfig-x071-201804 (attached as .config)

compiler: gcc-7 (Debian 7.2.0-12) 7.2.1 20171025

reproduce:

```
# save the attached .config to linux build tree
make ARCH=i386
```

All warnings (new ones prefixed by >>):

```
mm/pmalloc.c: In function 'pmalloc_pool_show_avail':
>> mm/pmalloc.c:71:25: warning: format '%lu' expects argument of type 'long unsigned int', but
argument 3 has type 'size_t {aka unsigned int}' [-Wformat=]
    return sprintf(buf, "%lu\n", gen_pool_avail(data->pool));
```

```

    ~^  ~~~~~
    %u
mm/pmalloc.c: In function 'pmalloc_pool_show_size':
mm/pmalloc.c:81:25: warning: format '%lu' expects argument of type 'long unsigned int', but
argument 3 has type 'size_t {aka unsigned int}' [-Wformat=]
    return sprintf(buf, "%lu\n", gen_pool_size(data->pool));
    ~^  ~~~~~
    %u

```

vim +71 mm/pmalloc.c

```

63
64 static ssize_t pmalloc_pool_show_avail(struct kobject *dev,
65                                         struct kobj_attribute *attr,
66                                         char *buf)
67 {
68     struct pmalloc_data *data;
69
70     data = container_of(attr, struct pmalloc_data, attr_avail);
> 71     return sprintf(buf, "%lu\n", gen_pool_avail(data->pool));
72 }
73

```

---
0-DAY kernel test infrastructure Open Source Technology Center
<https://lists.01.org/> Intel Corporation

<< Attached: the .config.gz >>

**Initramfs / initrd**

**Rationale behind initrd**

An ‘initrd’ – now called *initramfs* – is a block device in RAM (a ramdisk).

initrd – initial RAM disk / initramfs :

Difference between initrd and initramfs?
<ul style="list-style-type: none"><li>• initrd is the older style tech</li><li>• to build an initrd image with cur dir content, can do: find .   sudo cpio -R root:root   gzip -9 &gt; initrd.img</li><li>• to build an initramfs image with cur dir content, can do (uses the ‘newc’ format):</li></ul>


```
find . | sudo cpio -o --format=newc -R root:root | gzip -9 >
initramfs.img
```

mount “/” : kernel at boot  
... init ... ; runs /sbin/init

Kernel must ‘understand’ the root fs : implies it has the filesystem driver !


*ext4* ; fs ‘driver’

Builtin:



mount *will succeed*

Else, it's a kernel module:



Now, mount of rootfs will *not* succeed until the ext4 kernel module is loaded into RAM.

insmod / modprobe <pathname.ko>

BUT: the kernel module is here : /lib/modules/4.14.183/ - this location is **within** the root filesystem !

In order to mount the rootfs, we require the *ext4.ko* module  
BUT it's in the rootfs !

*Thus we have a classic 'Chicken and Egg' problem !!!*

*How to solve this??*



**initramfs !**

*Recall, an 'initrd' – now called initramfs –  
is a block device in RAM (a ramdisk).*

Contents of / (rootfs) are scaled-down;  
/bin; /dev/ ; /etc/ ; /lib ; /usr ; ...

Essentially:

- the kernel mounts initramfs as a temporary rootfs (in RAM)
- scripts within load up all required drivers etc
- kernel then mounts the real root filesystem
- once done, it unmounts the temporary initramfs.





&lt;&lt;

From the excellent material from [Bootlin](#):



## Root filesystem in memory: *initramfs*

It is also possible to boot the system with a filesystem in memory: *initramfs*

- ▶ Either from a compressed CPIO archive integrated into the kernel image
- ▶ Or from such an archive loaded by the bootloader into memory
- ▶ At boot time, this archive is extracted into the Linux file cache
- ▶ It is useful for two cases:
  - Fast booting of very small root filesystems. As the filesystem is completely loaded at boot time, application startup is very fast.
  - As an intermediate step before switching to a real root filesystem, located on devices for which drivers not part of the kernel image are needed (storage drivers, filesystem drivers, network drivers). This is always used on the kernel of desktop/server distributions to keep the kernel image size reasonable.
- ▶ Details (in kernel documentation):  
[filesystems/ramfs-rootfs-initramfs](#)



## External *initramfs*

- ▶ To create one, first create a compressed CPIO archive:

```
cd rootfs/  
find . | cpio -H newc -o > ../initramfs.cpio  
cd ..  
gzip initramfs.cpio
```
- ▶ If you're using U-Boot, you'll need to include your archive in a U-Boot container:

```
mkimage -n 'Ramdisk Image' -A arm -O linux -T ramdisk -C gzip \  
-d initramfs.cpio.gz uInitramfs
```
- ▶ Then, in the bootloader, load the kernel binary, DTB and *uInitramfs* in RAM and boot the kernel as follows:

```
bootz kernel-addr initramfs-addr dtb-addr
```



## Built-in initramfs

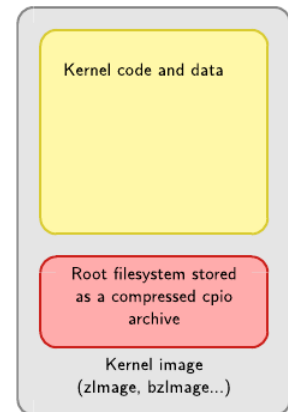
To have the kernel Makefile include an initramfs archive in the kernel image: use the `CONFIG_INITRAMFS_SOURCE` option.

- ▶ It can be the path to a directory containing the root filesystem contents
- ▶ It can be the path to a ready made cpio archive
- ▶ It can be a text file describing the contents of the initramfs

See the kernel documentation for details:

[driver-api/early-userspace/early\\_userspace\\_support](#)

**WARNING:** only binaries from GPLv2 compatible code are allowed to be included in the kernel binary using this technique. Otherwise, use an external initramfs.



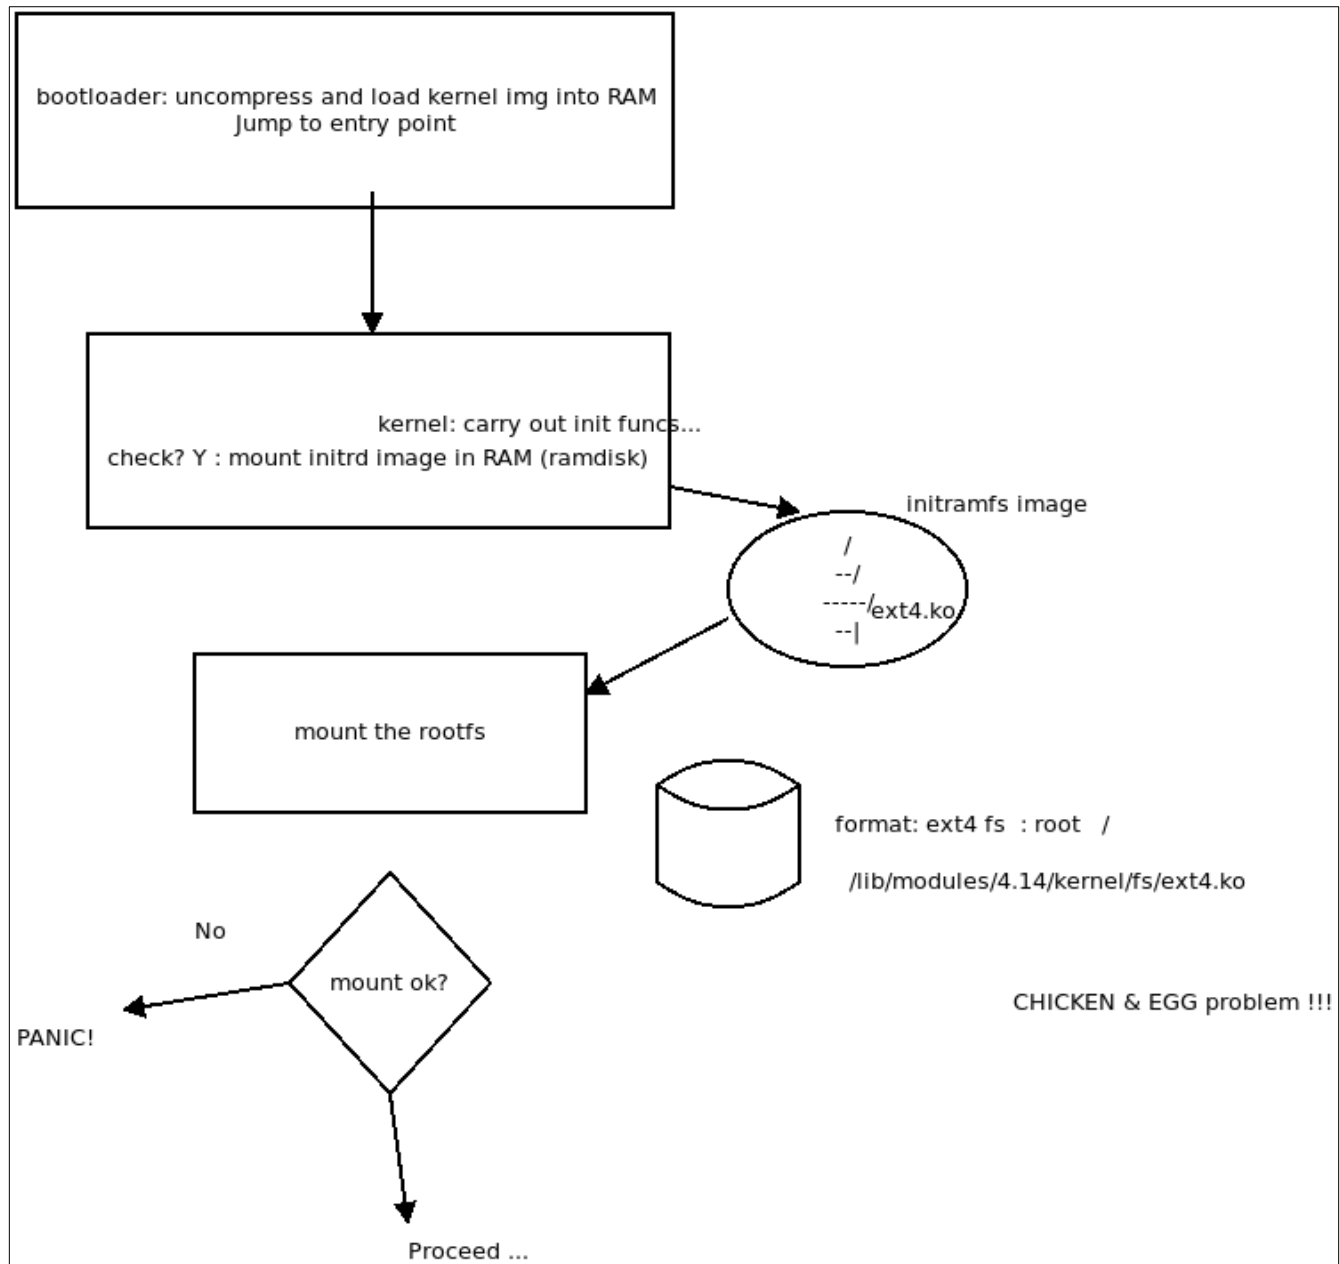
>>

From the “Kernel Rebuild Guide” by Kwan Lowe, Digital Hermit  
[ <http://www.digitalhermit.com/linux/Kernel-Build-HOWTO.html> ] :

If you have built your main boot drivers as modules (e.g., SCSI host adapter, filesystem, RAID drivers) then you will need to create an initial RAMdisk image. The `initrd` is a way of sidestepping the chicken and egg problem of booting -- drivers are needed to load the root filesystem but the filesystem cannot be loaded because the drivers are on the filesystem. As the manpage for **mkinitrd** states:

**mkinitrd** creates filesystem images which are suitable for use as Linux initial ramdisk (`initrd`) images. Such images are often used for preloading the block device modules (such as IDE, SCSI or RAID) which are needed to access the root filesystem. `mkinitrd` automatically loads filesystem modules (such as `ext3` and `jbd`), IDE modules, all `scsi_hostadapter` entries in `/etc/modules.conf`, and raid modules if the systems root partition is on raid, which makes it simple to build and use kernels using modular device drivers.

<< P.T.O. >>



### Initramfs is very useful for stuff like:

- running an app *before* the kernel fully comes up (for a **time critical app** too!)
  - get password for encrypted block device(s)
    - `man systemd-ask-password-console.service`  
 “systemd-ask-password-console.service is a system service that queries the user for system passwords (such as hard disk encryption keys and SSL certificate passphrases) on the console. It is intended to be used during boot to ensure proper handling of passwords necessary for boot. ...”  
[more details here](#)

- set console font
- set keyboard map
- sometimes *only* an initrd image is required
  - no / WORM storage device (kiosk, etc)
  - kdump (second) kernel – boots into an initrd environment (simply to send the `/proc/vmcore` file over the network to a server system)

Also see a good article on unpacking, changing and packing back an initrd image here:  
<http://www.alexonlinux.com/opening-and-modifying-the-initrd>

From the LKP 2E book:

At boot, and, of course, assuming the `initramfs` feature is enabled, the bootloader will, as part of its work, perform an extra step: after uncompressing and loading the kernel image in RAM, it will also load the specified `initramfs` image file into RAM. Now, when the kernel runs and detects the presence of the `initramfs` image, it uncompresses it, and using its content (via scripts), it loads up the required kernel modules into RAM (see Figure 3.2):

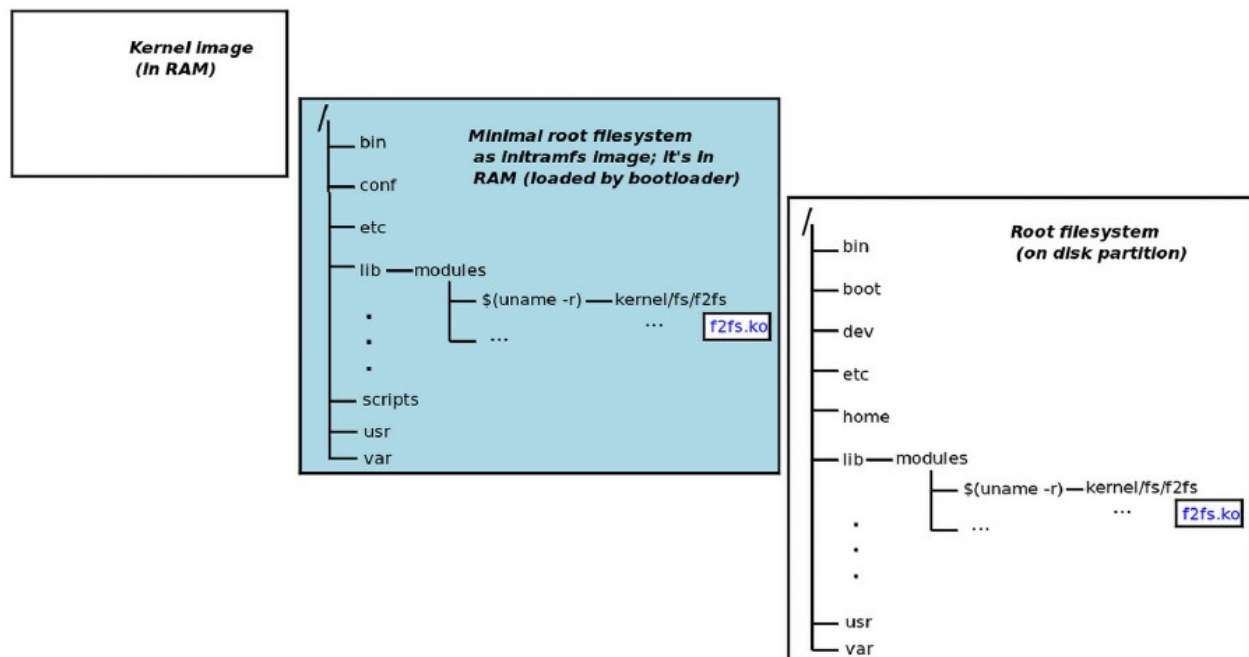


Figure 3.2: The `initramfs` image serves as a middleman between early kernel and actual root filesystem availability

...

***Listing the initramfs contents***

*[On Fedora 27 x86\_64 ; on Ubuntu, use `lsinitramfs !`]*

```
$ sudo lsinitrd
```

```
[sudo] password for xxx:
```

```
Image: /boot/initramfs-4.15.10-300.fc27.x86_64.img: 21M
```

```
=====
```

```
Early CPIO image
```

```
=====
```

drwxr-xr-x	3	root	root	0	Jan	5	14:48	.
-rw-r--r--	1	root	root	2	Jan	5	14:48	early_cpio
drwxr-xr-x	3	root	root	0	Jan	5	14:48	kernel
drwxr-xr-x	3	root	root	0	Jan	5	14:48	kernel/x86
drwxr-xr-x	2	root	root	0	Jan	5	14:48	kernel/x86/microcode
-rw-r--r--	1	root	root	99328	Jan	5	14:48	

```
kernel/x86/microcode/GenuineIntel.bin
```

```
=====
```

```
Version: dracut-046-8.git20180105.fc27
```

```
Arguments: -f
```

```
dracut modules:
```

```
bash
```

```
systemd
```

```
systemd-initrd
```

```
[...]
```

```
shutdown
```

```
=====
```

drwxr-xr-x	12	root	root	0	Jan	5	14:48	.
crw-r--r--	1	root	root	5,	1	Jan	5	14:48
crw-r--r--	1	root	root	1,	11	Jan	5	14:48
crw-r--r--	1	root	root	1,	3	Jan	5	14:48
crw-r--r--	1	root	root	1,	8	Jan	5	14:48
crw-r--r--	1	root	root	1,	9	Jan	5	14:48
lrwxrwxrwx	1	root	root	7	Jan	5	14:48	bin -> usr/bin
drwxr-xr-x	2	root	root	0	Jan	5	14:48	dev
drwxr-xr-x	11	root	root	0	Jan	5	14:48	etc
drwxr-xr-x	2	root	root	0	Jan	5	14:48	etc/cmdline.d
drwxr-xr-x	2	root	root	0	Jan	5	14:48	etc/conf.d
-rw-r--r--	1	root	root	124	Jan	5	14:48	etc/conf.d/systemd.conf
-rw-r--r--	1	root	root	303	Jan	5	14:48	etc/dhclient.conf

```
...
```

-rw-r--r--	1	root	root	1377	Aug	4	2017	usr/share/terminfo/v/vt220
lrwxrwxrwx	1	root	root	20	Jan	5	14:48	usr/share/unimaps ->
/usr/lib/kbd/unimaps								
drwxr-xr-x	3	root	root	0	Jan	5	14:48	var
lrwxrwxrwx	1	root	root	11	Jan	5	14:48	var/lock -> ../run/lock
lrwxrwxrwx	1	root	root	6	Jan	5	14:48	var/run -> ../run
drwxr-xr-x	2	root	root	0	Jan	5	14:48	var/tmp

=====

\$

*Script to extract and thus see initramfs content (on Ubuntu)*


Note: the newer version of mkinitrd is *mkinitramfs*.

A good article on [Initramfs](#).

### **kaiwanTECH Linux OS Corporate Training Programs**

Please do check out our current offering of world-class, seriously-valuable, high on returns, technical Linux OS corporate training programs here: <https://kaiwantech.com>

**Linux Operating System  
Specialized**



The highest quality Training on:

Linux Fundamentals, CLI and Scripting  
Linux Systems Programming  
Linux Kernel Internals  
Linux Device Drivers  
Embedded Linux  
Linux Debugging Techniques  
**New! Linux OS for Technical Managers**

Please do visit our website for details:  
<http://kaiwantech.in>

